

# Interprocess Communication (IPC)

Computer Operating Systems  
BLG 312E

2014-2015 Spring

## Types of Interaction

- between concurrent processes
  - resource sharing
  - communication
  - synchronization

## Levels of Interaction

- interaction between processes on three levels
  - processes not aware of each other (competing): using system resources (moderated by operating system)
  - processes indirectly aware of each other (sharing): resource sharing through mutual exclusion and synchronization
  - processes directly aware of each other (communicating)

## Resource Sharing


- mutual exclusion
  - two types of resources
    - can be used by more than one process at a time (e.g. reading from a file)
    - can be used by only one process at a time
      - due to physical constraints (e.g. some I/O units)
      - if the actions of one process interferes with those of another (e.g. writing to a shared memory location)
- synchronization
  - a process needs to proceed after another process completes some actions

## Example

2 processes: Observer and Reporter  
counter shared variable

```
observer:                reporter:
while TRUE {              while TRUE {
  observe;                 print_counter;
  counter ++;              counter=0;
}
```

## Example – Possible Errors



## Example – Possible Errors

```
counter++  LOAD ACC, COUNTER
           INC  ACC
           SAVE COUNTER, ACC
```

### Race:

- when processes access a shared variable
- outcome depends on order and running speed of processes
- may be different for different runs

## Example – Possible Errors

**P1:** `k=0 (initial value)`  
`while TRUE`  
`k=k+1;`  
**what about the values of k depending on the order of P1 and P2 executions?**

**P2:** `while TRUE`  
`k=k+1;`  
**SOLUTION: mutual exclusion**

## Sharing

- two types of sharing:
  - READ (no need for mutual exclusion)
  - WRITE (mutual exclusion needed)
- for consistency
  - mutual exclusion
  - synchronization

## Synchronization

- programs should not be dependent on running order of processes
- programs working together may need to be synchronized at some points
  - e.g. a program uses output calculated by another program

## Mutual Exclusion

**critical section (CS):** Part of code in a process in which operations on shared resources are performed.

**mutual exclusion:** only one process can execute a CS for a resource at a time


## Example

<p><b><u>P1:</u></b></p> <pre>while TRUE {   &lt;non-CS&gt;   mx_begin   &lt;CS ops&gt;   mx_end   &lt;non-CS&gt; }</pre>	<p><b><u>P2:</u></b></p> <pre>while TRUE {   &lt;non-CS&gt;   mx_begin   &lt;CS ops&gt;   mx_end   &lt;non-CS&gt; }</pre>
---	---

## Mutual Exclusion – Possible Problems

- deadlock
  - more than one process requires the same resources
  - each process does not release the resource required by the other

**Example:** 3 processes and 3 resources



P1() P2() P3()  
 req(R1); req(R2); req(R3);  
 req(R2); req(R3); req(R1);

## Mutual Exclusion

- `mx_begin`
  - is there a process in its CS which has not yet executed `mx_end`?
  - if NOT
    - allow process to proceed into CS
    - leave mark for other processes
- `mx_end`
  - allow any process waiting to go into its CS to proceed
  - if there aren't any, then leave mark (empty)

## Mutual Exclusion Implementation

- only one process may be in its CS
- if a process wants to enter its CS and if there are no others executing their CS, it shouldn't wait
- any process not executing its CS should not prevent another process from entering its own CS
- no assumptions should be made about the order and speed of execution of processes
- no process should stay in its CS indefinitely
- no process should wait to enter its own CS indefinitely

## Mutual Exclusion Solutions

- software based solutions
- hardware based solutions
- software and hardware based solutions

## A Software Based Solution

- use a shared flag that shows whether a process is in its CS or not: *busy*  
`busy ← TRUE` : process in CS  
`busy ← FALSE` : no process in CS
- `mx_begin`:
  - `while (busy);`
  - `busy = TRUE;`
  - wait until process in CS is finished
  - enter CS
- `mx_end`: `busy = FALSE;`

## A Software Based Solution

- a possible error
  - *busy* is also a shared variable!
  - Example:
    - P1 checks and finds `busy=FALSE`
    - P1 interrupted
    - P2 checks and finds `busy=FALSE`
    - both P1 and P2 enter CS

## Solutions Requiring *Busy Waiting*

```

shared variable turn = 1;

Process 1:      Process 2:
local variables  local variables
my_turn=1;        my_turn=2;
others_turn=2;    others_turn=1;

mx_begin: while (turn != my_turn);
mx_end  : turn = others_turn;
    
```

## Solutions Requiring *Busy Waiting*

- use up CPU time
- works properly but has limitations:
  - processes enter their CS in turn
  - depends on speed of process execution
  - depends on number of processes

## Solutions Requiring *Busy Waiting*

- first correct solution: Dekker algorithm
- Peterson algorithm (1981)
  - similar approach
  - simpler

## Peterson Algorithm

- shared variables:
  - req\_1, req\_2: bool and initialized to FALSE
  - turn: integer and initialized to "P1" or "P2"

```

P1:      P2:
mx_begin:  mx_begin:
req_1 = TRUE;  req_2 = TRUE;
turn = P2;    turn = P1;
while (req_2 && turn==P2);  while (req_1 && turn==P1);
< CS >      < CS >
mx_end: req_1 = FALSE;  mx_end: req_2 = FALSE;
    
```

## Peterson Algorithm

- different scenarios:
  - P1 is active, P2 is passive
    - req\_1=TRUE and turn=P2
    - req\_2=FALSE so P1 proceeds after while loop
  - P1 in CS, P2 wants to enter CS
    - req\_2=TRUE and turn=P1;
    - req\_1=TRUE so P2 waits in while loop
    - P2 continues after P1 executes mx\_end

## Peterson Algorithm

- (*different scenarios cntd.*):
    - P1 and P2 want to enter CS at the same time
- ```

P1:      P2:
req_1=TRUE;  req_2=TRUE;
turn=P2;    turn=P1;

⇒ order depends on which process assigns value to the
   turn variable first.
    
```

## Hardware Based Solutions

- with uninterruptable machine code instructions completed in one machine cycle
  - e.g.: test\_and\_set
    - busy waiting used
    - when a process exits CS, no mechanism to determine which other process enters next
      - indefinite waiting possible
- disabling interrupts
  - interferes with scheduling algorithm of operating system

## Hardware Based Solutions

```
test_and_set(a):  cc ← a
                  a ← TRUE
```

- with one machine instruction, contents of “a” copied into condition code register and “a” is assigned TRUE

```
mx_begin: test_and_set(busy);
           while (cc) {
             test_and_set(busy);
           }
mx_end:   busy=FALSE;
```

busy: shared variable  
cc: local condition code

## Semaphores

- hardware and software based solution
- no busy waiting
- does not waste CPU time
- **semaphore** is a special variable
  - only access through using two special operations
  - special operations cannot be interrupted
  - operating system carries out special operations

## Semaphores

- s: semaphore variable
  - special operations:
    - P (wait): when entering CS: mutex\_begin
    - V (signal): when leaving CS: mutex\_end
- ```
P(s) :
if (s > 0)
    s=s-1;
else
    wait_on_s;

V(s) :
if (anyone_waiting_on_s)
    activate_next_in_line;
else
    s=s+1;
```

## Semaphores

- take on integer values ( $\geq 0$ )
- created through a special system call
- assigned an initial value
- binary semaphore:
  - can be 0/1
  - used for CS
- counting semaphore:
  - can be integers  $\geq 0$

## Example: Observer – Reporter

```
global shared variables:
counter: integer;
sem: semaphore;

process P1:
    observe;
    P(sem);
    counter++;
    V(sem);
    ....

process P2:
    ...
    P(sem);
    print(counter);
    counter=0;
    V(sem);

main_program:
    sem=1; counter=0;
    activate(P1);
    activate(P2);
```

## Example: Observer – Reporter

sample run:

P1: P(sem) ... sem=0;

P2: P(sem) ... sem=0 so P2 is suspended

P1: V(sem) ... P2 is waiting for sem; activate P2

P2: V(sem) ... no one waiting; sem=1

## Synchronization with Semaphores

- a process may require an event to proceed – process is suspended
    - e.g. process waiting for input
  - another process detecting the occurrence of event wakes up suspended process
- ⇒ “suspend – wake-up” synchronization

## Synchronization with Semaphores

- solution:

event: semaphore; event=0;

<b>process P1:</b>	<b>process P2:</b>
...	...
P(event);	...
...	V(event);
	...

- more than two processes may be synchronized

## Semaphores

Initial value for semaphore:

- =1 for mutual exclusion
- =0 for synchronization

## Semaphores

- possible deadlock scenario:

x, y: semaphore;

x=1; y=1;

<b>process 1:</b>	<b>process 2:</b>
...	...
P(x);	P(y);
...	...
P(y);	P(x);
...	...
V(x);	V(y);
V(y);	V(x);
...	...

