

# Computer Operating Systems, Practice Session 5

## Semaphore Operations in Unix

Mustafa Ersen (ersenm@itu.edu.tr)

Istanbul Technical University  
34469 Maslak, İstanbul

18 March 2015

Today

# Computer Operating Systems, PS 5

Semaphore Operations

Signal Mechanism in Linux

Examples

# Semaphore Creation

- ▶ Header files in Unix to be used in semaphore operations:
  - ▶ `sys/ipc.h`
  - ▶ `sys/sem.h`
  - ▶ `sys/types.h`
- ▶ Semaphore Creation:

```
int semget(key_t key, int nsems, int semflg);
```

If successful, a nonnegative integer is returned as the semaphore set identifier, otherwise -1 is returned, with `errno` of the error.

`semflg`: `IPC_CREAT|0700` (Last 9 bits: permission flags)

A semaphore set including `nsems` semaphores is created and associated with `key` if one of the following holds:

- ▶ `key` is equal to `IPC_PRIVATE`
- ▶ `IPC_CREAT` & No semaphore set exists associated with `key` value

# Semaphore Operations

- ▶ `int semop(int semid, struct sembuf *sops, unsigned nsops);`
  - ▶ `semop` operates on semaphores selected from semaphore set associated with `semid`
  - ▶ Each of the `nsops` elements, pointed by `sops`, determines operation on a specific semaphore (each element is of type: `sembuf` )
- ▶ `struct sembuf{`
  - `unsigned short sem_num; // semaphore number starts with 0`
  - `short sem_op; // semaphore operation`
  - `short sem_flg; // operation flags``};`
- ▶ The operations contained in `sops` are performed in array order **atomically** (i.e., the operations are performed either as a complete unit, or not at all)
- ▶ `sem_flg`
  - ▶ `SEM_UNDO`: Allows individual operations in the array to be automatically undone when the process exits.
  - ▶ `IPC_NOWAIT`: (Do not allow to wait) If you can not decrease, give error message and return
- ▶ `sem_op`
  - ▶ `== 0`: wait for it to be 0 (Must have read permission)
  - ▶ `!= 0`: value is added to the semaphore value (The process must have alter permission on the semaphore set)

# Semaphore Control

- ▶ Control of the Value

```
int semctl(int semid, int semnum, int cmd, arg);
```

- ▶ cmd

- ▶ IPC\_RMID : Remove the semaphore set, awakening all processes blocked
- GETVAL : Return the value of `semval` for the corresponding semaphore
- SETVAL : Set the value of `semval` of the corresponding semaphore to `arg.val`
- SETALL : Set `semval` values for all semaphores of the set using `arg.array`
- GETALL : Return all of the `semval` values for all semaphores of the set into `arg.array`

# Basic Semaphore Operations

```
1 // increment operation
2 void sem_signal(int semid, int val){
3     struct sembuf semaphore;
4     semaphore.sem_num=0;
5     semaphore.sem_op=val;
6     semaphore.sem_flg=1; // relative: add sem_op to value
7     semop(semid, &semaphore, 1);
8 }
9
10 // decrement operation
11 void sem_wait(int semid, int val){
12     struct sembuf semaphore;
13     semaphore.sem_num=0;
14     semaphore.sem_op=(-1*val);
15     semaphore.sem_flg=1; // relative: add sem_op to value
16     semop(semid, &semaphore, 1);
17 }
```

# Handling Signals

► Necessary header files for handling signals:

- signal.h
- sys/types.h

```
1 // signal-handling function
2 void mysignal(int signum){
3     printf("Received signal with num=%d\n", signum);
4 }
5
6 void mysigset(int num){
7     struct sigaction mysigaction;
8     mysigaction.sa_handler=(void *)mysignal;
9     // using the signal-catching function identified by sa_handler
10    mysigaction.sa_flags=0;
11    // sigaction() system call is used to change the action taken by
12    // a process on receipt of a specific signal(specified with num)
13    sigaction(num,&mysigaction ,NULL);
14 }
```

# Handling Signals

- ▶ Sending a signal (specified with num=sig) from a process to another process (with given pid):

```
int kill(pid_t pid, int sig);
```

- ▶ Waiting for a signal:

```
int pause(void);
```



# Example 1

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/wait.h>
5 #include <sys/ipc.h>
6 #include <sys/sem.h>
7 #include <sys/types.h>
8 #include <signal.h>    // sigaction
9
10 #define SEMKEY 8
11 int sem_id;
12
13 // increment operation
14 void sem_signal(int semid, int val){
15     struct sembuf semaphore;
16     semaphore.sem_num=0;
17     semaphore.sem_op=val;
18     semaphore.sem_flg=1;    // relative: add sem_op to value
19     semop(semid, &semaphore, 1);
20 }
```

# Example 1

```
1 // decrement operation
2 void sem_wait(int semid, int val){
3     struct sembuf semaphore;
4     semaphore.sem_num=0;
5     semaphore.sem_op=(-1*val);
6     semaphore.sem_flg=1; // relative: add sem_op to value
7     semop(semid, &semaphore, 1);
8 }
9
10 // signal-handling function
11 void mysignal(int signum){
12     printf("Received signal with num=%d\n", signum);
13 }
14
15 void mysigset(int num){
16     struct sigaction mysigaction;
17     mysigaction.sa_handler=(void *)mysignal;
18     // using the signal-catching function identified by sa_handler
19     mysigaction.sa_flags=0;
20     // sigaction() system call is used to change the action taken by
21     // a process on receipt of a specific signal(specified with num)
22     sigaction(num,&mysigaction, NULL);
23 }
```

# Example 1

```
1  int main(void){
2      // signal handler with num=12
3      mysigset(12);
4      int f=1, i, children[10];
5      // creating 10 child processes
6      for(i=0; i<10; i++){
7          if (f>0)
8              f=fork();
9          if (f==-1){
10             printf("fork error....\n");
11             exit(1);
12         }
13         if (f==0)
14             break;
15         else
16             children[i]=f; // get pid of each child process
17     }
```

# Example 1

```
1  // parent process
2  if (f > 0) {
3      // creating a semaphore with key=SEMKEY
4      sem_id = semget(SEMKEY, 1, 0700|IPC_CREAT);
5      // setting value of the 0th semaphore of the set identified
6      // with sem_id to 0
7      semctl(sem_id, 0, SETVAL, 0);
8      // waiting for a second
9      sleep(1);
10     // sending the signal 12 to all child processes
11     for (i=0; i<10; i++)
12         kill(children[i], 12);
13     // decrease semaphore value by 10 (i.e., wait for all children
14     // to increase semaphore value)
15     sem_wait(sem_id, 10);
16     printf("ALL CHILDREN HAS Finished ...\n");
17     // remove the semaphore set identified with sem_id
18     semctl(sem_id, 0, IPC_RMID, 0);
19     exit(0);
20 }
```

# Example 1

```
1  // child process
2  else{
3      // wait for a signal
4      pause();
5      // returning the sem_id associated with SEMKEY
6      sem_id = semget(SEMKEY, 1, 0);
7      printf("I am the CHILD Process created in %dth order. My PROCESS ID: %d\n",
8             i, getpid());
9      // getting value of the 0th semaphore of the set identified
10     // with sem_id
11     printf("SEMAPHORE VALUE: %d\n", semctl(sem_id, 0, GETVAL, 0));
12     // increase semaphore value by 1
13     sem_signal(sem_id, 1);
14 }
15 return 0;
16 }
```

## Output of Example 1

```
1 Received signal with num=12
2 I am the CHILD Process created in 0th order. My PROCESS ID: 1629
3 SEMAPHORE VALUE: 0
4 Received signal with num=12
5 I am the CHILD Process created in 1th order. My PROCESS ID: 1630
6 SEMAPHORE VALUE: 1
7 Received signal with num=12
8 I am the CHILD Process created in 2th order. My PROCESS ID: 1631
9 Received signal with num=12
10 I am the CHILD Process created in 3th order. My PROCESS ID: 1632
11 SEMAPHORE VALUE: 2
12 SEMAPHORE VALUE: 3
13 Received signal with num=12
14 I am the CHILD Process created in 4th order. My PROCESS ID: 1633
15 SEMAPHORE VALUE: 3
16 Received signal with num=12
17 I am the CHILD Process created in 6th order. My PROCESS ID: 1635
18 SEMAPHORE VALUE: 5
```

## Output of Example 1 (Continues)

```
1 Received signal with num=12
2 Received signal with num=12
3 I am the CHILD Process created in 9th order. My PROCESS ID: 1638
4 SEMAPHORE VALUE: 6
5 Received signal with num=12
6 Received signal with num=12
7 I am the CHILD Process created in 5th order. My PROCESS ID: 1634
8 SEMAPHORE VALUE: 7
9 I am the CHILD Process created in 7th order. My PROCESS ID: 1636
10 SEMAPHORE VALUE: 8
11 I am the CHILD Process created in 8th order. My PROCESS ID: 1637
12 SEMAPHORE VALUE: 9
13 ALL CHILDREN HAS Finished ...
```

## Example 2 - Deadlock

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/wait.h>
5 #include <sys/ipc.h>
6 #include <sys/sem.h>
7 #include <sys/types.h>
8 #include <signal.h>
9
10 #define SEMKEY_A 1
11 #define SEMKEY_B 2
12 #define SEMKEY_C 3
13
14 // increment operation
15 void sem_signal(int semid, int val){
16     struct sembuf semaphore;
17     semaphore.sem_num=0;
18     semaphore.sem_op=val;
19     semaphore.sem_flg=1; // relative: add sem_op to value
20     semop(semid, &semaphore, 1);
21 }
```



## Example 2 - Deadlock

```
1 // decrement operation
2 void sem_wait(int semid, int val){
3     struct sembuf semaphore;
4     semaphore.sem_num=0;
5     semaphore.sem_op=(-1*val);
6     semaphore.sem_flg=1; // relative: add sem_op to value
7     semop(semid, &semaphore, 1);
8 }
9
10 // signal-handling function
11 void mysignal(int signum){
12     printf("Received signal with num=%d\n", signum);
13 }
14
15 void mysigset(int num){
16     struct sigaction mysigaction;
17     mysigaction.sa_handler=(void *)mysignal;
18     // using the signal-catching function identified by sa_handler
19     mysigaction.sa_flags=0;
20     // sigaction() system call is used to change the action taken by
21     // a process on receipt of a specific signal(specified with num)
22     sigaction(num,&mysigaction, NULL);
23 }
```

## Example 2 - Deadlock

```
1  int main(void){
2      // signal handler with num=12
3      mysigset(12);
4      int semA,semB,semC, c[2], f=1,i ,myOrder;
5      // creating 2 child processes
6      for(i=0; i<2; i++){
7          if (f>0)
8              f=fork();
9          if (f==-1){
10             printf("fork error ....\n");
11             exit(1);
12         }
13         if (f==0)
14             break;
15         else
16             c[i]=f; // get pid of each child process
17     }
```

## Example 2 - Deadlock

```
1 // parent process
2 if (f!=0){
3     printf("PARENT is starting to CREATE RESOURCES....\n");
4     // creating 3 semaphores and setting two of them as 1 and the other as 0
5     semA=semget(SEMKEY_A,1,0700|IPC_CREAT);
6     semctl(semA, 0, SETVAL, 1);
7     semB=semget(SEMKEY_B,1,0700|IPC_CREAT);
8     semctl(semB, 0, SETVAL, 1);
9     semC=semget(SEMKEY_C,1,0700|IPC_CREAT);
10    semctl(semC, 0, SETVAL, 0);
11    sleep(2);
12    printf("PARENT is starting CHILD Processes ..... \n");
13    // sending the signal 12 to all child processes
14    for (i=0; i<2; i++)
15        kill(c[i],12);
16    // decrease semaphore value by 2 (i.e., wait for all children)
17    sem_wait(semC,2);
18    printf("PARENT: Child processes has done, resources are removed back.\n");
19    // remove the created semaphore sets
20    semctl(semC,0,IPC_RMID,0);
21    semctl(semA,0,IPC_RMID,0);
22    semctl(semB,0,IPC_RMID,0);
23    exit(0);
24 }
```

## Example 2 - Deadlock

```
1 // child process
2 else{
3     myOrder=i;
4     printf("CHILD %d: waiting permission from PARENT ....\n", myOrder);
5     // wait for a signal
6     pause();
7     // returning the sem_ids associated with SEMKEY_A, SEMKEY_B and SEMKEY_C
8     semA=semget(SEMKEY_A,1,0);
9     semB=semget(SEMKEY_B,1,0);
10    semC=semget(SEMKEY_C,1,0);
11    printf("CHILD %d has permission from PARENT, is starting ...\n", myOrder);
12    if (myOrder==0){
13        printf("CHILD %d: DECREASING sem A.\n", myOrder);
14        sem_wait(semA, 1);
15        sleep(1);
16        printf("CHILD %d: sem A is completed, DECREASING sem B.\n", myOrder);
17        sem_wait(semB, 1);
18        printf("CHILD %d: I am in the CRITICAL REGION.\n", myOrder);
19        sleep(5); /* Critical Region Operations */
20        // increase all the semaphore values by 1
21        sem_signal(semB, 1);
22        sem_signal(semA, 1);
23        sem_signal(semC, 1);
24    }
```

## Example 2 - Deadlock

```
1  else if (myOrder==1){
2      printf("CHILD %d: DECREASING sem B.\n", myOrder);
3      sem_wait(semB, 1);
4      sleep(1);
5      printf("CHILD %d: sem B is completed, DECREASING sem A.\n", myOrder);
6      sem_wait(semA, 1);
7      printf("CHILD %d: I am in the CRITICAL REGION.\n", myOrder);
8      sleep(5); /* Critical Region Operations */
9      // increase all the semaphore values by 1
10     sem_signal(semA,1);
11     sem_signal(semB,1);
12     sem_signal(semC,1);
13 }
14 }
15 return 0;
16 }
```

## Output of Example 2 - Deadlock

```
1 PARENT is starting to CREATE RESOURCES....  
2 CHILD 1: waiting permission from PARENT ....  
3 CHILD 0: waiting permission from PARENT ....  
4 PARENT is starting CHILD Processes .....  
5 Received signal with num=12  
6 CHILD 1 has permission from PARENT, is starting ...  
7 CHILD 1: DECREASING sem B.  
8 Received signal with num=12  
9 CHILD 0 has permission from PARENT, is starting ...  
10 CHILD 0: DECREASING sem A.  
11 CHILD 1: sem B is completed , DECREASING sem A.  
12 CHILD 0: sem A is completed , DECREASING sem B.
```

## Example 3 - Preventing Deadlock

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <sys/wait.h>
5 #include <sys/ipc.h>
6 #include <sys/sem.h>
7 #include <sys/types.h>
8 #include <signal.h>
9 #include <sys/errno.h>
10
11 #define SEMKEY_AB 5
12 #define SEMKEY_C 6
```

## Example 3 - Preventing Deadlock

```
1 // increment operation
2 void sem_signal(int semid, int val){
3     struct sembuf semaphore;
4     semaphore.sem_num=0;
5     semaphore.sem_op=val;
6     semaphore.sem_flg=1; // relative: add sem_op to value
7     semop(semid, &semaphore, 1);
8 }
9
10 // increment operation using two semaphores
11 void sem_multi_signal(int semid, int val, int nsems){
12     struct sembuf semaphore[2];
13     int i;
14     for (i=0; i<nsems; i++){
15         semaphore[i].sem_num=i;
16         semaphore[i].sem_op=val;
17         semaphore[i].sem_flg=1;
18     }
19     // TWO Operations are performed on SAME SEMAPHORE SET
20     semop(semid, semaphore, 2);
21     for (i=0; i<nsems; i++){
22         printf("SIGNAL : SEM %d IS NOW: .... %d\n", i, semctl(semid, i, GETVAL, 0));
23     }
24 }
```



## Example 3 - Preventing Deadlock

```
1 // decrement operation
2 void sem_wait(int semid, int val){
3     struct sembuf semaphore;
4     semaphore.sem_num=0;
5     semaphore.sem_op=(-1*val);
6     semaphore.sem_flg=1; // relative: add sem_op to value
7     semop(semid, &semaphore, 1);
8 }
9
10 // decrement operation using two semaphores
11 void sem_multi_wait(int semid, int val, int nsems){
12     struct sembuf semaphore[2];
13     int i;
14     for (i=0; i<nsems; i++){
15         semaphore[i].sem_num=i;
16         semaphore[i].sem_op=(-1*val);
17         semaphore[i].sem_flg=1;
18     }
19     //TWO Operations are performed on SAME SEMAPHORE SET:
20     semop(semid, semaphore, 2);
21     for (i=0; i<nsems; i++){
22         printf("WAIT : SEM %d is NOW .... %d\n", i, semctl(semid, i, GETVAL, 0));
23     }
24 }
```

## Example 3 - Preventing Deadlock

```
1 // signal-handling function
2 void mysignal(int signum){
3     printf("Received signal with num=%d\n", signum);
4 }
5
6 void mysigset(int num){
7     struct sigaction mysigaction;
8     mysigaction.sa_handler=(void *)mysignal;
9     // using the signal-catching function identified by sa_handler
10    mysigaction.sa_flags=0;
11    // sigaction() system call is used to change the action taken by a
12    // process on receipt of a specific signal (specified with num)
13    sigaction(num,&mysigaction ,NULL);
14 }
```

## Example 3 - Preventing Deadlock

```
1  int main(void){
2      // signal handler with num=12
3      mysigset(12);
4      int semAB, semC, c[2], f=1, i, myOrder;
5      // creating 2 child processes
6      for(i=0; i<2; i++){
7          if (f>0)
8              f=fork();
9          if (f==-1){
10             printf("fork error ....\n");
11             exit(1);
12         }
13         if (f==0)
14             break;
15         else
16             c[i]=f; // get pid of each child process
17     }
```

## Example 3 - Preventing Deadlock

```

1  if (f!=0){ // parent process
2      printf("PARENT is starting to CREATE RESOURCES...\n");
3      // creating a set of 2 semaphores and setting their values as 1
4      semAB=semget(SEMKEY_AB, 2, 0700|IPC_CREAT);
5      if(semAB == -1)
6          printf("SEMGET ERROR on SEM SET, Error Code: %d \n", errno);
7      if (semctl(semAB, 0, SETVAL, 1) == -1)
8          printf("SMCTL ERROR on SEM A, Error Code: %d \n", errno);
9      if (semctl(semAB, 1, SETVAL, 1) == -1)
10         printf("SMCTL ERROR on SEM B, Error Code: %d \n", errno);
11     printf("PARENT: SEM A is NOW .... %d\n", semctl(semAB,0,GETVAL,0));
12     printf("PARENT: SEM B is NOW .... %d\n", semctl(semAB,1,GETVAL,0));
13     //creating another semaphore and setting its value as 0
14     semC=semget(SEMKEY_C,1,0700|IPC_CREAT);
15     semctl(semC, 0, SETVAL, 0);
16     printf("PARENT: SEM C is NOW .... %d\n", semctl(semC,0,GETVAL,0));
17     sleep(2);
18     printf("PARENT is starting CHILD Processes ..... \n");
19     for (i=0; i<2; i++)
20         kill(c[i],12);
21     sleep(5);
22     sem_wait(semC,2); //decrease sem. value by 2 (i.e., wait for all children)
23     printf("PARENT: SEM C is NOW .... %d\n", semctl(semC,0,GETVAL,0));
24     printf("PARENT: Child processes has done, resources are removed back.\n");
25     semctl(semC,0,IPC_RMID,0);
26     semctl(semAB,0,IPC_RMID,0);
27     exit(0);
}

```



## Example 3 - Preventing Deadlock

```
1  else{    // child process
2      myOrder=i;
3      printf("CHILD %d: waiting permission from PARENT ....\n", myOrder);
4      // wait for a signal
5      pause();
6      // returning the sem_ids associated with SEMKEY_AB and SEMKEY_C
7      semAB=semget(SEMKEY_AB,2,0);
8      semC=semget(SEMKEY_C,1,0);
9      printf("CHILD %d has permission from PARENT, is starting ....\n",myOrder);
10     printf("CHILD %d: DECREASING sem AB.\n", myOrder);
11     // decrease two semaphores in the set specified by semAB by 1
12     sem_multi_wait(semAB,1,2);
13     printf("CHILD %d: I am in the CRITICAL REGION.\n", myOrder);
14     sleep(5);
15     // increase two semaphores in the set specified by semAB by 1
16     sem_multi_signal(semAB,1,2);
17     // increase the third semaphore by 1
18     sem_signal(semC,1);
19 }
20 return 0;
21 }
```

## Output of Example 3 - Preventing Deadlock

```
1 PARENT is starting to CREATE RESOURCES....
2 PARENT: SEM A is NOW .... 1
3 PARENT: SEM B is NOW .... 1
4 PARENT: SEM C is NOW .... 0
5 CHILD 1: waiting permission from PARENT ....
6 CHILD 0: waiting permission from PARENT ....
7 PARENT is starting CHILD Processes .....
8 Received signal with num=12
9 CHILD 1 has permission from PARENT, is starting ....
10 CHILD 1: DECREASING sem AB.
11 WAIT : SEM 0 is NOW .... 0
12 WAIT : SEM 1 is NOW .... 0
13 CHILD 1: I am in the CRITICAL REGION.
14 Received signal with num=12
15 CHILD 0 has permission from PARENT, is starting ....
16 CHILD 0: DECREASING sem AB.
17 SIGNAL : SEM 0 IS NOW: .... 0
18 SIGNAL : SEM 1 IS NOW: .... 0
19 WAIT : SEM 0 is NOW .... 0
20 WAIT : SEM 1 is NOW .... 0
21 CHILD 0: I am in the CRITICAL REGION.
22 SIGNAL : SEM 0 IS NOW: .... 1
23 SIGNAL : SEM 1 IS NOW: .... 1
24 PARENT: SEM C is NOW .... 0
25 PARENT: Child processes has done, resources are removed back.
```