


8 Cache Memory (The internal memory system)

8.1 The performance gap between the CPU and the main memory:

Main memories are implemented based on DRAM (dynamic RAM) technology.

The speed of processors are much higher than the speeds of main memories.

Therefore main memory can not meet bandwidth requirements of the CPUs.



Memory requests per second, on average
Single processor

Source: Hennesy, Patterson, *Computer Architecture*, 5/e


Access per second, 1/latency

The multicore processors increase the bandwidth requirements.

Intel i7 with 4 cores and 3.2 GHz clock rate can achieve a total peak bandwidth of about 409.6 GB/s.

The peak bandwidth to DRAM main memory is only 6% of this (25 GB/s).

8.2 The memory hierarchy:



There are memories of different speeds, costs and sizes. Faster memory is more expensive.

To increase the average speed and (at the same time) to decrease the cost of the memory system different types of memories are organized in the form of multilevel memory hierarchy.

A faster memory (SRAM) namely the cache memory is installed between the main memory and the CPU.

In cache systems we focus on main memory and cache memory not on disk.

Main memory and cache memory can be accessed by the CPU directly (internal). But the data in the disk must be first fetched in to the main memory (or cache).

The goal: To provide a memory system with high speed close to the fastest level of memory and cost per byte low as the cheapest level.

This solution takes the advantage of locality of computer programs.

8.3 The principle of locality (Locality of reference):

Programs tend to reuse instructions and data they have used recently.

Observation: Most programs spend 90% of its execution time in only 10% of the code.

We can predict what instructions and data a program will use in the near future based on its access in the recent past.

There are two types of locality:

Temporal Locality: Recently accessed addresses are likely to be accessed in the near future.

Spatial Locality: After an access to a memory address the next access will be likely to a near address.

Reasons for locality:

Structure of the program: Generally, related data is stored in nearby locations.

Loops

Arrays

8.4 Memory Technologies

RAM (Random Access Memory):

Actually the word random access is a misuse of the term, because all of the semiconductor (electronic) memories used in computer systems are random access (not sequential).

Distinguishing characteristic of RAM:

- It is possible both to read data from the memory and to write new data into the memory easily and rapidly.

These operations are accomplished through the use of electrical signals.

It could be better to name this type of memory as "Read Write Memory".

- RAM is *volatile*. If the power is interrupted, then the data are lost. Thus, RAM can be used only as temporary storage.


Memory latency measures:

- **Access time:** Time between read request and when desired word arrives
- **Cycle time:** Minimum time between two unrelated requests to memory

There are two types of RAM; DRAM and SRAM.

DRAM (Dynamic RAM):

- It is used to implement the **main memory**.
- A dynamic RAM (DRAM) is made with cells that store data as charge on capacitors.
- Because capacitors have a natural tendency to discharge, dynamic RAMs require periodic charge **refreshing** to maintain data storage (Every $\sim 8\text{ms}$).
- Called *dynamic*, because the stored charge leaks away, even with power continuously applied.
- During refreshing memory is unavailable (latency). (-)
- Must be re-written after being read. Reading destroys the information (latency). (-)
- Difference between access time and cycle time. Cycle time $>$ access time (-)
- Cheap and dense: one transistor/bit. More bits can be placed in one chip. (+)




Some improvements:


- SDRAM (Synchronous DRAM): Added clock to DRAM interface. Burst mode.
- DDRAM (Double data rate DRAM): Data is transferred on both the rising edge and falling edge.

SRAM (Static RAM):

- It is used to implement the **cache memory**.
- SRAM uses flip-flops (or latches) to store bits (see Digital Circuits course).
- It requires 6 transistors/bit (more expensive, less bits per chip). (-)
- Refreshment is not necessary. (+)
- Access time \approx cycle time. (+)

**DRAM vs. SRAM:**

- Both are volatile. Power must be continuously supplied.
- DRAM is more dense (smaller cells, more cells per unit area) and less expensive than a corresponding SRAM.
- DRAM requires the supporting refresh circuitry.
- Latency of DRAM is high.



Associative Memory, Content Addressable Memory (CAM):

- Random access memory (+ circuitry for search)
- A word is retrieved based on a portion of its contents rather than its address.
- The user supplies a data word (Argument A) (not the address) and the CAM searches its entire memory simultaneously (not sequential) to see if that word is stored anywhere in it.
- The user also supplies a key value (K) to determine the portion of data to search for.
- If the search data (or the required portion) is found in a row of the memory then the corresponding match bit is set and the output is the stored data in this row.
- It is used in high speed searching applications.
We will use it to search a data in the cache memory.
- It consists of SRAM to store data and digital circuits for parallel search.

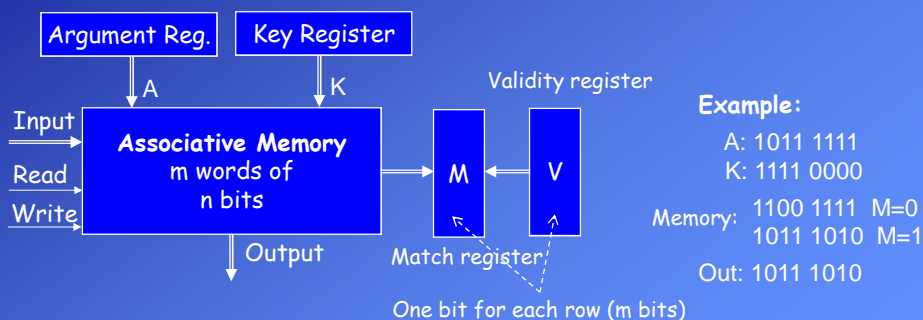
Associative Memory, Content Addressable Memory (CAM) (cont'd):

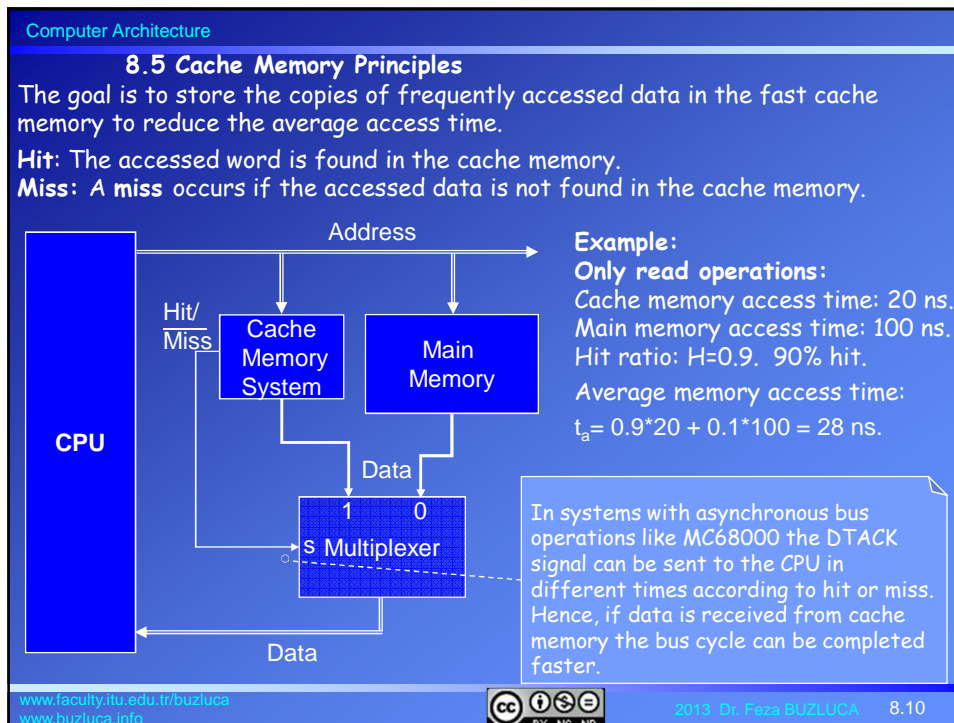

On power-up the memory may include random values.

An additional flag (valid bit V) is necessary to indicate whether or not a memory line has been loaded with valid data.

On power-up, the hardware sets all the valid bits of the all lines to "invalid" ($V_i=0$).

When data is written to line i then corresponding valid bit is set ($V_i=1$).





8.5 Cache Memory Principles (cont'd)

Block transfer:

Upon a cache miss (the requested element is not found in the cache) a *block* of elements that contains the requested element is brought from the main memory to the cache memory.


Reason: spatial locality.

It is expected that the next requested element will be residing in the neighboring locality of the current requested element.

The disadvantage of block transfer can be the long duration.

To decrease the block transfer time between main and cache memories *memory interleaving technique* is used.

Data are stored in different memory modules, that is, consecutive memory addresses are stored in successive memory modules.



8.5 Cache Memory Principles (cont'd)

Replacement Techniques:

When the cache memory is full, a *replacement algorithm* must be applied to select the block in the cache which is to be replaced by the new block from the main memory.

There are different replacement techniques; the most common techniques:

- **FIFO (First In First Out):** The block that has been in the cache the longest is replaced.
- **LRU (Least Recently Used):** The block that has been used the least while residing in the cache is replaced.

The history of block usage is taken into consideration

Aging counters are necessary to keep track of references to blocks in cache.

Cache operations are performed by a **hardware unit** called *Cache Memory Controller* or *Cache Memory Management Unit*.

8.6 Cache Memory Mapping Techniques

Which the contents of the main memory are in the cache memory?
In which location of the cache memory are the referenced data?

1. Full Associative Mapping

a) Without blocks:

Actually all mapping techniques are applied on data blocks.

Firstly, just for simplicity I will show the technique without blocks.

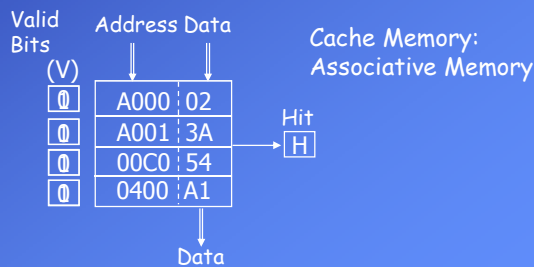
Method: The most referenced addresses and data are kept in a an associative memory.

The address generated by the CPU is searched in the cache (content addressable memory).
If there is a hit data is read from cache.

If a miss occurs data is read from the main memory and data in the cache is replaced.


Without blocks the technique only benefits from the temporal locality but not from the spatial locality.

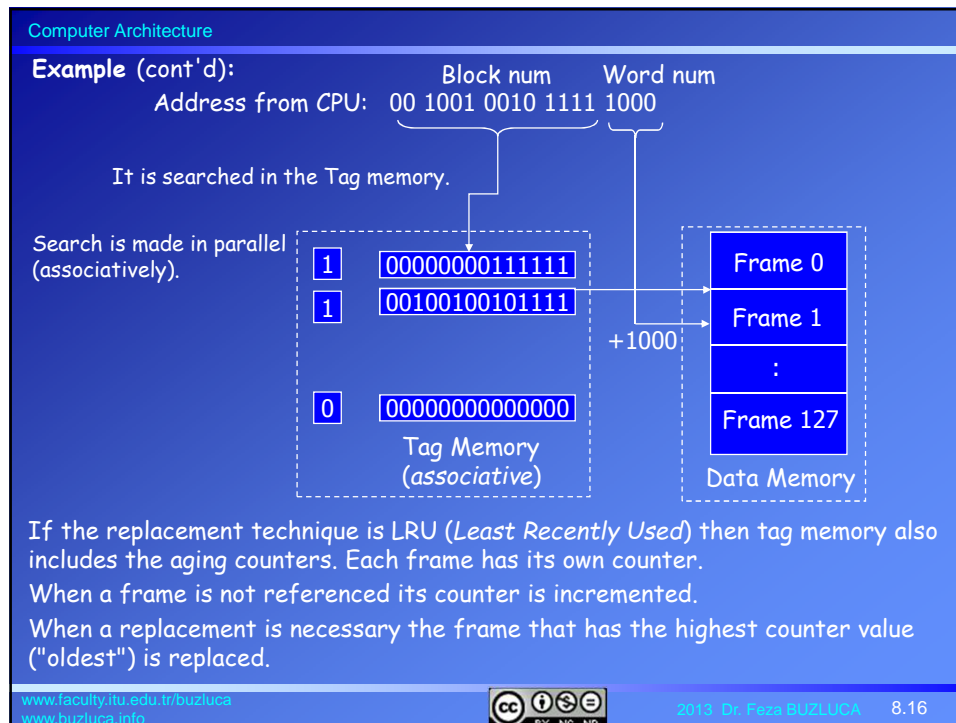
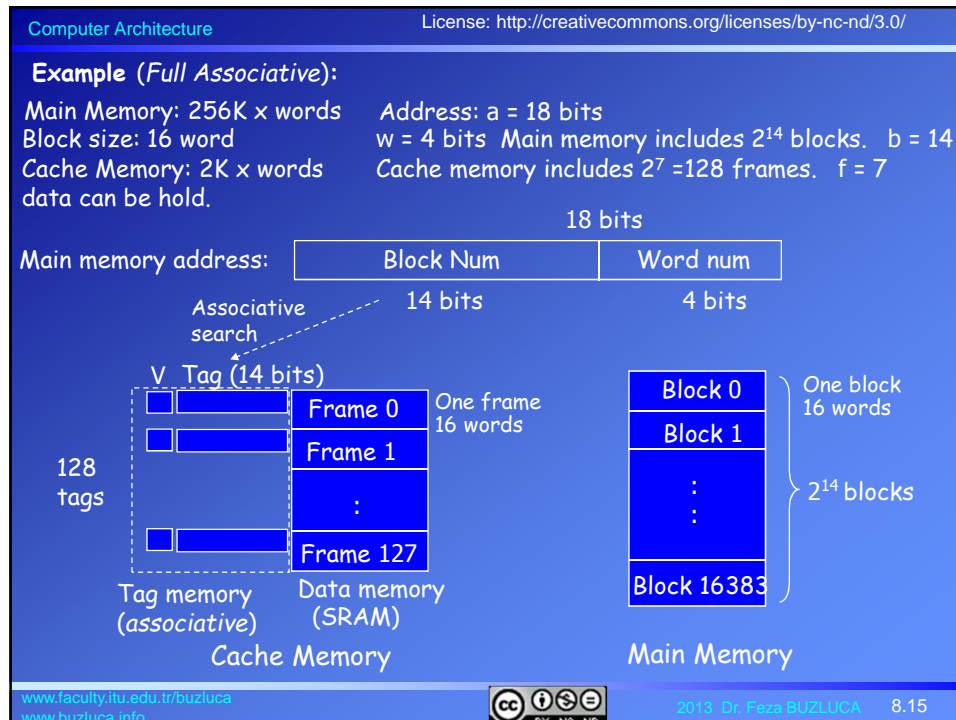
Therefore in practice data blocks are moved between main and cache memories.



b) Full Associative Mapping (with blocks)

The CPU generates a main memory address and the cache controller unit (CCU) divides this address into two fields:





2. Direct Mapping

An incoming main memory block is always placed into a specific fixed cache block location.

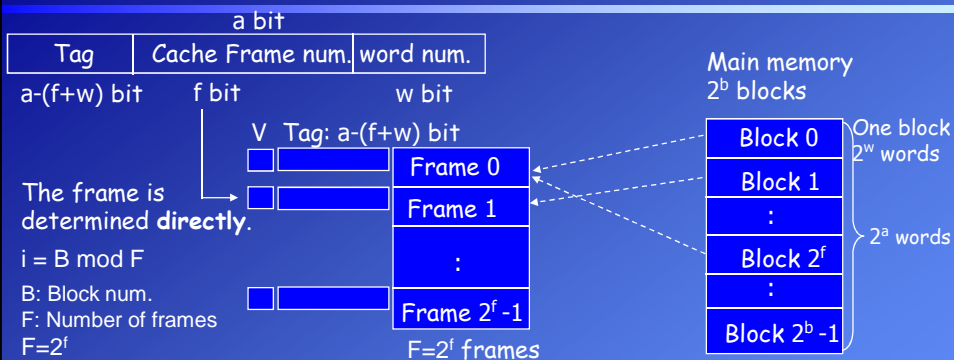

It is not necessary to search the location of a block in the cache, because it is predetermined and fixed.

Therefore associative memory is not necessary.

As the size of the main memory is greater than the cache, several blocks of the main memory try to reside in the same cache frame.

It is necessary to determine which main memory block is currently residing in a frame.

Cache memory control unit divides the address from CPU into three fields:



A main memory block can only reside in the cache frame with the number that is determined by the "Cache Frame num." field of the address.

Even if there are unused frames in the cache, two blocks with the same "Cache Frame Num" field can not be in the cache at the same time.

During replacement a decision algorithm is not necessary.

Because the frame of the incoming block is fixed. This frame will be replaced.

Associative memory is not necessary, because there is no need to search in the cache.

Computer Architecture

Example (Direct Mapping):

Main Memory: 256K x word
 Block size: 16 words
 Cache Memory: 2K x word
 data capacity

Address: $a = 18$ bits
 $w = 4$ bits Main memory contains 2^{14} blocks. $b = 14$
 Cache memory contains $2^7 = 128$ frames. $f = 7$

18 bits		
Tag	Frame num.	word num.
7 bits	7 bits	4 bits

In this system, the data in the following two different addresses try to reside in the same cache frame.

Tag	Frame num	Word num
0000000	0000000	XXXX
0000001	0000000	XXXX

The "Frame number" fields of both addresses are same: 0000000. They will be placed in Frame 0.

At a specific point in time only one of these data can be in the cache memory.

To determine which data is currently in the Frame 0 of the cache memory, the tag value of the address is compared with the tag value stored in cache memory.

Computer Architecture

Example (Direct Mapping):

Main address from CPU:

Tag	Frame	Word
0010001	0000001	1000


Hit/Miss ←

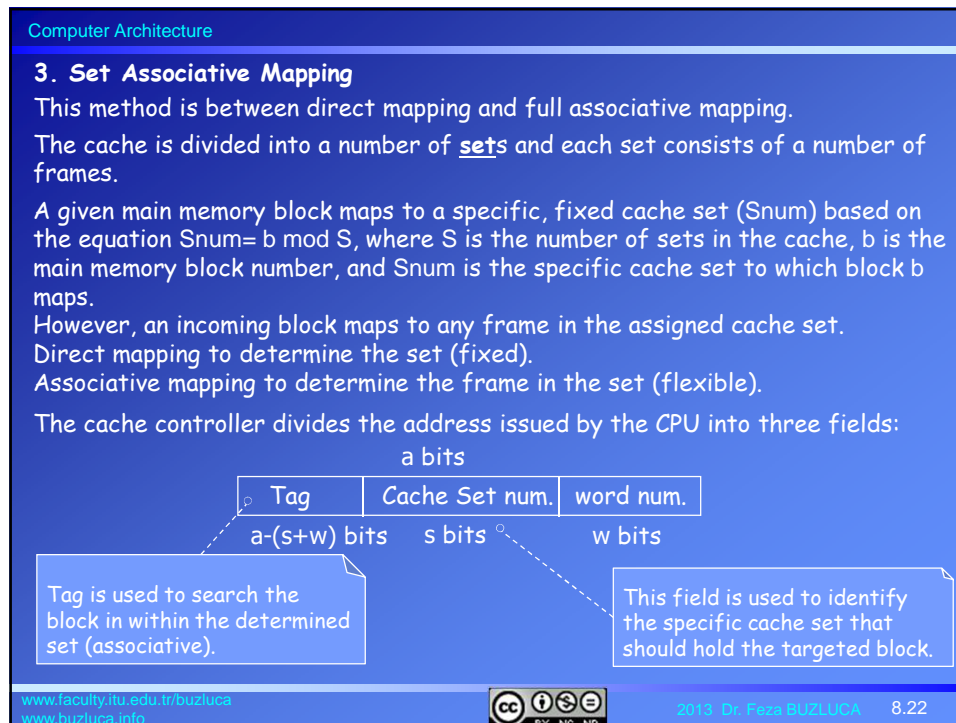
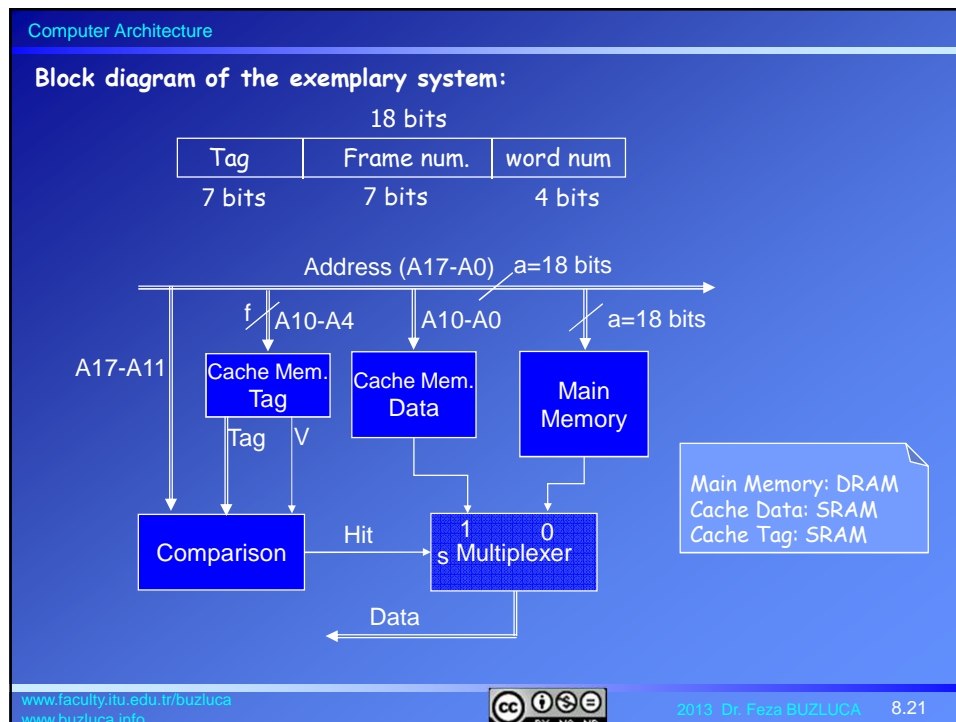
Compare

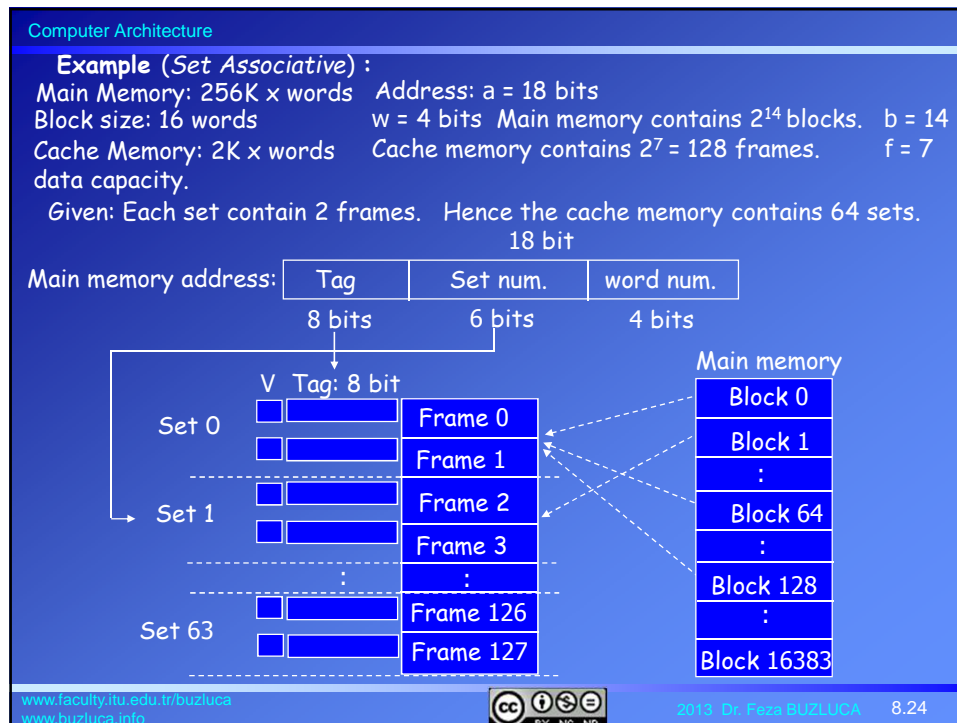
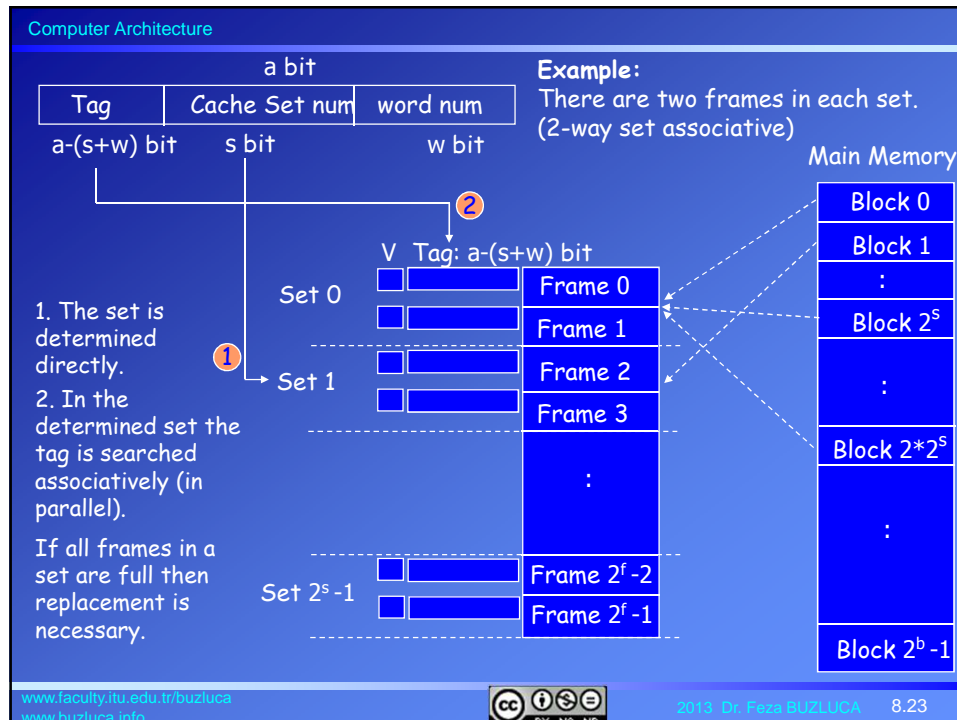
Data memory is addressed directly.

Tag memory is addressed directly.

If the tag in the address is equal to the tag in the cache then the data of the referenced address is in the cache memory.







Summary of mapping techniques:

- **Associative mapping** is the most *flexible* and *efficient* technique because an incoming main memory block can reside in any frame of the cache.
The disadvantage of this technique is the high *hardware cost* due to the associative memory.
- The main disadvantage of the **direct mapping** is the *inefficient* use of the cache, because a number of main memory blocks may compete for a cache frame even if there exist other empty frames.
This disadvantage decreases the hit ratio.
The main advantage of this technique is its simplicity; no search is needed. It is also simple in terms of the replacement mechanism.
- The cache utilization efficiency of the **set associative mapping** technique is expected to be *moderate*; namely between fully associative technique and the direct mapping technique.
However, the technique inherits the simplicity of the direct mapping technique in terms of determining the target set.
By changing the number of frames in a set we can make it close to one of the other techniques.

8.7 Cache Memory - Main Memory Interactions

→ **Read (Hit):** Data is read from the cache memory.

→ **Read (Miss):**

- Read Through (RT)** : While the data (block) is being brought from the main memory to the cache, it is also read by the CPU simultaneously.
Cache memory and the main memory are accessed in parallel.
- No Read Through (NRT)**: Data are first brought from the main memory to the cache memory and then the CPU reads data from the cache.

→ **Write (Hit):**

- Write Through (WT)**: In each write operation data is written to the cache and also to the main memory.
The write-through policy increases the access time but provides coherence between the cache frames and their counterparts in the main memory.
- Write Back (WB)**: All writes are made only to the cache.
A block is written back to the main memory only when a replacement is needed.
There are two types of write-back policy: *Simple write back* and *flagged write back*.

→ **Write (Hit):****b) Write Back (WB) (cont'd):**• **Simple Write Back (SWB):**

The replaced frame is always written back to the main memory.
It is not checked whether the frame was changed or not.

• **Flagged Write Back (FWB):**

Every cache frame is assigned a bit, called the *dirty bit*, to indicate that at least one write operation has been made to the block while residing in the cache.

At replacement time, the dirty bit is checked; if it is set, then the block is written back to the main memory, otherwise, it is simply overwritten by the incoming block.

→ **Write (Miss):**

a) Write Allocate (WA): The main memory block is updated and brought to the cache.

b) No Write Allocate (NWA): The missed main memory block is updated in the main memory and not brought to the cache.

If later an attempt is made to read this block, a miss will occur and data will be brought to the cache.

The write-through (WT) policy can be used together with write-allocate (WA) or no-write-allocate (NWA) methods. WTWA, WTNWA

In write-back (WB) policy to maintain coherence between cache and main memory the write-allocate (WA) method is used. WBWA

Information hold in Tag memory:

In addition to Valid (V) and tag bits, dependent on the used method, following data must be also kept in the tag memory:

If LRU is used *aging counters*, if flagged Write-Back (FWB) method is used "dirty" bit (D).


One line of a tag memory:

V	D	Counter	Tag
---	---	---------	-----

8.8 CPU - Cache Memory - Main Memory Interconnection

CPU - cache memory - main memory interconnection can be implemented in two ways:


a) Parallel connection



- **Read Through (RT):** While a block is brought from the main memory to the cache it is also read by the CPU simultaneously (in parallel).
- **Load Through (LT):** A data is written to the cache and to the main memory simultaneously.

Parallel connection is suitable for the *Write-Through (WT)* method. It can also be used with the *Write-Back (WB)* method.

b) Serial connection



- **No Read Through (NRT):** The block is first brought from the main memory to the cache and then the CPU reads the data from the cache.
- **No Load Through (NLT):** First the data in the cache is updated, then the data is written to the main memory.

Serial connection is suitable for the *Write-Back (WB)* method.

8.9 Access Time:

- t_a : Average Memory Access Time
 w : Write ratio (number of write accesses / total number of all accesses)
 h : Hit ratio
 t_{cache} : Cache memory access time
 t_{main} : Main memory access time
 t_{trans} : Time to transfer a block between main memory and the cache
 w_d : The probability that a block in a cache is updated

Probability	WT, RT/LT (Write-through, Parallel read/write)		WB, WA, NRT/NLT (Write-back, Serial read/write)	
	NWA	WA	SWB	FWB
Read Hit $(1-w)h$	Access Time t_{cache}		Access Time t_{cache}	
Read Miss $(1-w)(1-h)$	t_{trans}	t_{trans}	$2t_{\text{trans}}+t_{\text{cache}}$	$w_d(2t_{\text{trans}}+t_{\text{cache}}) + (1-w_d)(t_{\text{trans}}+t_{\text{cache}})$
Write Hit wh	t_{main}	t_{main}	t_{cache}	t_{cache}
Write Miss $w(1-h)$	t_{main}	$t_{\text{main}}+t_{\text{trans}}$	$2t_{\text{trans}}+t_{\text{cache}}$	$w_d(2t_{\text{trans}}+t_{\text{cache}}) + (1-w_d)(t_{\text{trans}}+t_{\text{cache}})$

Computer Architecture

Access Time Calculation:

• Write Through with Write Allocate, Read/Load Through (WTWA, RT/LT):

Read Hit + Read Miss + Write Hit + Write Miss

$$t_a = (1-w)h t_{\text{cache}} + (1-w)(1-h)t_{\text{trans}} + w \cdot h \cdot t_{\text{main}} + w(1-h)(t_{\text{main}} + t_{\text{trans}})$$

$$t_a = (1-w)h t_{\text{cache}} + (1-h)t_{\text{trans}} + w \cdot t_{\text{main}}$$

• Write Through with No Write Allocate, Read/Load Through (WTNWA, RT/LT)

$$t_a = (1-w)h t_{\text{cache}} + (1-w)(1-h)t_{\text{trans}} + w \cdot h \cdot t_{\text{main}} + w(1-h)t_{\text{main}}$$

$$t_a = (1-w) t_{\text{cache}} + (1-w)(1-h)t_{\text{trans}} + w \cdot t_{\text{main}}$$

• Simple Write Back with Write Allocate, No Read Through (SWBWA, NRT/NLT)

Read Hit + Read Miss + Write Hit + Write Miss

$$t_a = (1-w)h t_{\text{cache}} + (1-w)(1-h)(2t_{\text{trans}} + t_{\text{cache}}) + w \cdot h \cdot t_{\text{cache}} + w(1-h)(2t_{\text{trans}} + t_{\text{cache}})$$

$$t_a = t_{\text{cache}} + (1-h) \cdot 2 \cdot t_{\text{trans}}$$

One t_{trans} is necessary to transfer a frame from the cache to the main memory and the second one to bring the new block from the main memory to the cache.

• Flagged Write Back, Write Allocate, No Read Through (FWBWA, NRT/NLT):

$$t_a = t_{\text{cache}} + (1-h)t_{\text{trans}} + w_d \cdot (1-h)t_{\text{trans}}$$

$$t_a = t_{\text{cache}} + (1-h)(1+w_d)t_{\text{trans}}$$

www.faculty.itu.edu.tr/buzluca
www.buzluca.info



2013 Dr. Feza BUZLUCA 8.31

Computer Architecture

Exemplary processors with cache memories:

- **Intel386™**: Cache memory is outside of the CPU chip. SRAM memory.
- **Intel486™ (1989)**
8-KByte on-chip (L1)
- **Intel® Pentium® (1993)**
L1 on-chip: 8 KB instruction, 8 KB data cache (Harvard architecture)
- **Intel P6 Family: (1995-1999)**
 - Intel Pentium Pro:
L1 on-chip: 8 KB instruction, 8 KB data cache (Harvard architecture)
First L2 cache memory in the CPU chip.
L2 on-chip: 256 KB. Different interconnections between L1, L2 and the CPU.
 - Intel Pentium II:
L1 on-chip: 16 KB instruction, 16 KB data cache (Harvard architecture)
L2 on-chip: 256 KB, 512 KB, 1 MB
- **Intel® Pentium® M (2003)**
L1 on-chip: 32 KB instruction, 32 KB data cache
L2 on-chip: up to 2 MByte
- **Intel® Core™ i7-980X Processor Extreme Edition (2010)**
Multicore: 6 cores. Private caches (L1) and shared caches (L2)
12 MB smartcache: All cores share this cache.

www.faculty.itu.edu.tr/buzluca
www.buzluca.info



2013 Dr. Feza BUZLUCA 8.32