

Analysis of Algorithms 1 (Fall 2013)

Istanbul Technical University Computer Eng. Dept.

Chapter 6: Heapsort



Course slides from
Susan Bridges @MS State
have been used in
preparation of these slides.

Last updated: October 11, 2011

Purpose

- Learn about the heap data structure and its operations
- Learn about the heapsort
- How to use heap for priority queues

Contents

- Heaps
- Maintaining the heap property
- Building a heap
- The heapsort algorithm
- Priority queues

Heapsort


- Running time of heapsort is $O(n \log_2 n)$
- It sorts in place
- It uses a data structure called a *heap*
- The heap data structure is also used to implement a priority queue efficiently

Full and Complete Binary Trees


- **Full binary tree:** binary tree in which each node is either a leaf node or has degree 2 (i.e., has exactly 2 children)
- **Complete binary tree:** full binary tree in which all leaves have the same depth
- **Nearly complete binary tree:** completely filled on all levels except possibly the lowest, which is filled *from the left* up to a point

Examples

Full binary tree:



Complete binary tree:



Representation of Nearly Complete Binary Tree

A nearly complete binary tree may be represented as an array (i.e., no pointers):

Number the nodes, beginning with the root node and moving from level to level, left to right within a level.

The number assigned to a node is its index in the array.


Additional Properties of Nearly Complete Binary Trees

- The root of the tree is $A[1]$.
- If a node has index i , we can easily compute the indices of its:
 - parent $\lfloor i/2 \rfloor$
 - left child $2i$
 - right child $2i + 1$

Numbering

Array:

16	14	10	8	7	9	3	2	4	1
----	----	----	---	---	---	---	---	---	---



Heap

- Implemented as an array object, $A[]$
- Array A that implements the heap has two attributes
 - $\text{length}(A)$
 - $\text{heap-size}(A)$

Heap

A binary tree with n nodes and of height h is **almost complete** iff its nodes correspond to the nodes which are numbered 1 to n in the complete binary tree of height h .

A **heap** is an *almost complete binary tree* that satisfies the **heap property**:

max-heap: For every node i other than the root:

$$A[\text{Parent}(i)] \geq A[i]$$

min-heap: For every node i other than the root:

$$A[\text{Parent}(i)] \leq A[i]$$

Max-Heap


A **max-heap** is an *almost complete binary tree* that satisfies the **heap property**:

For every node i other than the root,

$$A[\text{PARENT}(i)] \geq A[i]$$

What does this mean?

- the value of a node is at most the value of its parent
- the largest element in the heap is stored in the root
- subtrees rooted at a node contain smaller values than the node itself



16	14	10	8	7	9	3	2	4	1
1	2	3	4	5	6	7	8	9	10

Week 4: Heapsort

Height of a node in a heap

The *height* of a node in a heap is the number of edges on the longest simple downward path from the node to a leaf.

The height of a heap is the height of its root.

Since a heap of n elements is based on a complete binary tree, its height is $\Theta(\lg n)$.

Heaps have 5 basic procedures


- HEAPIFY: maintains the heap property
- BUILD-HEAP: builds a heap from an unordered array
- HEAPSORT: sorts an array in place
- EXTRACT-MAX: selects max element
- INSERT: inserts a new element

We will work with MAX heaps

MAX-HEAPIFY(A, i)

- Goal is to put the i^{th} element in the correct place in a portion of the array that “almost” has the heap property.
- The only element with index of i or greater that is out of place is $A[i]$.
- Assume that left and right subtrees of $A[i]$ have the heap property.
- “Sift” $A[i]$ down to the right position.


Array element 2,
the “4”, is out of
place



$\text{MAX-HEAPIFY}(A, 2)$

$\text{heap-size}[A] = 10$


Moving the 4 down



$\text{MAX-HEAPIFY}(A, 4)$

$\text{heap-size}[A] = 10$

The 4 is in the right spot, and the heap property of the tree has been restored.



$\text{MAX-HEAPIFY}(A, 9)$

$\text{heap-size}[A] = 10$

MAX-HEAPIFY

MAX-HEAPIFY(*A*, *i*)

1 *l* \leftarrow LEFT(*i*)

2 *r* \leftarrow RIGHT(*i*) ; *largest* \leftarrow *i*

3 if *l* \leq heap-size[*A*] and *A*[*l*] > *A*[*i*]

4 then *largest* \leftarrow *l*

5 else *largest* \leftarrow *i*

6 if *r* \leq heap-size[*A*] and *A*[*r*] > *A*[*largest*]

7 then *largest* \leftarrow *r*

8 if *largest* \neq *i*

9 then exchange *A*[*i*] \leftrightarrow *A*[*largest*]

10 MAX-HEAPIFY(*A*, *largest*)

Running time of MAX-HEAPIFY

- Run time of MAX-HEAPIFY(A, i)
 - Look at lines 1 – 9
 - Is there a loop? No.
 - Does the number of steps depend upon n ? No.
 - So the running time so far is $\Theta(1)$
 - How about line 10? We do not know yet.

Running time of MAX-HEAPIFY

The recursive call to MAX-HEAPIFY in line 10 implies a recurrence relation.


When we call MAX_HEAPIFY again, we already know that lines 1-9 cost $\Theta(1)$ steps.

But we may need to call MAX-HEAPIFY on a subtree rooted at one of the children of the current node, so we have to add the cost of doing that.

Running time of MAX-HEAPIFY

How many nodes might be involved?


In the case of a full binary tree, about half of the tree might be involved.



Running time of MAX-HEAPIFY

In a complete binary tree with 15 nodes, 8 of those nodes are leaves at the bottom level.


If we perform MAX-HEAPIFY on node i , 7 of the 15 nodes will be involved – about $\frac{1}{2}$ of the nodes.



Running time of MAX-HEAPIFY

What is the worst case?

When the last row of the tree is half full.



Here 7 out of 11 nodes are involved.

In general, $\leq 2/3^{\text{rds}}$ of the tree might be involved in the worst case.

Running time of MAX-HEAPIFY

Remember that, in a complete binary tree, *more than half* of the nodes in the entire tree are the leaf nodes on the bottom level of the tree.

But the only nodes involved in MAX-HEAPIFY are the descendants of $A[i]$, which must be in $A[i]$'s half of the tree.

So worst case is when the last row of the tree is half full on the left side and $A[i]$ is their ancestor.

Running time of MAX-HEAPIFY

The subtrees of the children of our current node have size at most $2n/3$.

The running time of MAX_HEAPIFY can be described by the recurrence:

$$T(n) \leq T(2n/3) + \Theta(1)$$

This is Case 2 by the master method, so:

$$T(n) = O(\lg n)$$

Running time of MAX-HEAPIFY

We could also describe the running time of MAX-HEAPIFY for a node of height h as $O(h)$. (This is useful only if we know the height of a specific node.)

BUILD-MAX-HEAP


- Use MAX-HEAPIFY in a bottom-up manner to convert an array $A[1..n]$ into a heap.
- Each leaf is initially a one-element heap. Elements $A[\lfloor n/2 \rfloor + 1..n]$ are leaves.
- MAX-HEAPIFY is called on all interior nodes.

BUILD-MAX-HEAP

BUILD-MAX-HEAP (A)

```
1  heap-size[A] ← length[A]
2  for i ← floor(length[A]/2) downto 1
    do
3      MAX-HEAPIFY (A, i)
```

a.




$\text{length}(A) = 10$

$\text{floor}(\text{length}(A)/2) = 5$

process from 5 to 1


4	1	3	2	16	9	10	14	8	7
1	2	3	4	5	6	7	8	9	10

b.




4	1	3	2	16	9	10	14	8	7
1	2	3	4	5	6	7	8	9	10

C.




4	1	3	14	16	9	10	2	8	7
1	2	3	4	5	6	7	8	9	10

d.




4	1	10	14	16	9	3	2	8	7
1	2	3	4	5	6	7	8	9	10

e.



4	16	10	14	7	9	3	2	8	1
1	2	3	4	5	6	7	8	9	10

f.



16	14	10	8	7	9	3	2	4	1
1	2	3	4	5	6	7	8	9	10

Running Time of BUILD-MAX-HEAP

- Simple upper bound:
 - each call to MAX-HEAPIFY costs $O(\lg n)$
 - $O(n)$ such calls
 - running time at most $O(n \lg n)$
- Previous bound is not tight:
 - lots of the elements are leaves
 - most elements are near leaves (small height)

Tighter Bound for BUILD-MAX-HEAP

$$\sum_{h=0}^{\lceil \lg n \rceil} \left\lfloor \frac{n}{2^{h+1}} \right\rfloor O(h) = O \left(n \sum_{h=0}^{\lceil \lg n \rceil} \frac{h}{2^h} \right)$$

By substituting $x = 1/2$ in the formula for differentiating infinite geometric series, we have:

$$\sum_{h=0}^{\infty} \frac{h}{2^h} = \frac{1/2}{(1-1/2)^2} = 2$$

Tighter Bound for BUILD-MAX-HEAP (continued)

Thus the running time is bounded by:

$$O\left(n \sum_{h=0}^{\lfloor \lg n \rfloor} \frac{h}{2^h}\right) = O\left(n \sum_{h=0}^{\infty} \frac{h}{2^h}\right) = O(n)$$

Therefore, we can build a heap from an unordered array in linear time.

Heapsort

- First build a heap.
- Then successively remove the biggest element from the heap and move it to the first position in the sorted array.
- The element currently in that position is then placed at the top of the heap and sifted to the proper position.

HEAPSORT

HEAPSORT (A)

1 BUILD-MAX-HEAP (A)

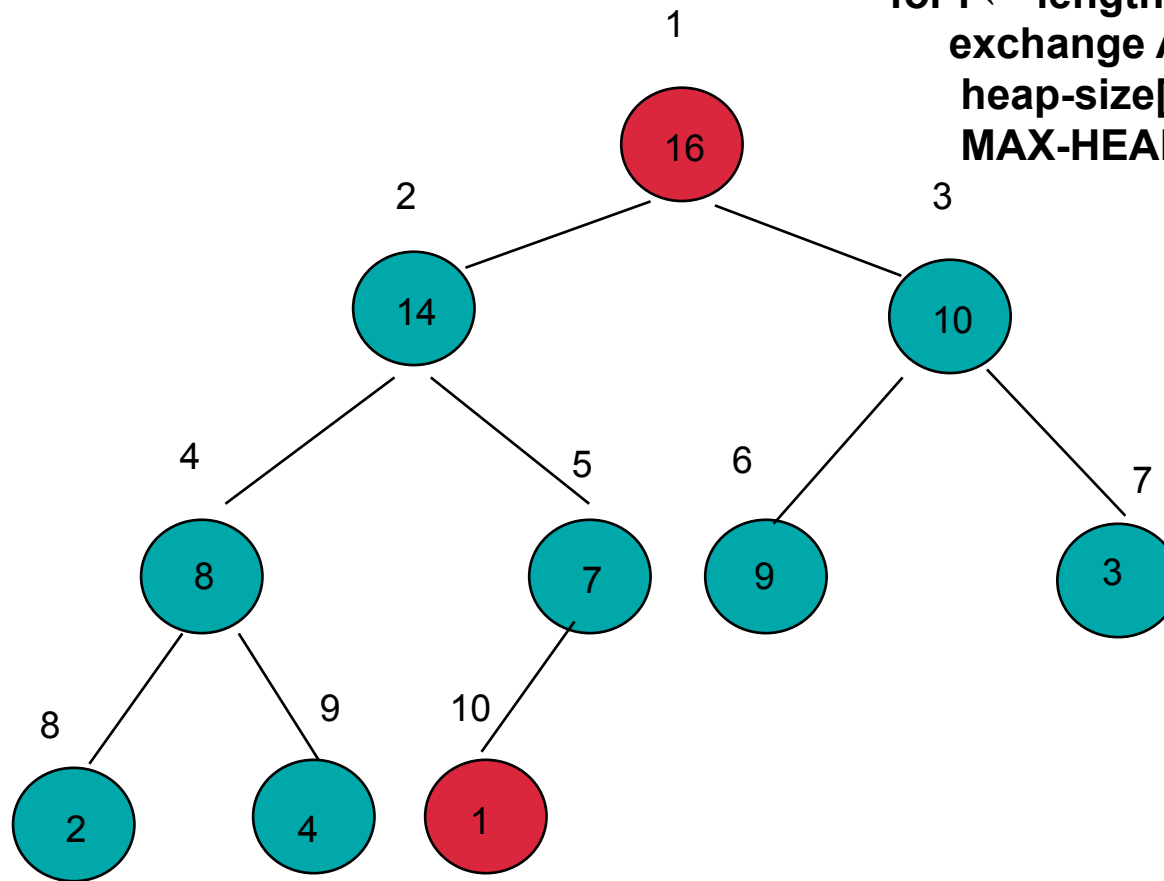
2 for $i \leftarrow \text{length}[A]$ downto 2 do

3 exchange $A[1] \leftrightarrow A[i]$

4 heap-size[A] \leftarrow heap-size[A] - 1

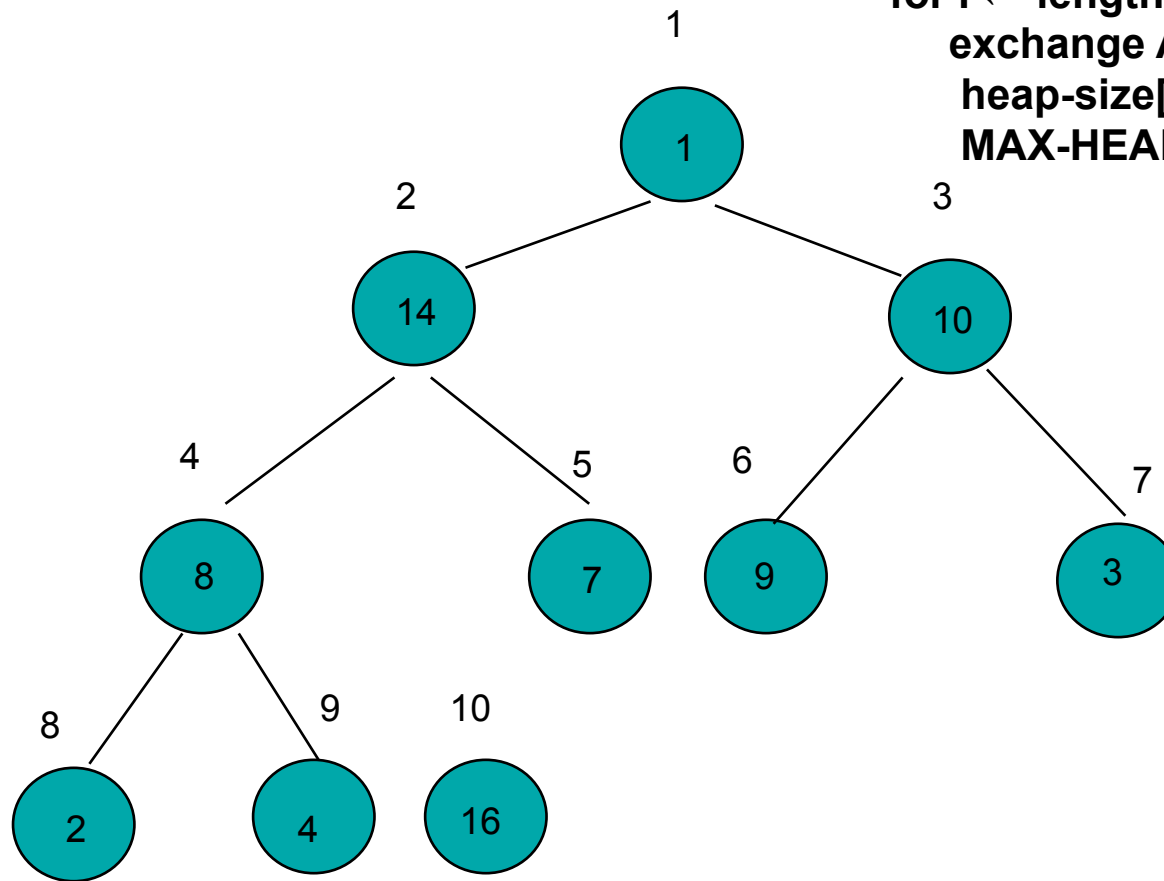
5 MAX-HEAPIFY (A, 1)

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



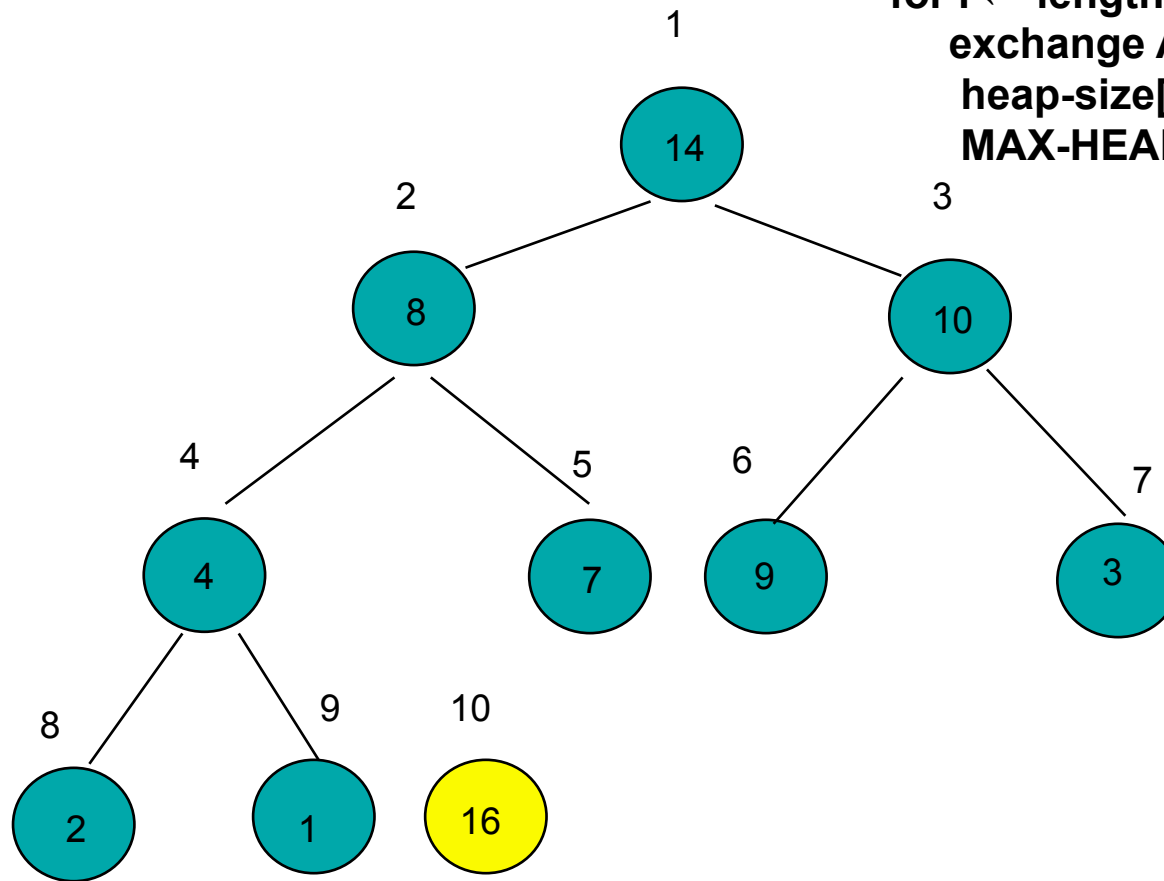
16	14	10	8	7	9	3	2	4	1
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



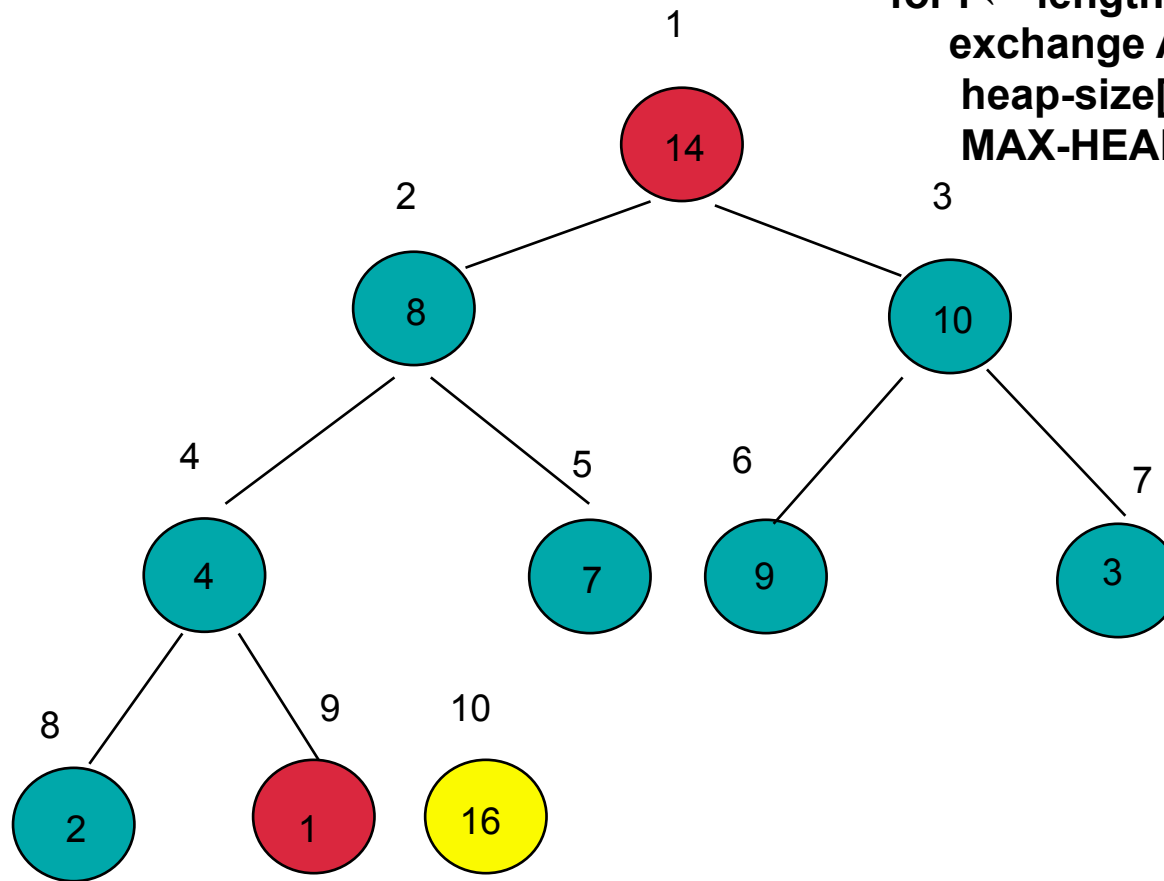
1	14	10	8	7	9	3	2	4	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



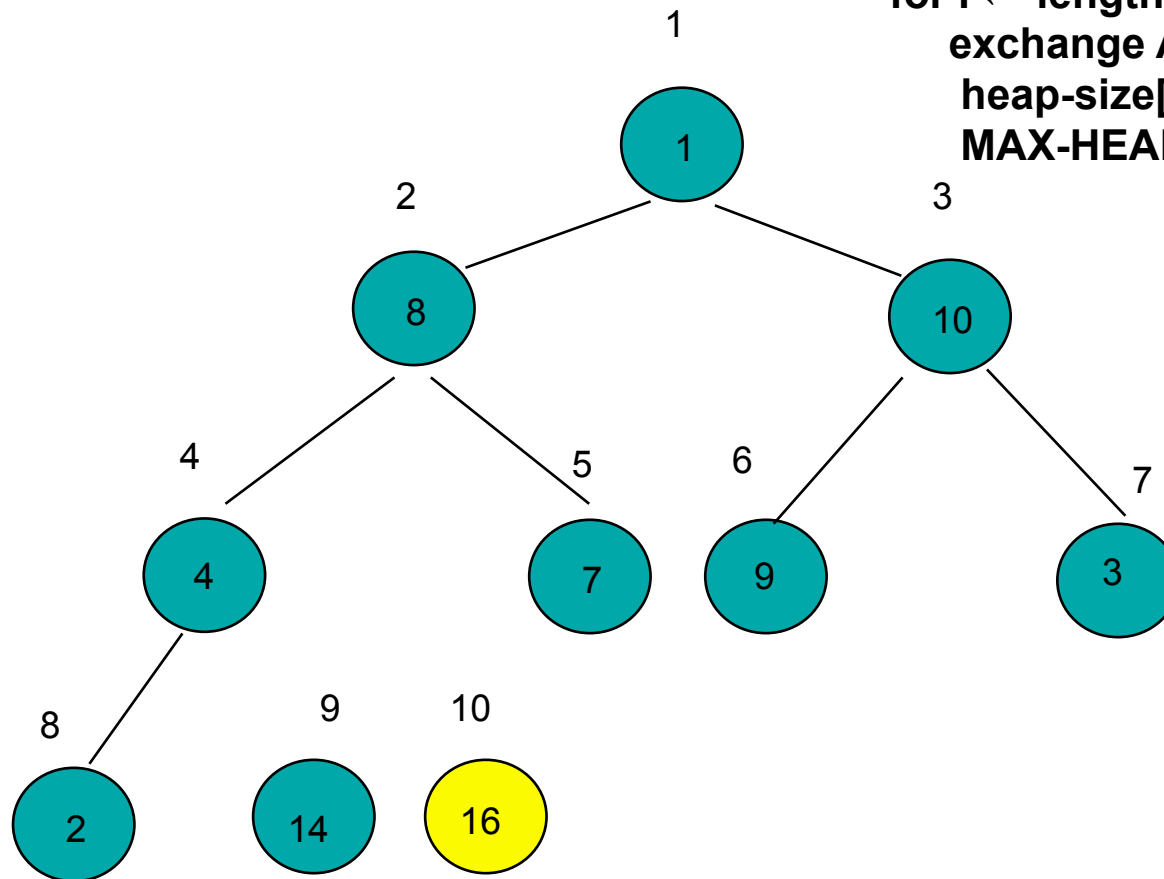
14	8	10	4	7	9	3	2	1	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)



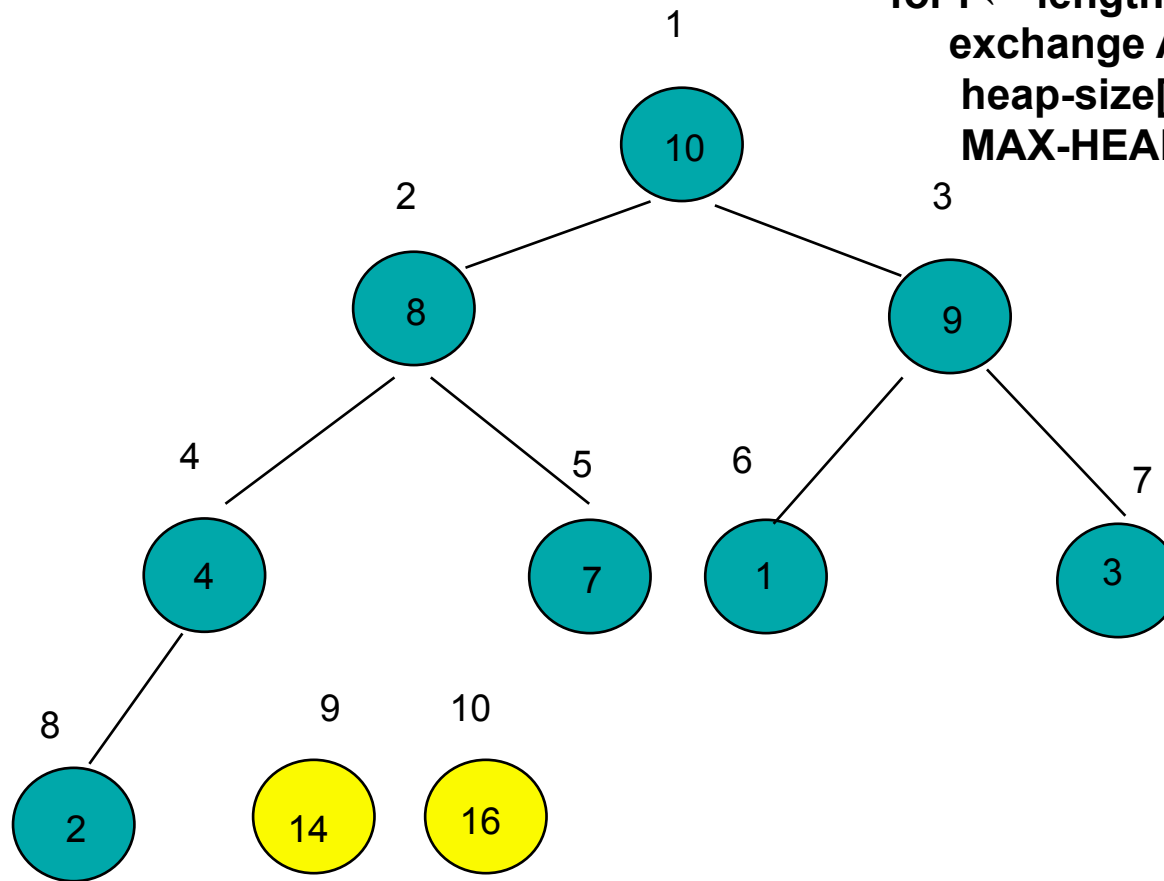
14	8	10	4	7	9	3	2	1	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)



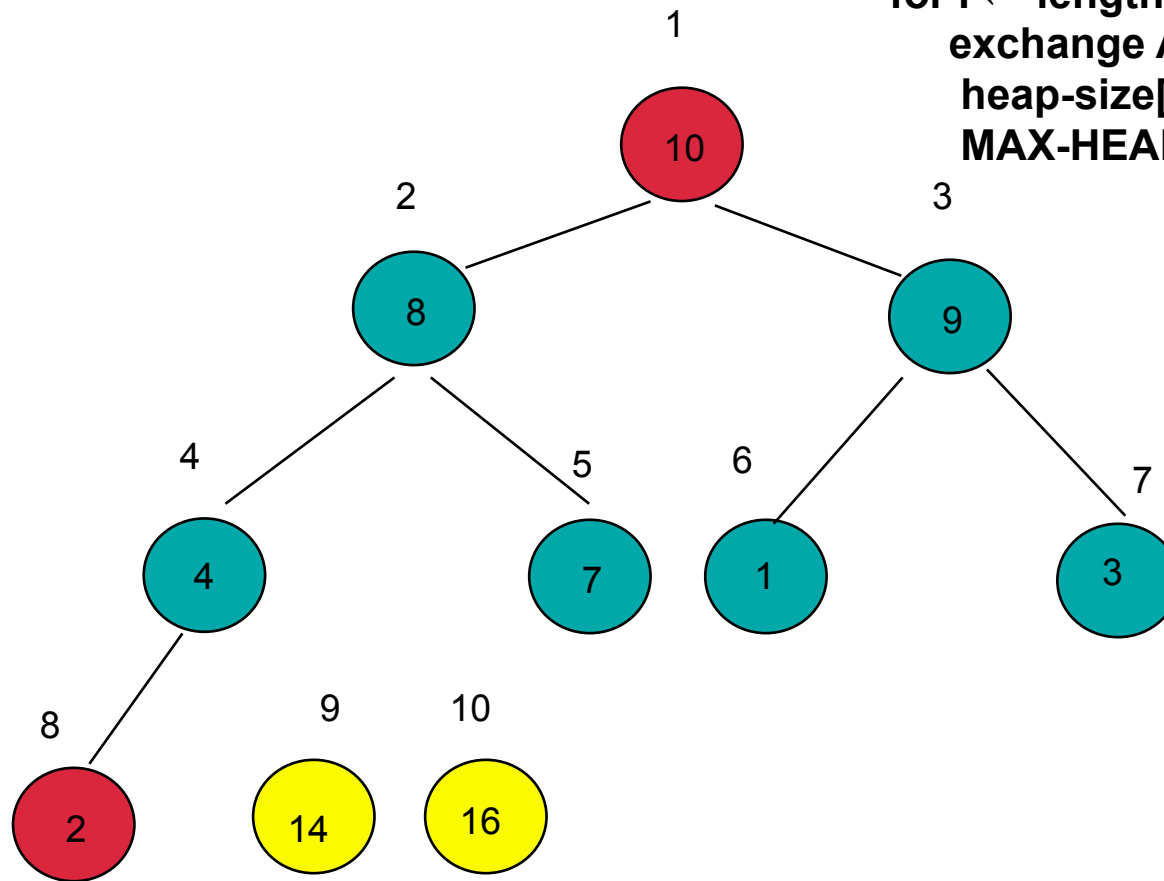
1	8	10	4	7	9	3	2	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)



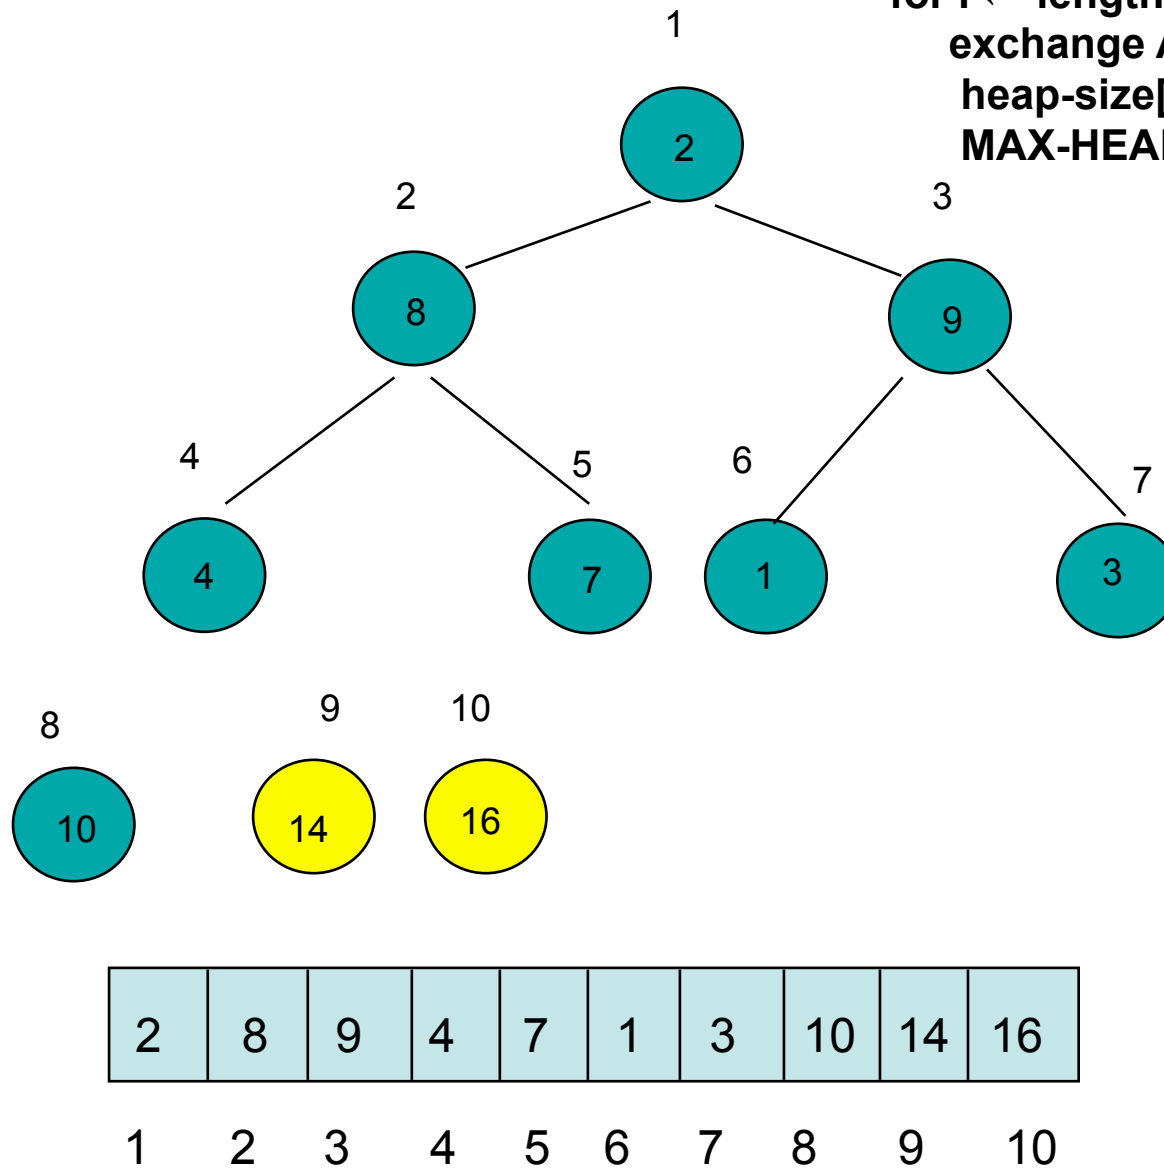
10	8	9	4	7	1	3	2	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)

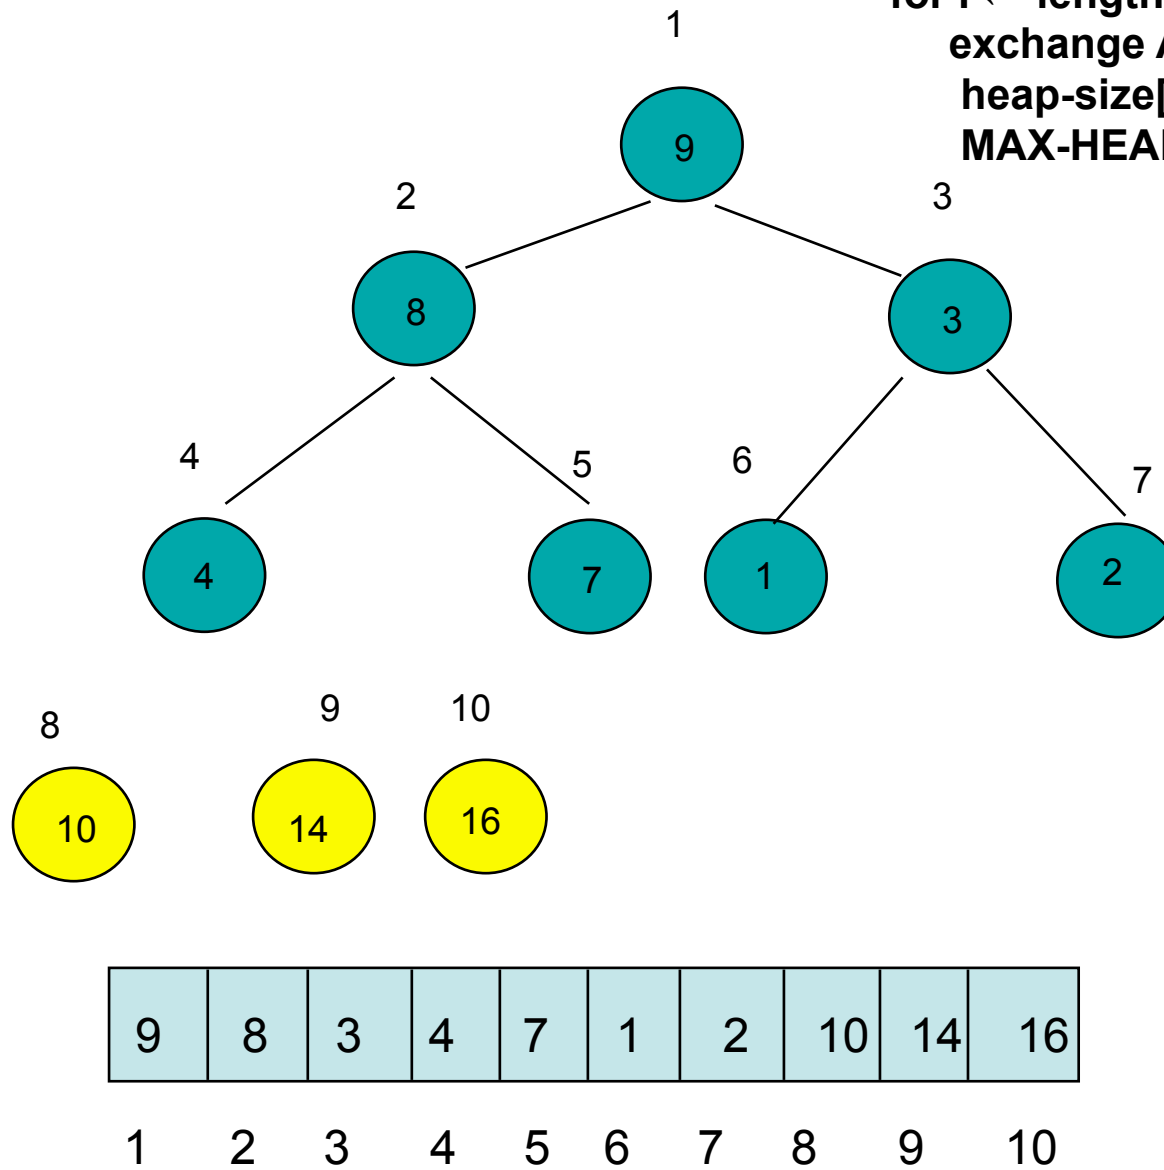


10	8	9	4	7	1	3	2	14	16
1	2	3	4	5	6	7	8	9	10

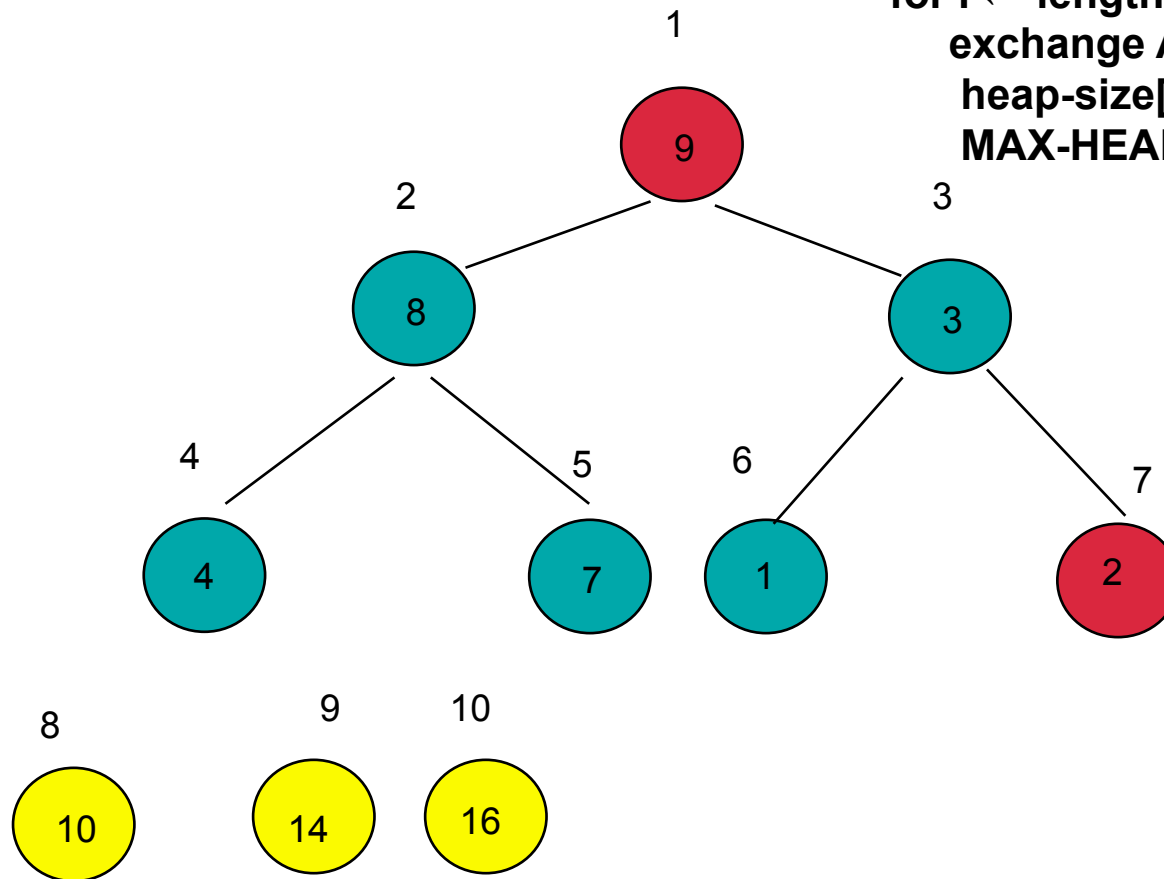
BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)

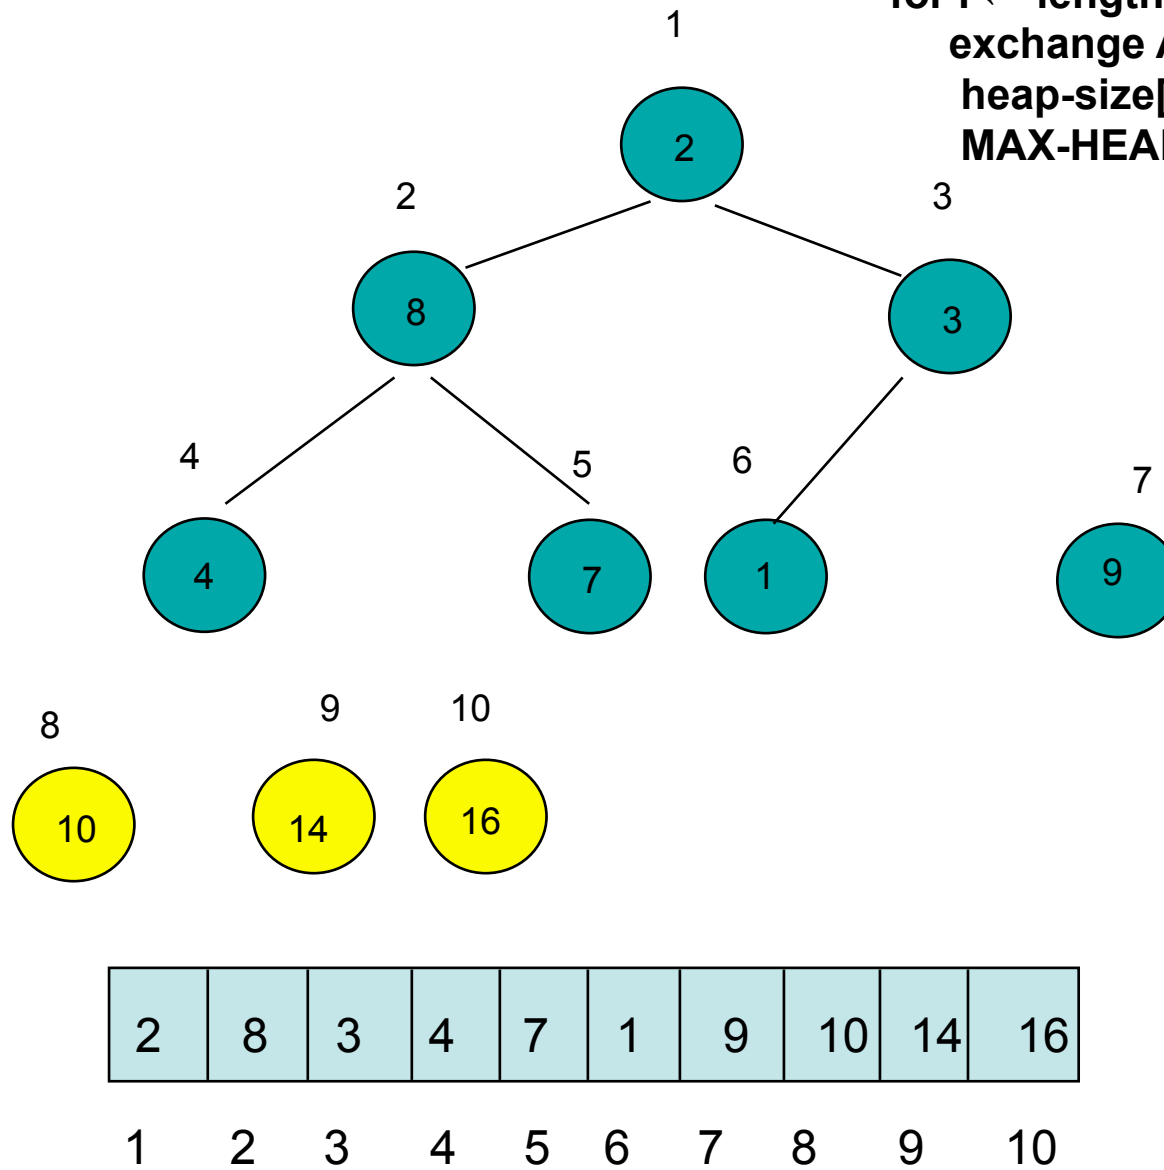


BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)

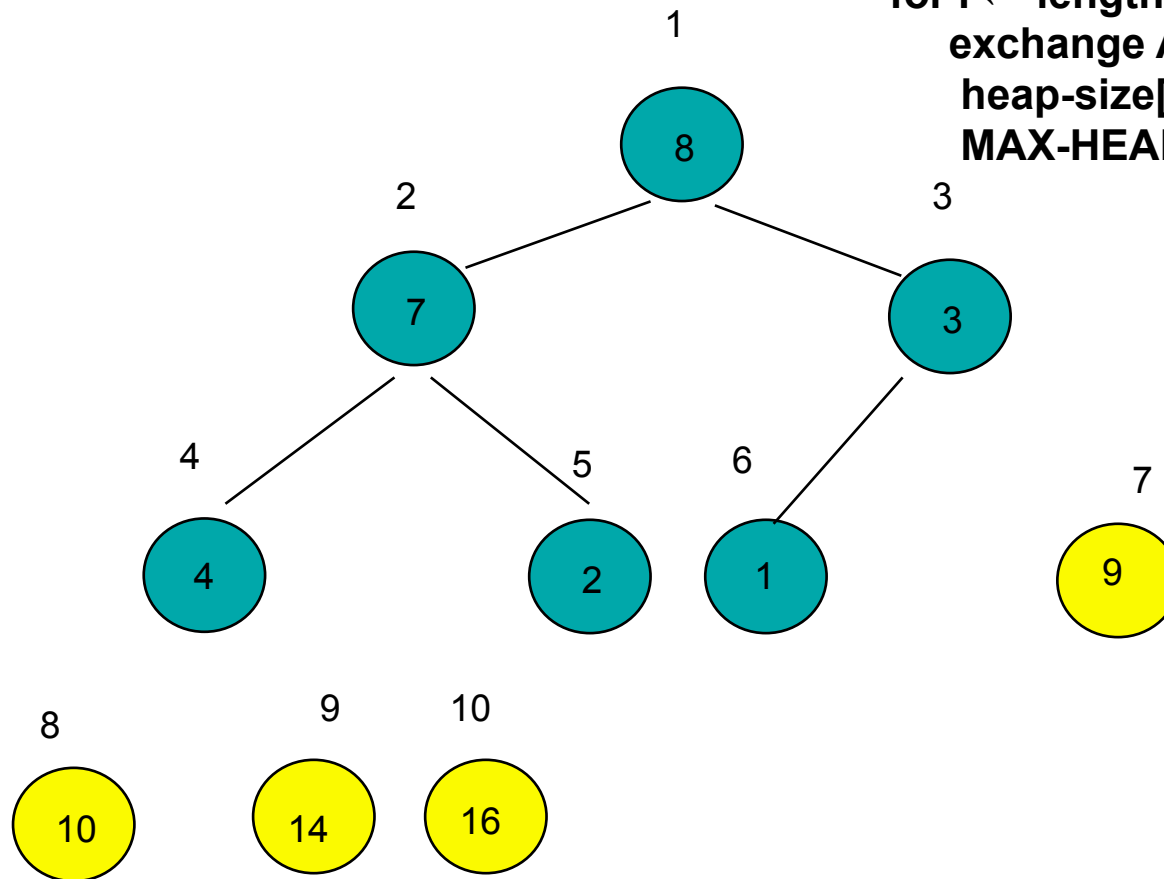


9	8	3	4	7	1	2	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)

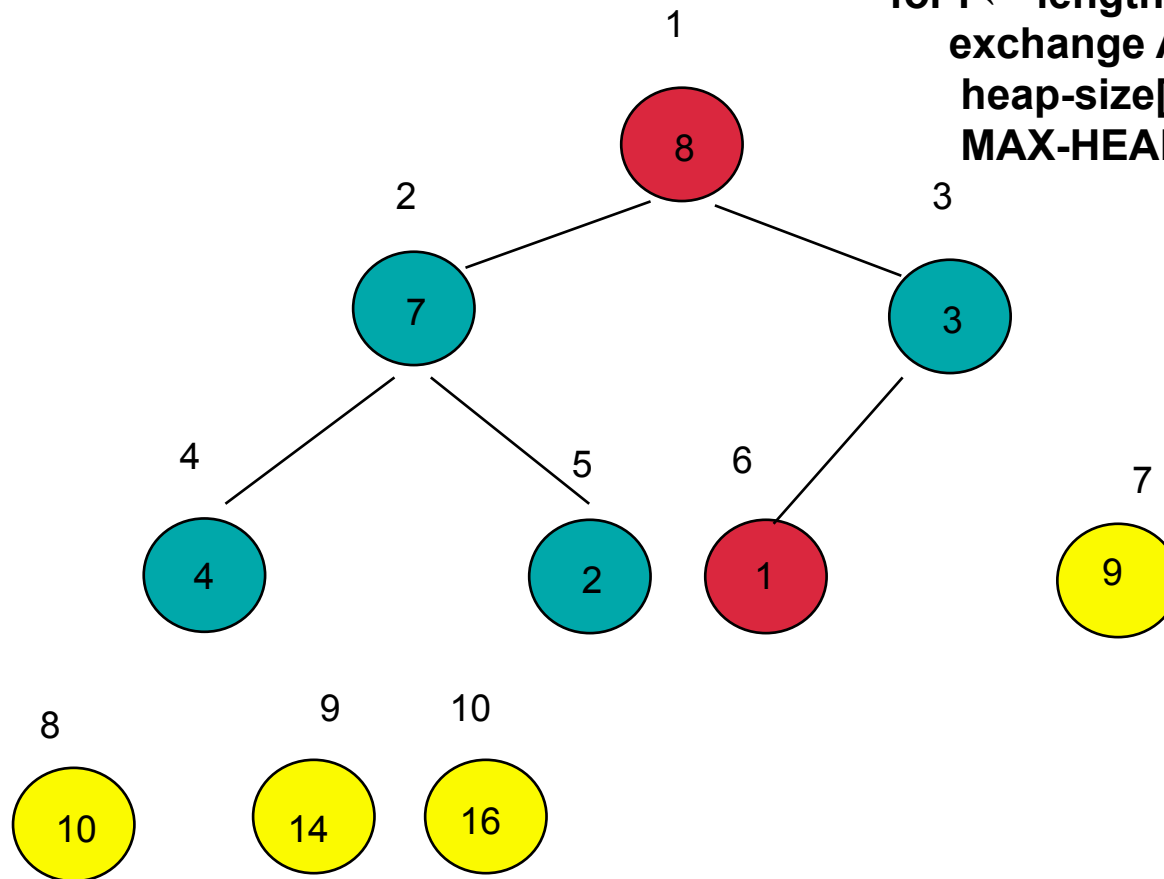


BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



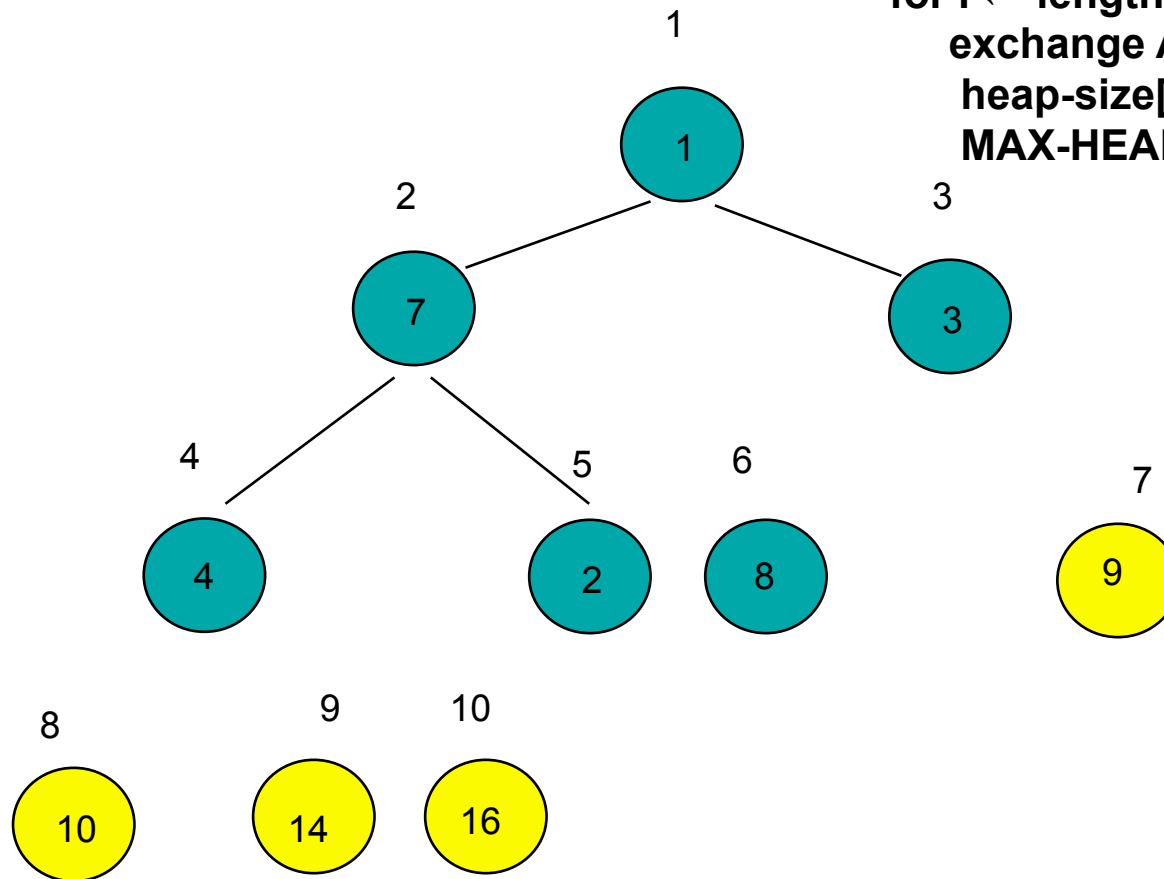
8	7	3	4	2	1	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



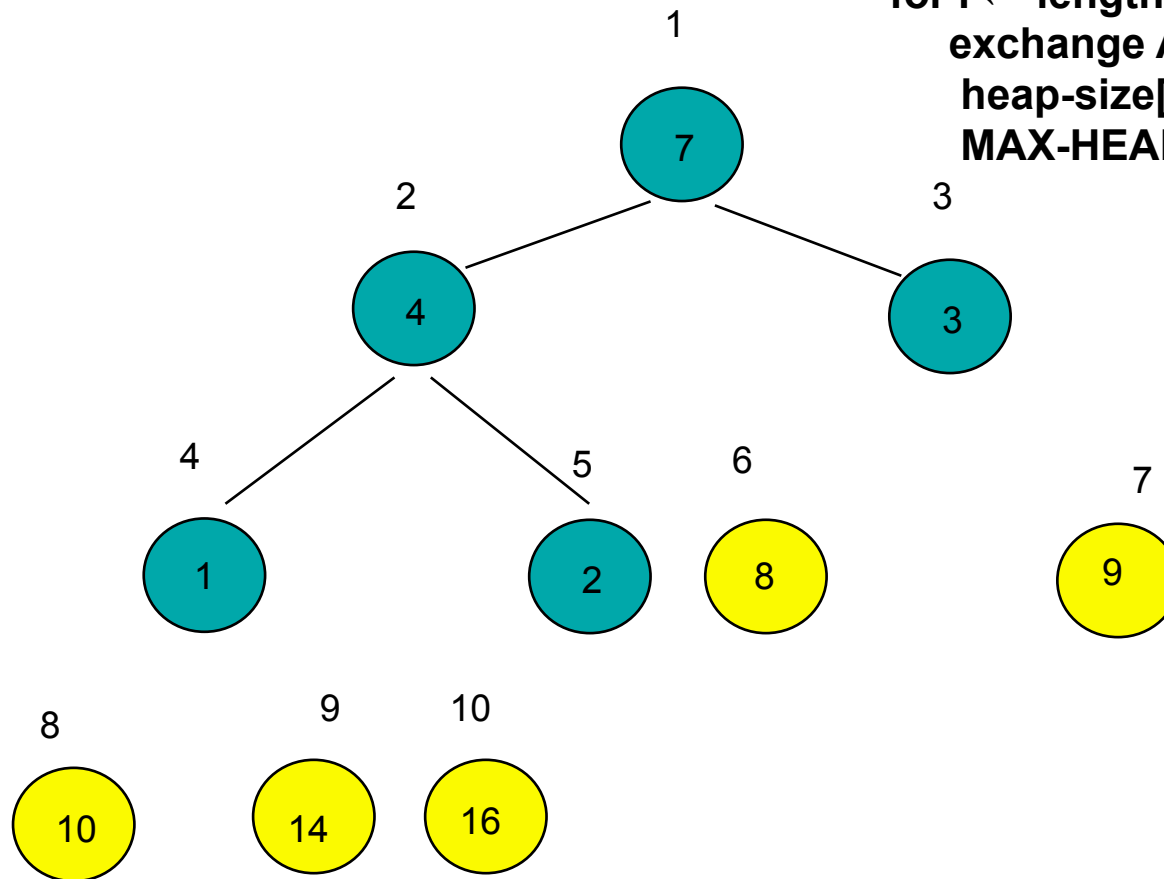
8	7	3	4	2	1	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)



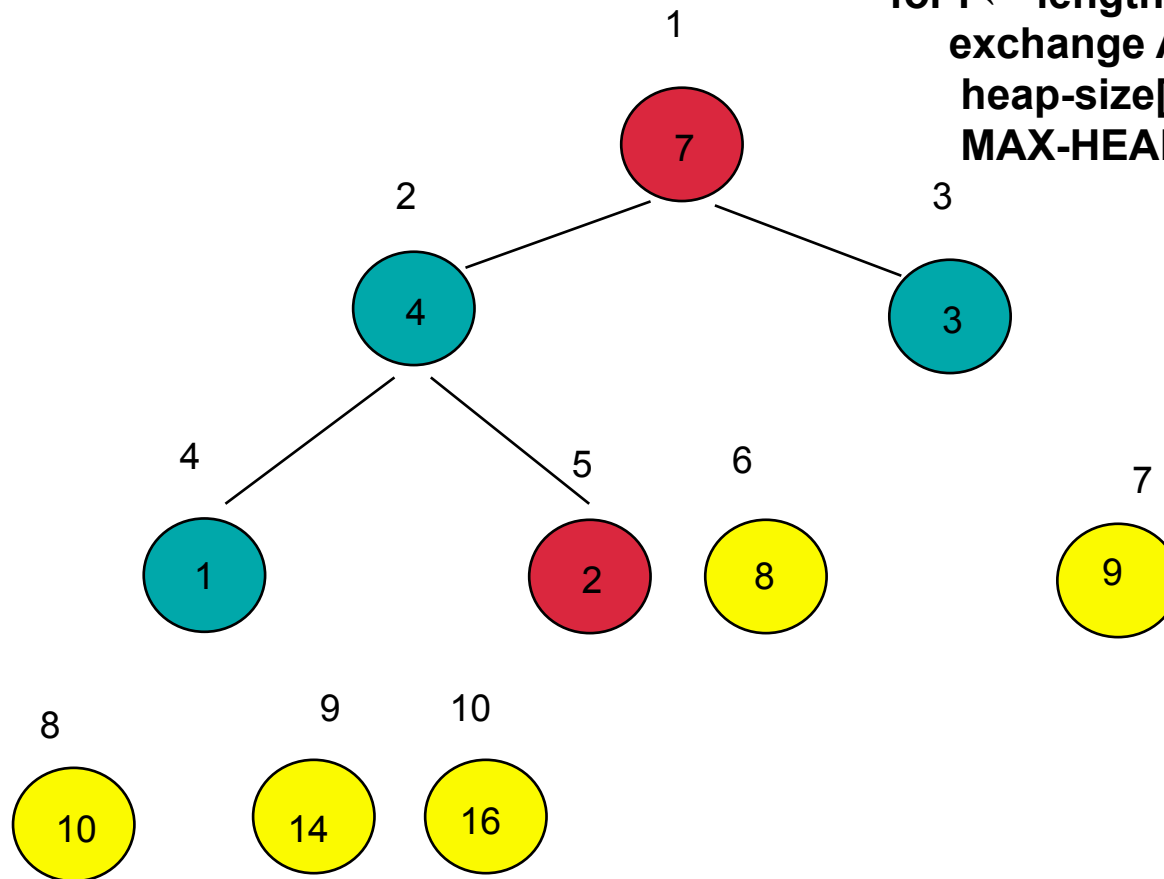
1	7	3	4	2	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)



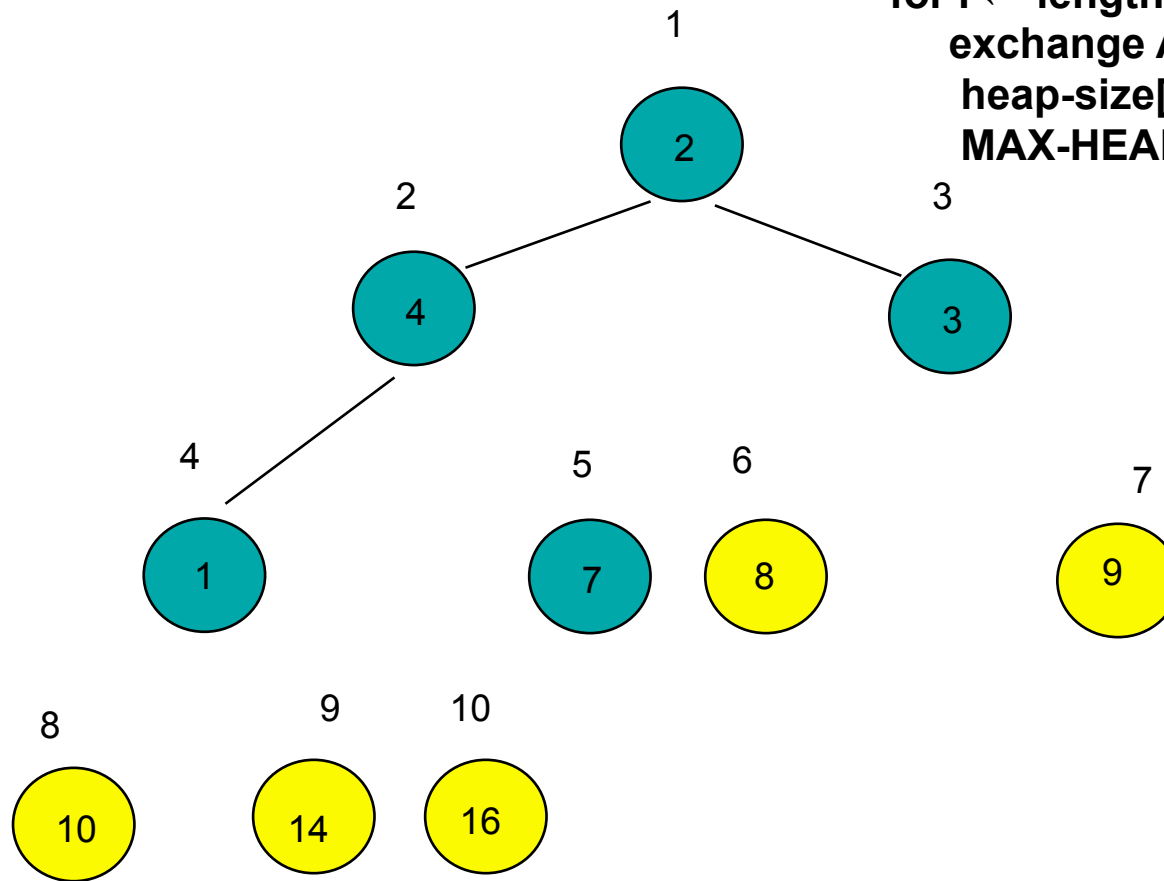
7	4	3	1	2	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)



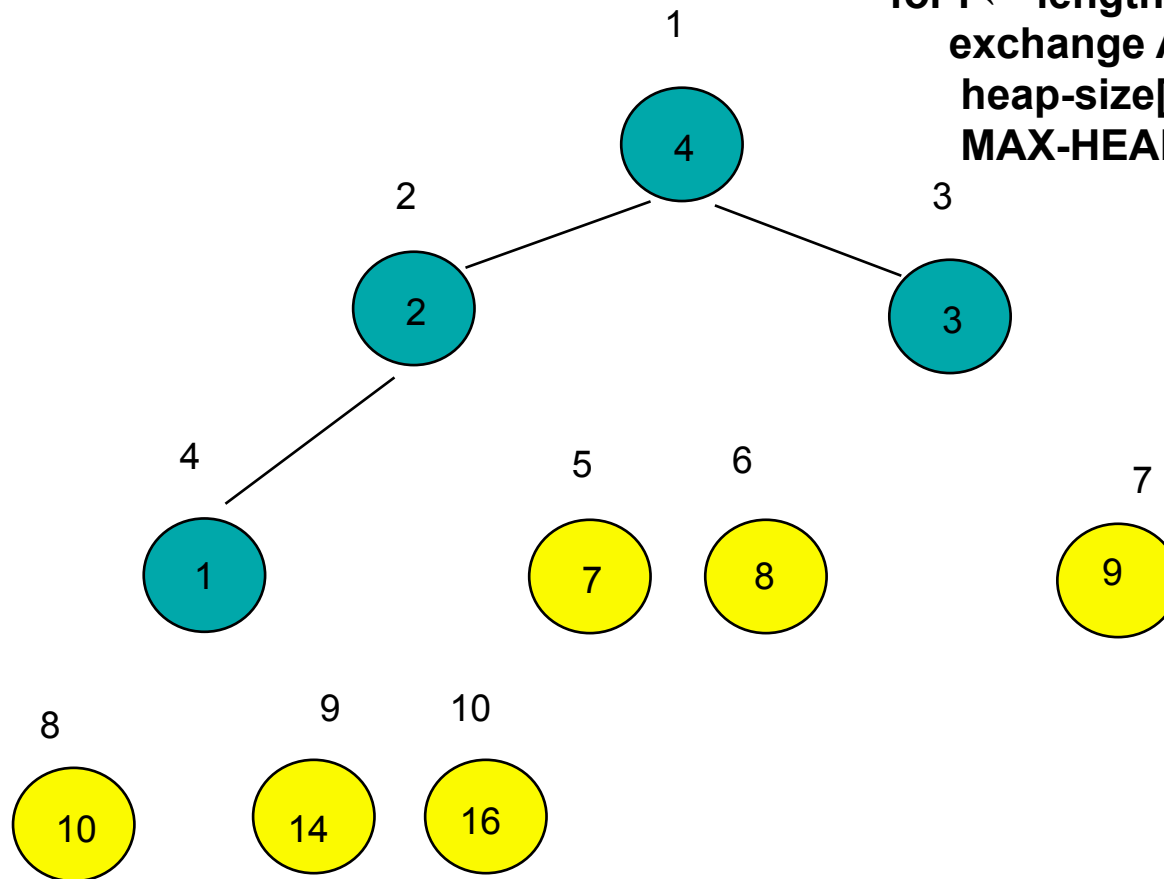
7	4	3	1	2	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



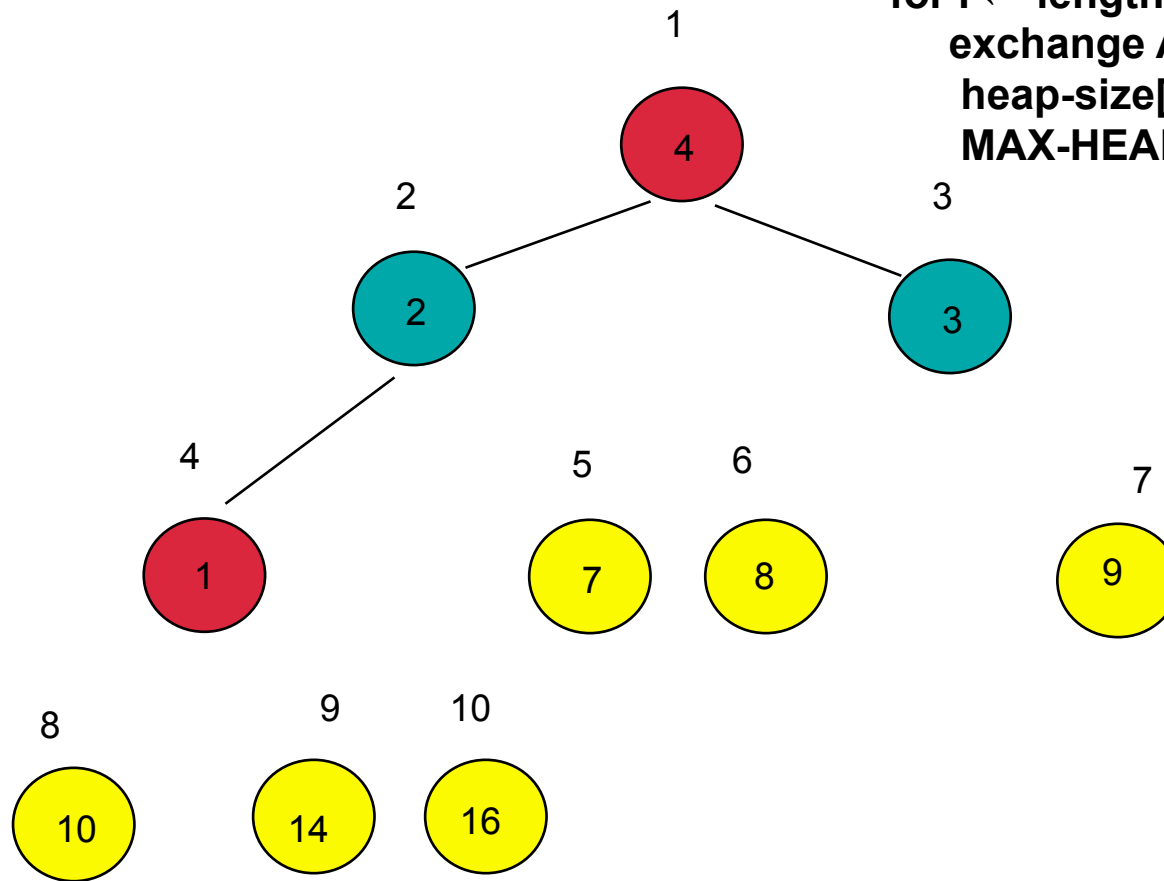
2	4	3	1	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)



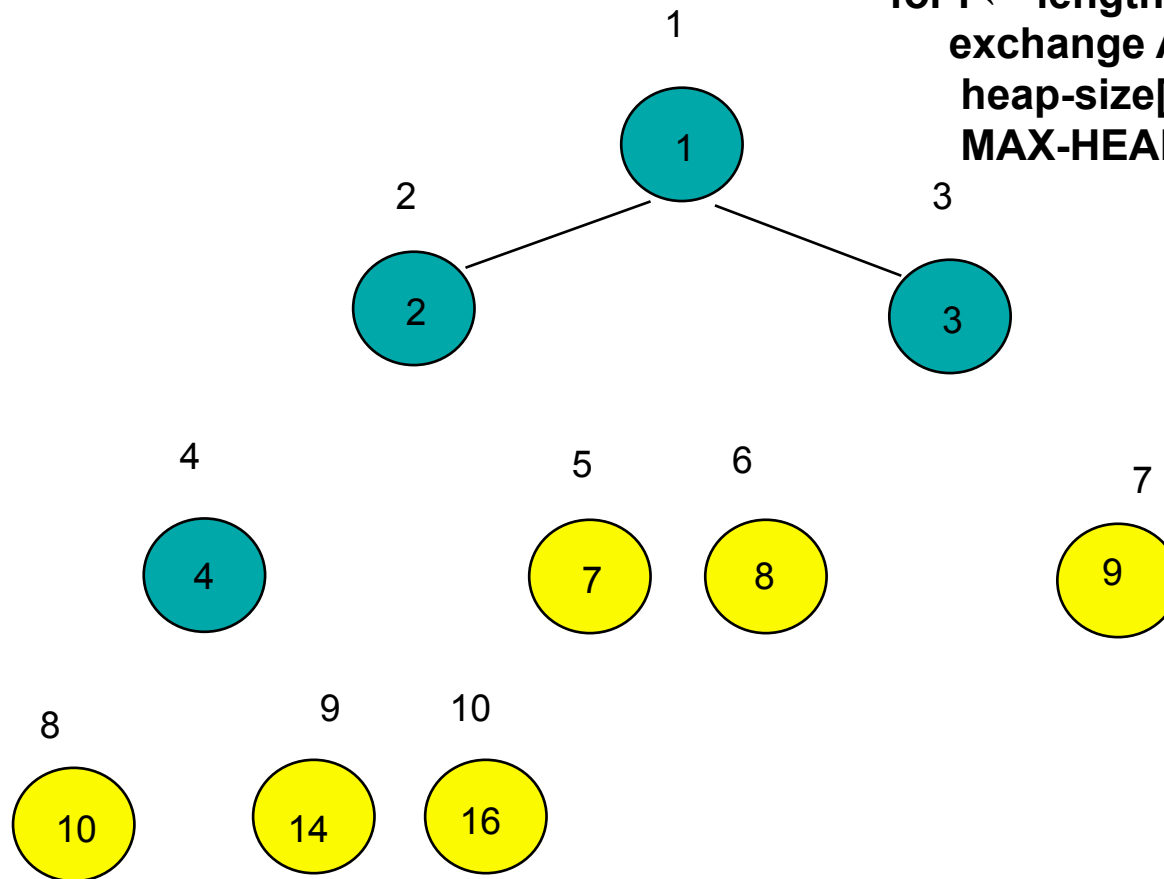
4	2	3	1	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



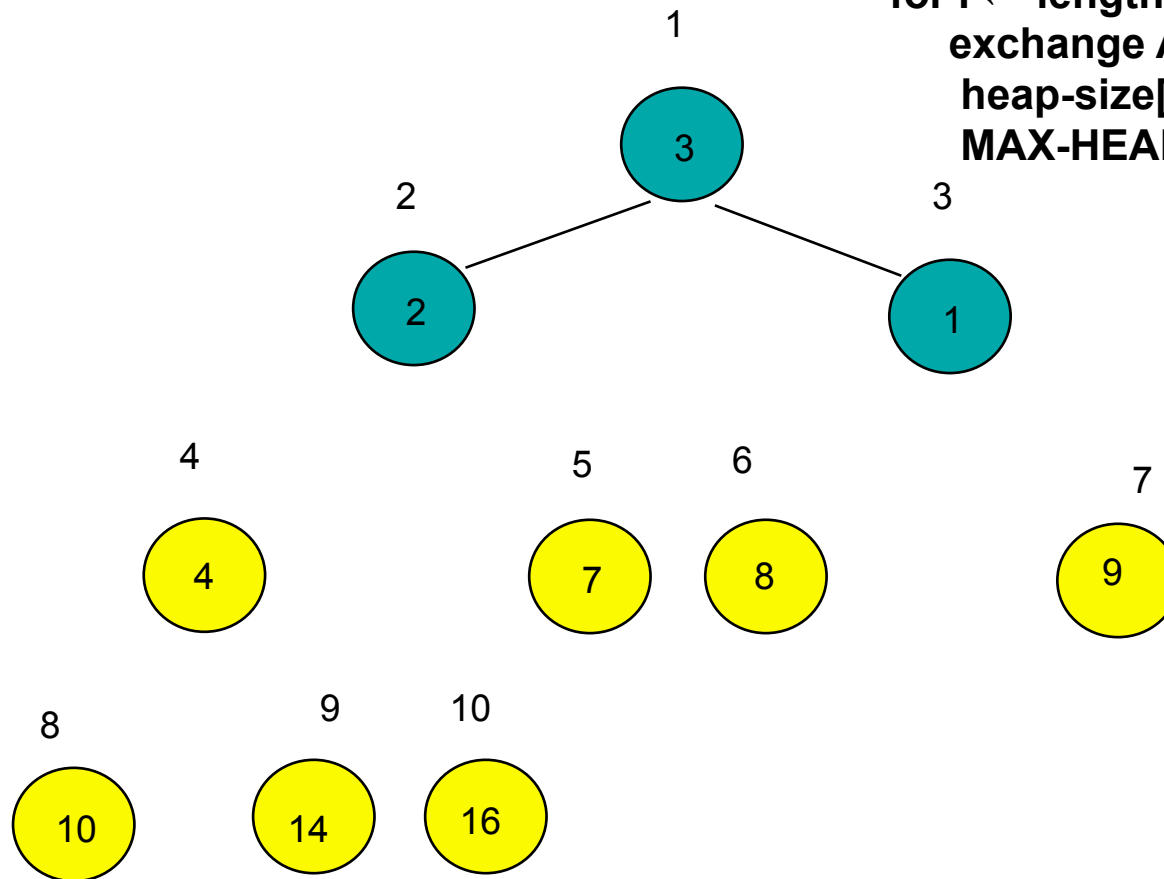
4	2	3	1	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



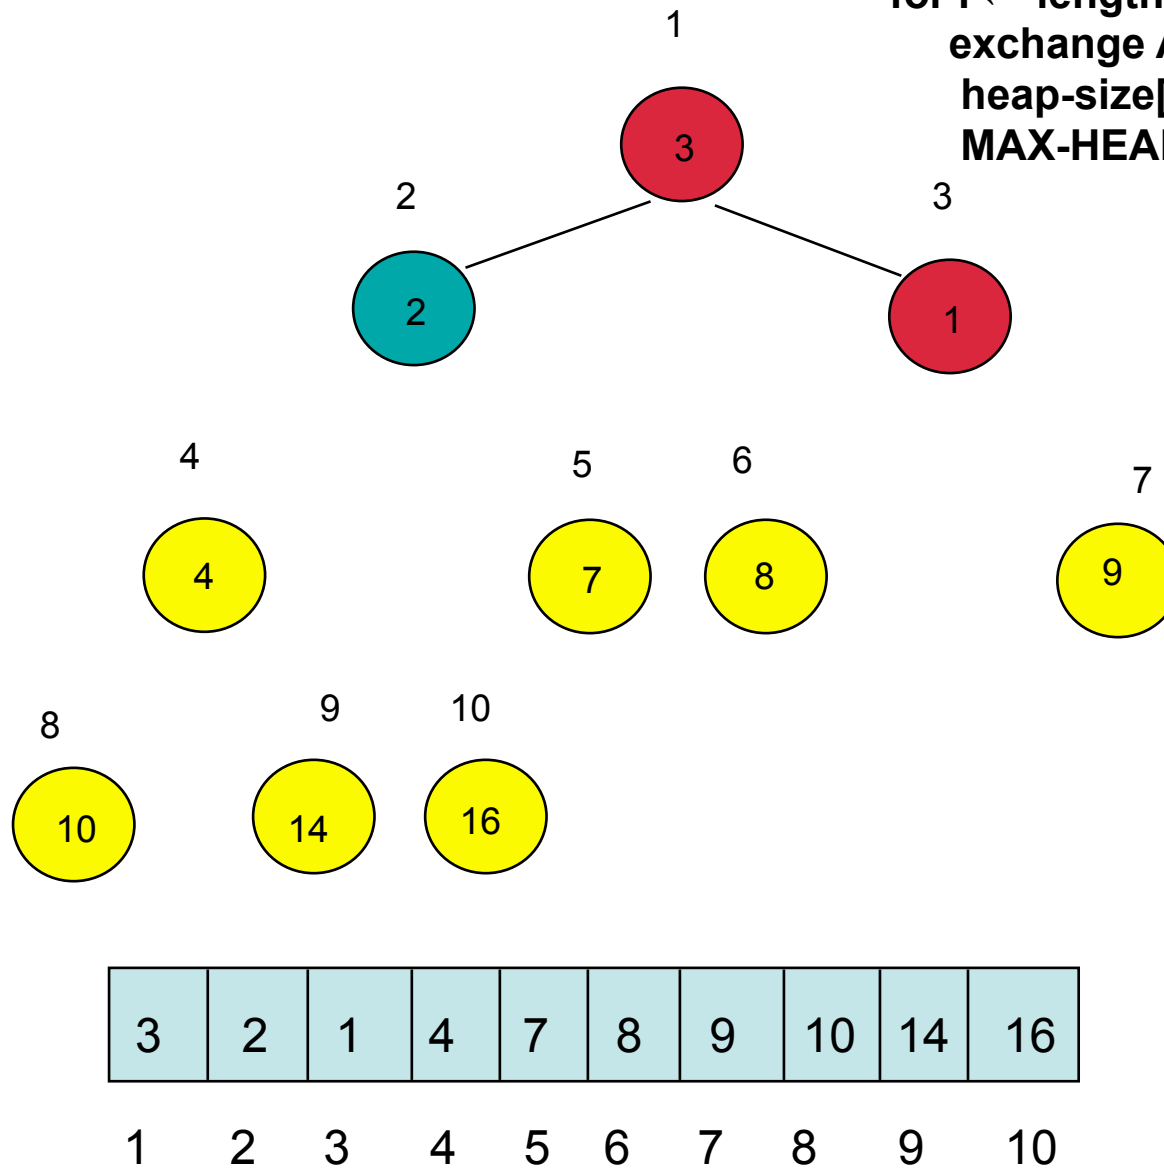
1	2	3	4	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)

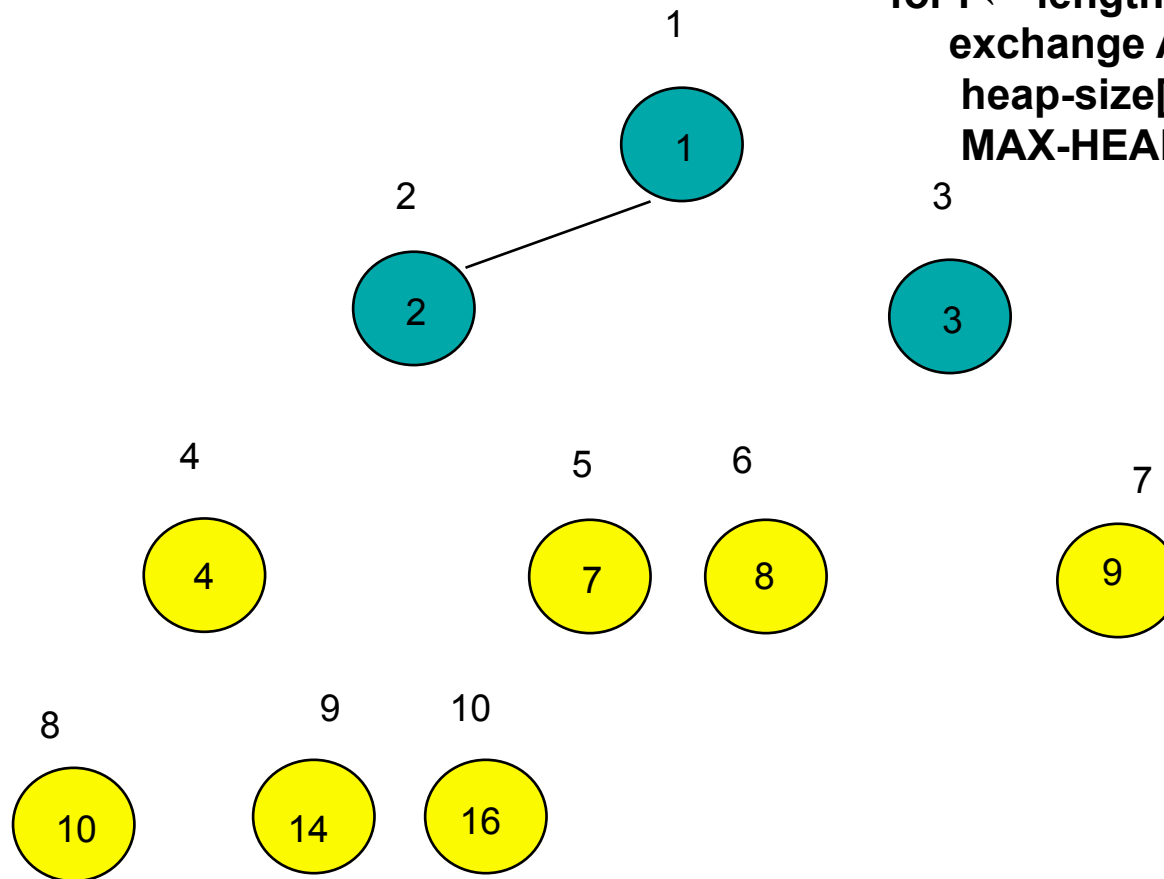


3	2	1	4	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)

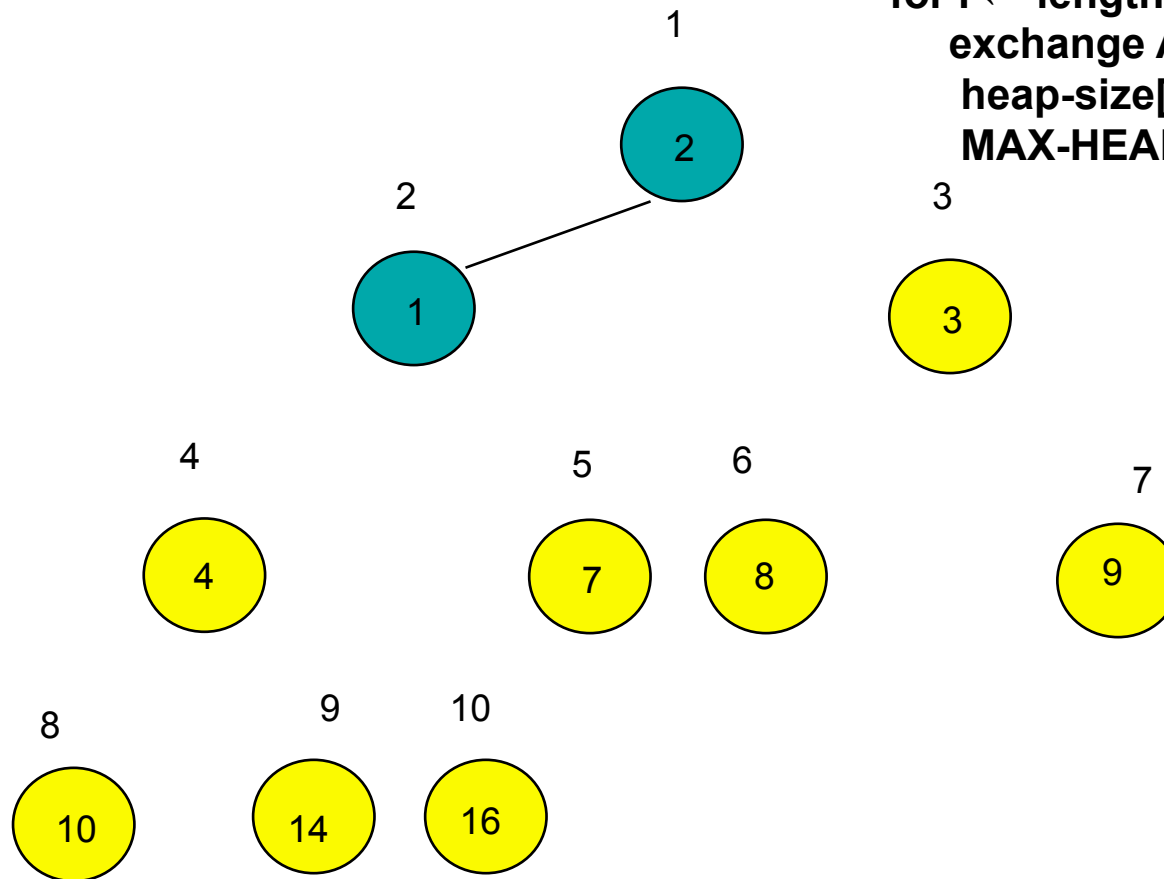


BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



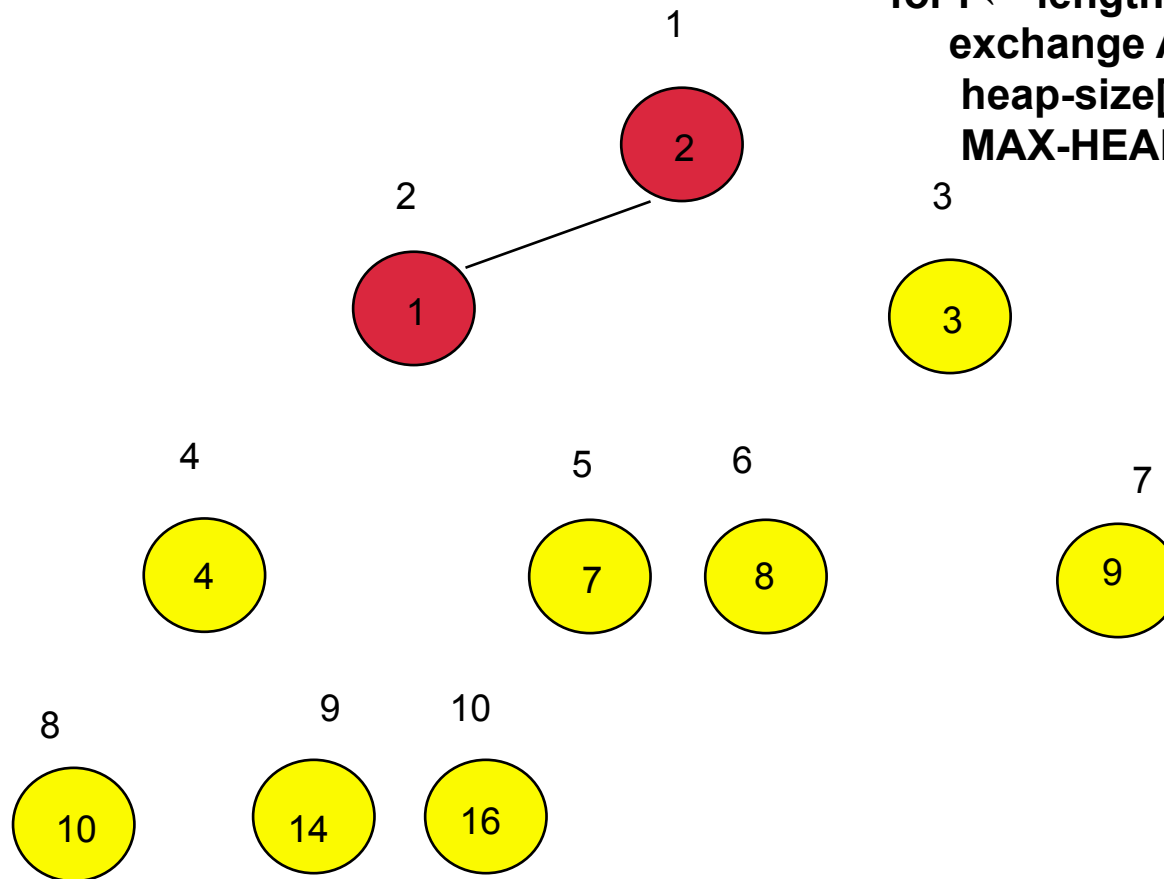
1	2	3	4	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



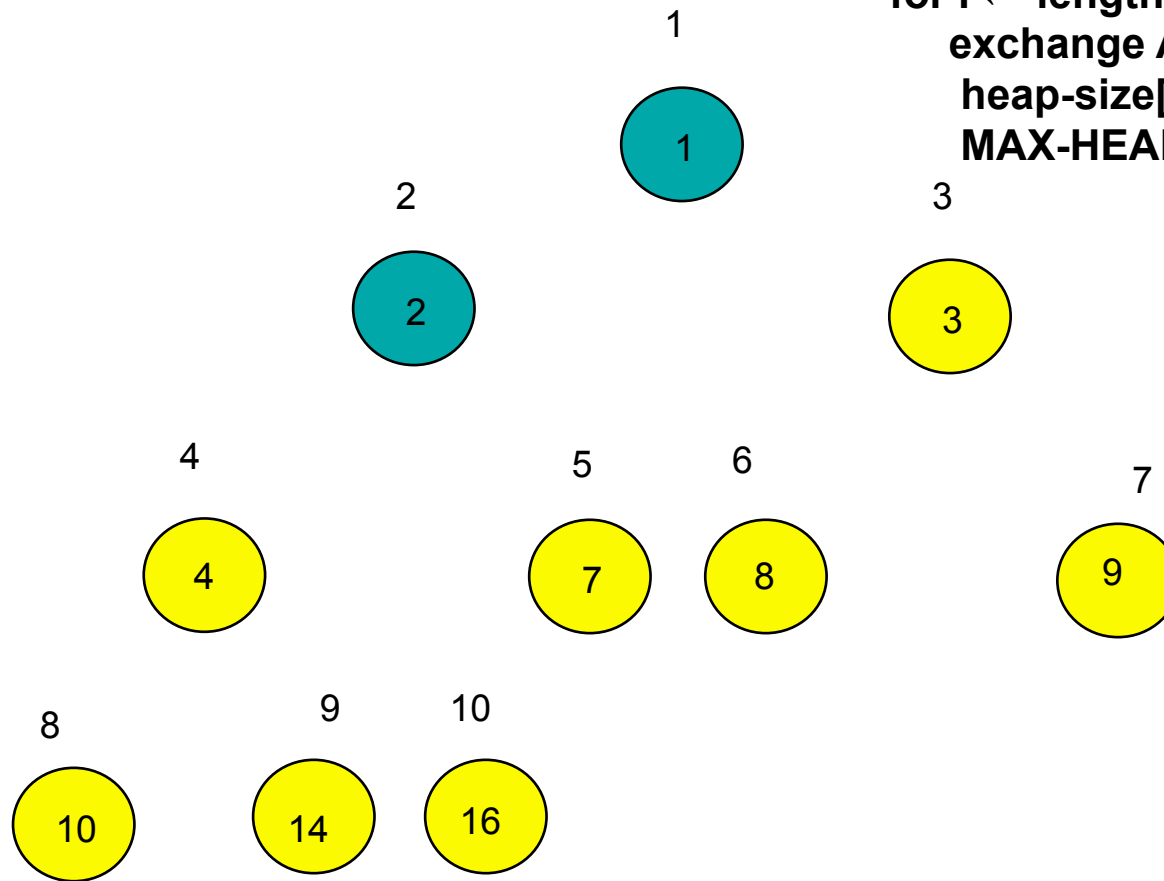
2	1	3	4	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



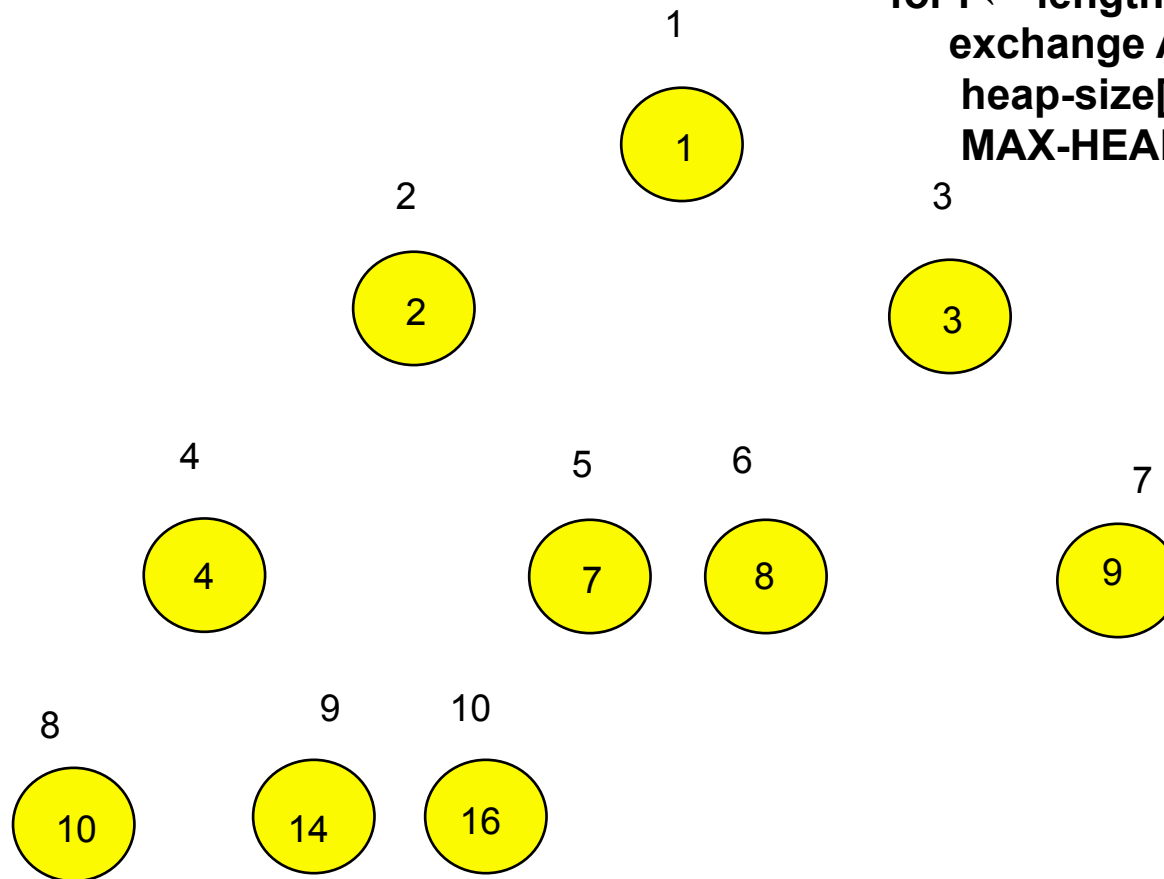
2	1	3	4	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 MAX-HEAPIFY(A, 1)



1	2	3	4	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

BUILD-MAX-HEAP(A)
 for $i \leftarrow \text{length}[A]$ downto 2 do
 exchange $A[1] \leftrightarrow A[i]$
 heap-size[A] \leftarrow heap-size[A] - 1
 MAX-HEAPIFY(A, 1)



1	2	3	4	7	8	9	10	14	16
1	2	3	4	5	6	7	8	9	10

Running time of Heapsort

HEAPSORT (A)

1 BUILD-MAX-HEAP (A)

2 for $i \leftarrow \text{length}[A]$ downto 2 do

3 exchange $A[1] \leftrightarrow A[i]$

4 heap-size[A] \leftarrow heap-size[A] - 1

5 MAX-HEAPIFY (A, 1)

Is there a loop? If so, how many times will it execute? What is the cost of one iteration of the loop?

Running time of Heapsort

HEAPSORT (A)

1	BUILD-MAX-HEAP (A)	$O(n)$
2	for $i \leftarrow \text{length}[A]$ downto 2 do	$O(n-1)$
3	exchange $A[1] \leftrightarrow A[i]$	$O(1)$
4	heap-size[A] \leftarrow heap-size[A] - 1	$O(1)$
5	MAX-HEAPIFY (A, 1)	$O(\lg n)$

Total time is:

$$O(n) + O(n-1) * [O(1) + O(1) + O(\lg n)]$$

which is approximately

$$O(n) + O(n \lg n)$$

or just

$$O(n \lg n)$$

Running time of Heapsort

- BUILD-MAX-HEAP takes $O(n)$.
- We have a loop. Each of the $n-1$ calls to MAX-HEAPIFY takes $O(\lg n)$ time.
- Total time is $O(n \lg n)$.
- Will heap sort always take $O(n \lg n)$ time?
Is there a best-case scenario? Is there a worst-case scenario? Why or why not?

Running time of Heapsort

- BUILD-MAX-HEAP takes $O(n)$.
- We have a loop. Each of the $n-1$ calls to MAX-HEAPIFY takes $O(\lg n)$ time.
- Total time is $O(n \lg n)$.
- Will heap sort always take $O(n \lg n)$ time? Is there a best-case scenario? Is there a worst-case scenario? Why or why not?

$$\sum_{i=1}^n \log(i)^c = \Theta(n \cdot \log(n)^c) \text{ for non-negative real } c$$

Space requirements of Heapsort

- Heapsort uses an array as its data structure.
- Heapsort sorts “in place”.
- Any extra storage needed?
- Only a negligible amount – one extra storage location is needed as temporary storage when swapping two array elements.

Priority Queues

- A priority queue is a data structure for maintaining a set S of elements, each with an associated value called a key.
- Applications include
 - scheduling jobs on a shared computer (max-priority queue)
 - event-driven simulators (min-priority queue)

Handles

- Elements of priority queue correspond to objects in application.
- We must be able to determine which application object corresponds to a given priority-queue element.
- We store a **handle** (pointer, integer, etc.) to the corresponding application object in each heap element.
- We also store a **handle** (array index) to the corresponding heap element in each application object.

Max-Priority Queue Operations

- $\text{INSERT}(S, x)$: insert element x into set S
- $\text{MAXIMUM}(S)$: return element of S with the largest key
- $\text{EXTRACT-MAX}(S)$: remove and return element of S with the largest key
- $\text{INCREASE-KEY}(S, x, k)$: increase value of x 's key to k , where k is at least as large as x 's current key value

Min-Priority Queue Operations

- $\text{INSERT}(S, x)$: insert element x into set S
- $\text{MINIMUM}(S)$: return element of S with the smallest key
- $\text{EXTRACT-MIN}(S)$: remove and return element of S with the smallest key
- $\text{DECREASE-KEY}(S, x, k)$: decrease value of x 's key to k , where k is at least as small as x 's current key value

Priority Queue Operations

- All operations can be done on a set of size n in $O(\lg n)$ time

HEAP-MAXIMUM

HEAP-MAXIMUM (A)

1 return A[1]

- Returns the item at the top of the heap
- Runs in $\Theta(1)$ time

HEAP-EXTRACT-MAX

HEAP-EXTRACT-MAX (A)

```
1  if heap-size[A] < 1
2      then error "heap underflow"
3  max ← A[1]
4  A[1] ← A[heap-size[A]]
5  heap-size[A] ← heap-size[A] - 1
6  MAX-HEAPIFY(A, 1)
7  return max
```


Running time of HEAP-EXTRACT-MAX

HEAP-EXTRACT-MAX (A)

1	if heap-size[A] < 1	$O(1)$
2	then error "heap underflow"	$O(1)$
3	max \leftarrow A[1]	$O(1)$
4	A[1] \leftarrow A[heap-size[A]]	$O(1)$
5	heap-size[A] \leftarrow heap-size[A] - 1	$O(1)$
6	MAX-HEAPIFY (A, 1)	$O(\lg n)$
7	return max	$O(1)$

Any loops? No. So just sum up the times: $O(6) + O(\lg n)$


The dominant term is $O(\lg n)$.

HEAP-INCREASE-KEY


HEAP-INCREASE-KEY(A, i, key)

```
1  if key < A[i]
2      then error "new key is smaller than
   current key"
3  A[i] ← key
4  while i > 1 and A[PARENT(i)] < A[i] do
5      exchange A[i] ↔ A[PARENT(i)]
6      i ← PARENT(i)
```


Example of HEAP-INCREASE-KEY




Example of HEAP-INCREASE-KEY (continued)



Example of HEAP-INCREASE-KEY (continued)



Example of HEAP-INCREASE-KEY (continued)



Running time of HEAP-INCREASE-KEY

HEAP-INCREASE-KEY(A, i, key)

1	if $\text{key} < A[i]$	$O(1)$
2	then error	$O(1)$
	“new key is smaller than current key”	
3	$A[i] \leftarrow \text{key}$	$O(1)$
4	while $i > 1$ and $A[\text{PARENT}(i)] < A[i]$ do	$O(\lg n)$
5	exchange $A[i] \leftrightarrow A[\text{PARENT}(i)]$	$O(3)$
6	$i \leftarrow \text{PARENT}(i)$	$O(1)$

Any loops? Yes. How many times will the loop execute? As many times as node i has ancestors, which = the depth of the tree. The depth of a binary tree is $O(\lg n)$. We do a constant amount of work in the loop. Cost is: $O(3) + O(4 \lg n)$, or just $O(\lg n)$


MAX-HEAP-INSERT

MAX-HEAP-INSERT (A, key)


1 heap-size[A] \leftarrow heap-size[A] + 1

2 A[heap-size] $\leftarrow -\infty$


3 HEAP-INCREASE-KEY (A, heap-size[A], key)




MAX-HEAP-INSERT(A,15)



HEAP-INSERT(A,15)



MAX-HEAP-INSERT(A,15)



MAX-HEAP-INSERT(A,15)

MAX-HEAP-INSERT

MAX-HEAP-INSERT (*A*, *key*)

1	<code>heap-size[A] ← heap-size[A] + 1</code>	$O(1)$
2	<code>A[heap-size] ← $-\infty$</code>	$O(1)$
3	<code>HEAP-INCREASE-KEY(A, heap-size[A], key)</code>	$O(\lg n)$

Any loops? No.

Add up the times: $O(1) + O(1) + O(\lg n) = O(2) + O(\lg n)$

Dominant term is $O(\lg n)$, so running time is just $O(\lg n)$.

Summary

- Definition of a heap
- How to build a heap
- How to use a heap for sorting
- How to analyze running time of heapsort
- Heapsort is in place (better than mergesort) and takes $O(n \lg n)$ (better than insertion)
- How to use a heap for priority queues

Notes

- For series sums, see the appendix at the end of the book or:
- <http://en.wikipedia.org/wiki/Summation>