



QUERY OPTIMIZATION




Query Optimization




- queries are composed using relational algebra operations and algorithms to implement these operations
- a user can express a query in a wide variety of ways
- query optimization: find a good plan

Relational Model: Query Optimization




- users do not worry about how best to state their queries
 - the DBMS is responsible for query optimization
 - a good optimizer has a wealth of information available to it that human users typically do not have
 - the number of distinct values of each type
 - the number of tuples currently appearing in each relation
 - the number of distinct values appearing in each attribute
 - a reoptimization may be required if the database statistics change over time
 - the optimizer is a program
 - much more patient than a typical user
 - embodying the skills and services of the best human programmers




Review

- alternative ways of evaluating a given query: query plans
 - access to the records
 - order of operations
 - algorithms to perform operations
- goal of query optimization: find a good query plan
 - finding absolute best plan usually not feasible
 - sufficient to find reasonably good plan




Two Approaches

- Heuristic Approach
 - apply heuristic rules to try to speed up query processing
- Cost-based Approach
 - estimate the cost of several execution plans
 - choose the one with the lowest cost




Heuristic Approach

- estimate statistics for intermediate results and reduce the size of intermediary tables
 - apply selection and projections before joining
 - apply most restrictive SELECT statements first
 - apply restricts and projections as early as possible



Example:

- find names of the directors of movies which have titles beginning with the letter Z
- assumption:
 - number of tuples in movie relation = 10000
 - number of tuples in person relation = 100
 - number of movies with titles beginning with Z = 50




Example (Cont.)

Total cost= 1 030 000 I/O operation

```
SELECT name FROM movie, person
WHERE movie.directorid = person.id
AND title like 'Z%'
```

- join movie and person relation:
 - read 10,000 tuples from movie relation
 - read each of the 100 person 10,000 times (=1,000,000)
 - construct an intermediate result consisting of 10,000 joined tuples and write them to disk
 - read the 10,000 joined tuples back into memory and find the movies with titles beginning with Z
- project the result over name
 - 50 tuples



Example (Cont.)


Total cost = 10 100 I/O operations

```
SELECT name FROM movie, person
```

```
WHERE title like 'Z%'
```


```
AND      movie.directorid = person.id
```

- restrict movie to just the tuples for movies with titles beginning with Z
 - read 10,000 tuples from movie relation
 - movies with titles beginning with Z: 50 tuples (can be kept in main memory)
- join the result to person relation
 - read 100 persons
 - produce a result 50 tuples (can be kept in main memory)
- project the result over name
 - 50 tuples




Example (Cont.)

- **total cost = 150 I/O operations**
 - unique index on title attribute of movie relation
 - read 50 tuples from movie relation to get the films that have a title starting with letter Z
- **total cost = 100 I/O operations**
 - unique index on title attribute of movie relation and id attribute of person relation
 - read 50 tuples from person relation



Query Optimization Process


- cast the query into internal form
- convert to canonical form
- choose candidate low-level procedures
- generate query plans and choose the cheapest



Query Optimization Process


- **cast the query into internal form**
- convert to canonical form
- choose candidate low-level procedures
- generate query plans and choose the cheapest

Cast the query into internal form



- conversion of the original query into some internal representation that is more suitable for machine manipulation
 - all queries in the external query language should be representable
 - it should not prejudice subsequent choices
 - abstract syntax tree or query tree

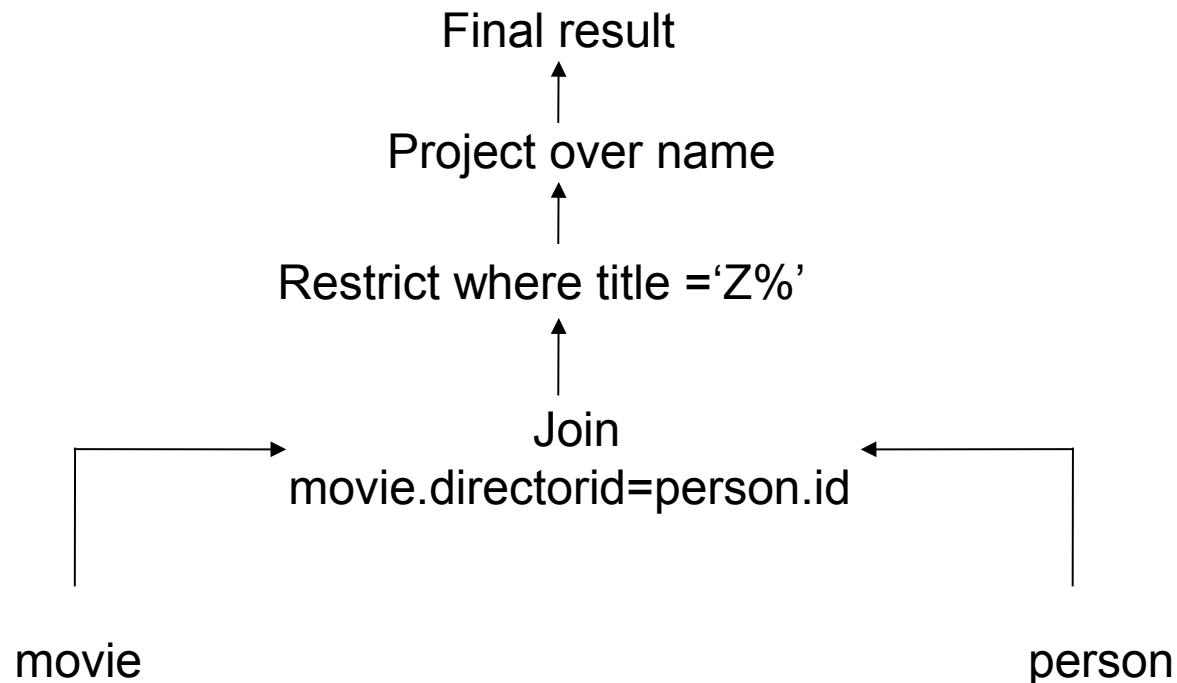
Query Tree




```
SELECT name FROM movie, person
```

```
WHERE movie.directorid = person.id
```


```
AND title like 'Z%'
```






Query Optimization Process

- cast the query into internal form
- **convert to canonical form**
- choose candidate low-level procedures
- generate query plans and choose the cheapest



Convert to canonical form

- queries can be expressed in a variety of ways
- convert the internal representation into some equivalent canonical form
 - guaranteed to be good
 - regardless of the actual data values and physical access paths



Expression Transformation

- a sequence of restrictions on the same relation can be transformed into a single *AND*ed restriction


```
SELECT ... FROM
```

```
    (SELECT ... FROM A WHERE p1) AS ...
```

```
    WHERE p2
```

```
SELECT ... FROM A WHERE (p1 AND p2)
```


Example: Expression Transformation



```
SELECT * FROM  
(SELECT * FROM movie WHERE score > 7) AS m  
WHERE year = 1985
```

- equivalent expression

```
SELECT * FROM movie  
WHERE score > 7 AND year = 1985
```




Expression Transformation

- in a sequence of projections against the same relation, all but the last can be ignored

```
SELECT ac12 FROM  
    (SELECT ac11 FROM A) AS ...  
SELECT ac12 FROM A
```


Example: Expression Transformation



```
SELECT title FROM  
(SELECT title, year FROM movie) AS m
```

- equivalent expression

```
SELECT title FROM movie
```




Expression Transformation

- a restriction of a projection can be transformed into a projection of a restriction


```
SELECT ... FROM  
    (SELECT acl FROM A) AS ...  
WHERE p
```

```
SELECT acl FROM  
    (SELECT * FROM A WHERE p) AS ...  
SELECT acl FROM A WHERE p
```



Distributivity

- f is said to be distributive over \circ if
$$f(A \circ B) \equiv f(A) \circ f(B)$$
- restriction distributes over union, intersection and difference
 - in some cases over join
- projection distributes over union and intersection
 - projection does not distribute over difference



Expression Transformation

- other transformation laws

- commutativity

$$A \circ B \equiv B \circ A$$

- associativity


$$A \circ (B \circ C) \equiv (A \circ B) \circ C$$

- idempotence

$$A \circ A \equiv A$$

- absorption

$$A \text{ UNION } (A \text{ INTERSECT } B) \equiv A$$




Semantic Transformations

```
SELECT MOVIEID FROM CASTING, MOVIE
WHERE (CASTING.MOVIEID=MOVIE.ID)
```

```
SELECT MOVIEID FROM CASTING
```


- due to the integrity constraint
 - foreign to matching candidate key join
 - every casting tuple does join to some movie tuple



Query Optimization Process


- cast the query into internal form
- convert to canonical form
- **choose candidate low-level procedures**
- generate query plans and choose the cheapest

Divide-and-Conquer Strategy




- break down complex queries into a sequence of smaller queries
 - detachment
 - tuple substitution
- example: names starting with 'A' of the lead actors of the movies with votes more than 500 and made after 1990

Example




```
SELECT DISTINCT ACTOR.NAME
FROM ACTOR,CASTING,MOVIE
WHERE (ACTOR.NAME LIKE 'A%')
      AND (ACTOR.ID=CASTING.ACTORID)
      AND (CASTING.ORD=1)
      AND (CASTING.MOVIEID=MOVIE.ID)
      AND (MOVIE.YEAR>1990)
      AND (MOVIE.VOTES>500)
```

Example




```
D1: SELECT ID FROM MOVIE
      WHERE (YEAR>1990)
           AND (VOTES>500)
SELECT DISTINCT ACTOR.NAME
FROM ACTOR,CASTING,D1
WHERE (ACTOR.NAME LIKE 'A%')
      AND (ACTOR.ID=CASTING.ACTORID)
      AND (CASTING.ORD=1)
      AND (CASTING.MOVIEID=D1.ID)
```

Example




```
D2: SELECT ACTORID,MOVIEID FROM CASTING
      WHERE (ORD=1)
SELECT DISTINCT ACTOR.NAME
FROM ACTOR,D2,D1
WHERE (ACTOR.NAME LIKE 'A%')
      AND (ACTOR.ID=D2.ACTORID)
      AND (D2.MOVIEID=D1.ID)
```

Example




```
D3: SELECT ID,NAME FROM ACTOR
      WHERE (NAME LIKE 'A%')
SELECT DISTINCT D3.NAME
FROM D3,D2,D1
WHERE (D3.ID=D2.ACTORID)
      AND (D2.MOVIEID=D1.ID)
```

Example




```
D4: SELECT D2.ACTORID FROM D2,D1
      WHERE (D2.MOVIEID=D1.ID)
SELECT DISTINCT D3.NAME
      FROM D3,D4
      WHERE (D3.ID=D4.ACTORID)
```




Query Optimization Process

- cast the query into internal form
- convert to canonical form
- choose candidate low-level procedures
- **generate query plans and choose the cheapest**

Generate query plans and choose the cheapest




- construct a set of candidate query plans
 - a set of candidate implementation procedures, one procedure for each of the low-level operation in the query
 - instead of generating all plans, generate a subset of plans using heuristic techniques
- assign a cost to any given plan
 - sum of the costs of the individual procedures
 - cost formulas depend on the size(s) of relations to be processed and intermediate results
 - use database statistics




Database Statistics

- for each base table:
 - number of tuples
- for each column of each table:
 - number of distinct values in this column
 - max./min. value, second max./min. value, average value
 - for indexed columns, the ten most frequently occurring values

Implementing the Relational Operators




- brute force
- index lookup
- hash lookup
- merge
- hash
- combinations of the above methods



Example

- join two relations: R, S
 - number of tuples in relation $R = m$
 - number of tuples of relation R on a page = pR
 - number of tuples in relation $S = n$
 - number of tuples of relation S on a page = pS
- C is the common attribute
- dCR : number of distinct values of the C attribute in relation R
- dCS : number of distinct values of the C attribute in relation S




Brute Force

```
do i := 1 to m
  do j := 1 to n
    if R[i].C = S[j].C then
      add joined tuple R[i]*S[j] to
result;
end;
end;
```

- **Cost:**

- $m + (m * n)$
- $(m / pR) + (m * n) / pS$



Index Lookup

- an index X on attribute C of S

- $X[1] \dots X[k]$ indexes

```
do i := 1 to m
```

```
  do j := 1 to k
```

```
    /* let tuple of S indexed by X[j] be S[j] */
```

```
      if  $R[i].C = S[j].C$  then
```

```
        add joined tuple  $R[i] * S[j]$  to
```

```
        result;
```


```
      end;
```

```
end;
```

- **Cost:**

- $(m / p_R) + ((m * n) / d_{CS})$

- $(m / p_R) + ((m * n) / d_{CS}) / p_S$



Hash Lookup

- S.C is a hash

```
do i := 1 to m
```

```
  k := hash (R[i].C);
```

```
  /* let there be h tuples stored at */
```

```
  do j := 1 to h
```

```
    if R[i].C = S[j].C then
```

```
      add joined tuple R[i] * S[j] to result;
```

```
    end;
```

```
end;
```