

**Istanbul Technical University  
Faculty of Computer and Informatics**



**BLG413E System Programming  
Project 2 Report  
Group 28**

**Cem Yusuf Aydoğdu  
150120251**

## 1. Introduction

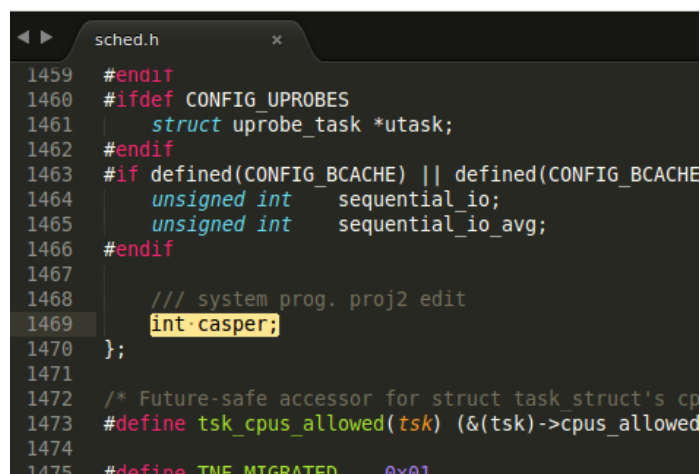
In this project, a new integer variable(**int casper**) which is used to control visibility of a process in **/proc** filesystem was added to the task descriptor and a **set\_casper** system call was implemented in order to change the control value of a given process.

Required visibility behaviours for each **casper** value are shown below:

Casper value	Required behaviour
0	Process is visible to all
1	Process is visible to only processes with same used id
2	Process is visible to only processes with same group id
3	Process is not visible for any process

## 2. Changes in kernel

Initially, **int casper** variable was added to the end of **struct task\_struct** in **/include/linux/sched.h** file. **struct task\_struct** is task descriptor which holds information(pid, state, flags...) about a process. For every process, there exists a task descriptor in the system.

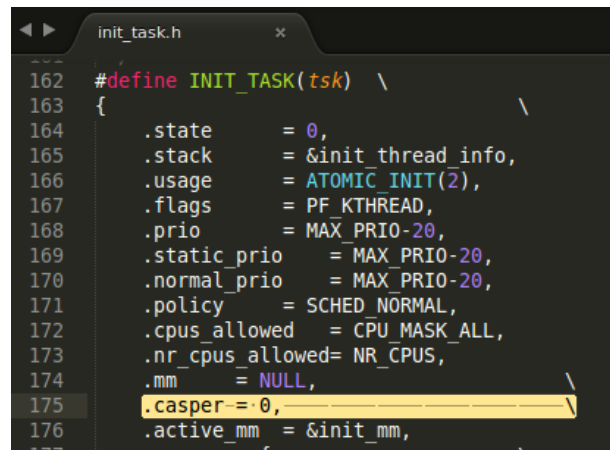


```
1459 #endif
1460 #ifdef CONFIG_UBROBES
1461     struct uprobe_task *utask;
1462 #endif
1463 #if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE)
1464     unsigned int sequential_io;
1465     unsigned int sequential_io_avg;
1466 #endif
1467
1468     /// system prog. proj2 edit
1469     int casper;
1470 };
1471
1472 /* Future-safe accessor for struct task_struct's cp
1473 #define tsk_cpus_allowed(tsk) (&(tsk)->cpus_allowed
1474
1475 #define TNE_MIGRATED 0x01
```

*Screenshot 1: /include/linux/sched.h*

After that, **INIT\_TASK** macro in `/include/linux/init_task.h` file was changed in order to initialize **casper** value to zero for Process 0.

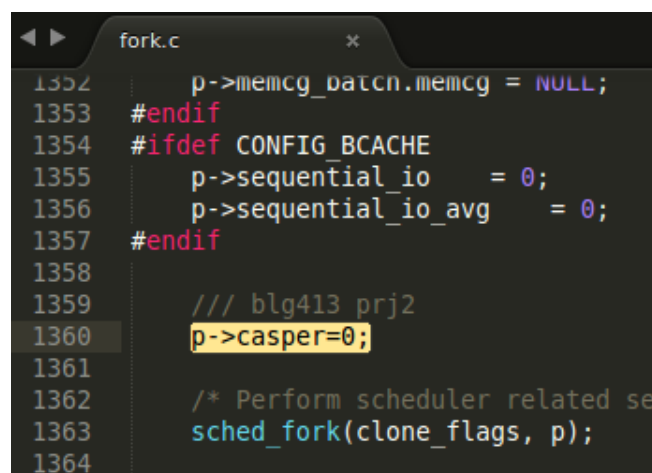
Process 0 is the first process created in kernel startup stage which is responsible for initialization of some core system functionalities. **INIT\_TASK** macro only initializes this process.



```
162 #define INIT_TASK(tsk) \
163 { \
164     .state      = 0, \
165     .stack      = &init_thread_info, \
166     .usage      = ATOMIC_INIT(2), \
167     .flags      = PF_KTHREAD, \
168     .prio       = MAX_PRIO-20, \
169     .static_prio = MAX_PRIO-20, \
170     .normal_prio = MAX_PRIO-20, \
171     .policy     = SCHED_NORMAL, \
172     .cpus_allowed = CPU_MASK_ALL, \
173     .nr_cpus_allowed = NR_CPUS, \
174     .mm         = NULL, \
175     .casper     = 0, \
176     .active_mm  = &init_mm, \
177     .sched_info = { \
```

*Screenshot 2: /include/linux/init\_task.h*

Then, **casper** was assigned to zero in **copy\_process** function in the **kernel/fork.c** file. This file contains functions about fork system call which is used to create processes. **copy\_process** function is responsible for the core functionality of the new process creation. This function creates a new **task\_struct** with information from a given old process, but it does not start the created process.

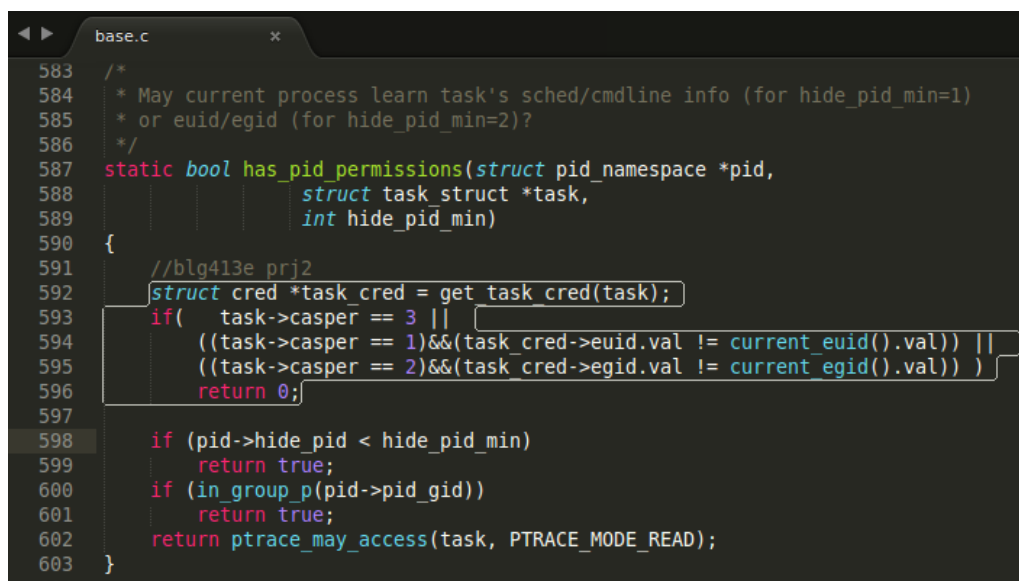


```
1352     p->memcg_data.memcg = NULL;
1353 #endif
1354 #ifdef CONFIG_BCACHE
1355     p->sequential_io      = 0;
1356     p->sequential_io_avg  = 0;
1357 #endif
1358
1359     /// blg413 prj2
1360     p->casper=0;
1361
1362     /* Perform scheduler related se
1363     sched_fork(clone_flags, p);
1364
```

*Screenshot 3: /kernel/fork.c*

Finally, **has\_pid\_permissions** function in **/fs/proc/base.c** file was modified to achieve different visibility requirements . This function checks whether the current process(given as **\*pid** in the first argument) can or can not see the task's(given as **\*task** in second argument) information. If the current process is not authorized, this function returns zero, vice versa.

In order to implement different visibility features, casper value of the task was obtained from **get\_task\_cred** function and compared with required values. If the **casper** value is 3, or **casper** value is 1 but user ids of task and current process does not match, or **casper** value is 2 but group ids of task and current process does not match; it returns 0 as false.

A screenshot of a code editor showing the implementation of the `has_pid_permissions` function in the file `base.c`. The function signature is `static bool has_pid_permissions(struct pid_namespace *pid, struct task_struct *task, int hide_pid_min)`. The function body starts with a comment `//blg413e prj2` and then calls `struct cred *task_cred = get_task_cred(task);`. It follows with a conditional check: `if( task->casper == 3 || ((task->casper == 1)&&(task_cred->euid.val != current_euid().val)) || ((task->casper == 2)&&(task_cred->egid.val != current_egid().val)) )`. If this condition is true, it returns 0. Otherwise, it checks `if (pid->hide_pid < hide_pid_min)` and returns true. Then it checks `if (in_group_p(pid->pid_gid))` and returns true. Finally, it returns `ptrace_may_access(task, PTRACE_MODE_READ);`.

```
583 /*
584  * May current process learn task's sched/cmdline info (for hide_pid_min=1)
585  * or euid/egid (for hide_pid_min=2)?
586  */
587 static bool has_pid_permissions(struct pid_namespace *pid,
588                                struct task_struct *task,
589                                int hide_pid_min)
590 {
591     //blg413e prj2
592     struct cred *task_cred = get_task_cred(task);
593     if( task->casper == 3 ||
594         ((task->casper == 1)&&(task_cred->euid.val != current_euid().val)) ||
595         ((task->casper == 2)&&(task_cred->egid.val != current_egid().val)) )
596         return 0;
597
598     if (pid->hide_pid < hide_pid_min)
599         return true;
600     if (in_group_p(pid->pid_gid))
601         return true;
602     return ptrace_may_access(task, PTRACE_MODE_READ);
603 }
```

*Screenshot 4: /fs/proc/base.c*

### 3. System Call

```
#include <linux/syscalls.h>
#include <linux/kernel.h>
#include <linux/sched.h>
#include <asm/errno.h>

asmlinkage long set_casper(pid_t pid, int value)
{
    if(value != 0 && value != 1 && value != 2 && value != 3)
    {
        return -EINVAL;
    }

    struct task_struct *p;
    read_lock(&tasklist_lock);
    p=find_task_by_vpid(pid);
    if(p == NULL)
    {
        read_unlock(&tasklist_lock);
        return -ESRCH;
    }
    read_unlock(&tasklist_lock);
    write_lock_irq(&tasklist_lock);
    p->casper = value;
    write_unlock_irq(&tasklist_lock);

    return 0;
}
```

The system call assigns *casper* value to a process with given pid. Initially, the system call returns “Invalid argument” error (**EINVAL**) if the given casper value is invalid.

Then, it uses a **read\_lock** in order to lock access to the critical section. After that, it finds the process using **find\_task\_by\_vpid** function. If there are no process with that given **pid**, it first unlocks the read spin lock and returns “No such process” error (**ESRCH**).

If it finds the process, it again unlocks the read lock and uses a write lock with **write\_lock\_irq** function which disables interrupts, in order to change *casper* value.

In order to add the system call, it was added to the end of the system call table (**arch/x86/syscalls/syscall\_32.tbl**) and system call header file which is in **include/linux/syscalls.h**.

## 4. Test Code

Finally, after all implementations the kernel was compiled and installed to the system. Then, a basic test program was written to test functionalities.

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#define NR_set_casper 355

int main(int argc, char **argv)
{
    if(argc!=3 || geteuid() != 0)
    {
        printf("usage: sudo ./executable pid value\n");
        return -1;
    }
    int pid = atoi(argv[1]), val=atoi(argv[2]);
    printf("casper value %d, pid %d\n",val, pid);
    long result = syscall(NR_set_casper, pid, val);
    printf("syscall result:%ld \n ",result);
    if(result == 0)
        printf("Success\n");
    else
    {
        printf("Error: ");
        if(errno== EINVAL)
            printf("Invalid value\n");
        if(errno==ESRCH)
            printf("Process not found\n");
    }
    return 0;
}
```

Number of the system call was denoted with **NR\_set\_casper** macro. Program initially checks for argument number and super user privilege conditions. Then, it parses first argument as **pid** and second argument as **casper** value.

After, it uses **syscall** function to call **set\_casper** system call and stores the result in a long integer. It informs about success and if an error occurs, it prints the required error message before exiting.

## 5. Results

In order to test the required functionalities, process id of bash was found as 1593 with \$\$ command. Also, an arbitrary program with pid number 1703 is called from bash. That process has the same user and group id with bash.

```
itucs@ubuntu
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ $
1593: command not found
127 ✗ itucs@ubuntu ~/Desktop/prj2 $
```

*Screenshot 5: Process ID of the bash*

```
itucs@ubuntu: ~/Desk
itucs@ubuntu: ~/Des
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ leafpad &
[1] 1703
0 ✓ itucs@ubuntu ~/Desktop/prj2 $
```

*Screenshot 6: A new process with pid=1703*

Since **casper=0** initially, any process should be visible on the /proc filesystem. Testing with different processes from different users proves that.

```
itucs@ubuntu: ~/Desktop/prj2
itucs@ubuntu: ~/Desktop/prj2 80
~/Desktop/prj2 $ ls /proc/1703
cpuset  latency  mountstats  pro
cwd     limits  net         roo
environ loginuid  ns          sch
exe     map_files oom_adj     sch
fd      maps    oom_score   ses
fdinfo  mem     oom_score_adj set
gid_map mountinfo pagemap     sma
io      mounts  personality sta
~/Desktop/prj2 $
```

*Screenshot 7: Process with same user and group id, when casper=0*

```
itucs@ubuntu: ~/Desktop/prj2
itucs@ubuntu: ~/Desktop/prj2 80
~/Desktop/prj2 $ ls /proc/1
symbolic link /proc/1/cwd: Permission d
symbolic link /proc/1/root: Permission
symbolic link /proc/1/exe: Permission d
cpuset  latency  mountstats  pro
cwd     limits  net         roo
environ loginuid  ns          sch
exe     map_files oom_adj     sch
fd      maps    oom_score   ses
fdinfo  mem     oom_score_adj set
gid map mountinfo pagemap     sma
```

*Screenshot 8: Process with different user and group id, when casper=0*

When **casper=1**, only processes with same user id should be visible. Testing with 1 and 1703 as process ids:

```

itucs@ubuntu: ~/Desktop/prj2
itucs@ubuntu: ~/Desktop/prj2 80x24
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ sudo ./testing 1 1
casper value 1, pid 1
syscall result:0
Success
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ sudo ./testing 1703 1
casper value 1, pid 1703
syscall result:0
Success
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ ls /proc/1
ls: cannot access /proc/1: No such file or directory
2 x itucs@ubuntu ~/Desktop/prj2 $ ls /proc/1703
attr          cpuset        latency       mountstats    projid
autogroup      cwd           limits        net           root
auxv           environ       loginuid      ns            sched
cgroup         exe           map_files     oom_adj       sched
clear_refs     fd            maps          oom_score     sessio
cmdline        fdinfo        mem           oom_score_adj setgro
comm           gid_map       mountinfo     pagemap       smaps
coredump_filter io            mounts        personality    stack
0 ✓ itucs@ubuntu ~/Desktop/prj2 $

```

*Screenshot 9: Testing for casper=1*

In order to test the condition where **casper=2**, a bash script was written which shows process groups for each process under the **/proc**.

```

## Script to print groups of processes in /proc
#!/bin/bash
PID_LIST=$(ls /proc/ | awk '/[0-9]/ {print}')      ## awk finds numbers in /proc

for PID in $PID_LIST
do
    echo -e "pid: $PID \t $(cat /proc/$PID/status | grep Groups)"
done

```

```

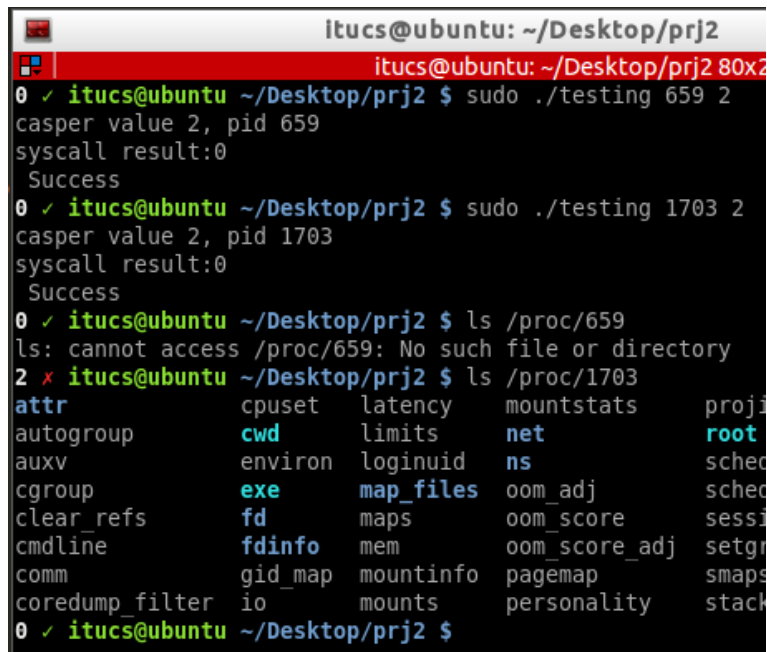
itucs@ubuntu: ~/Desktop
itucs@ubuntu: ~/Desktop
pid: 1541      Groups:      4 24 27 30 46 1
pid: 1545      Groups:      4 24 27 30 46 1
pid: 1549      Groups:      4 24 27 30 46 1
pid: 1556      Groups:      4 24 27 30 46 1
pid: 1584      Groups:      4 24 27 30 46 1
pid: 1590      Groups:      4 24 27 30 46 1
pid: 1592      Groups:      4 24 27 30 46 1
pid: 1593      Groups:      4 24 27 30 46 1
pid: 16        Groups:
pid: 17        Groups:
pid: 1703      Groups:      4 24 27 30 46 1
pid: 18        Groups:
pid: 19        Groups:

```

*Screenshot 10: Output of print\_groups.sh script*



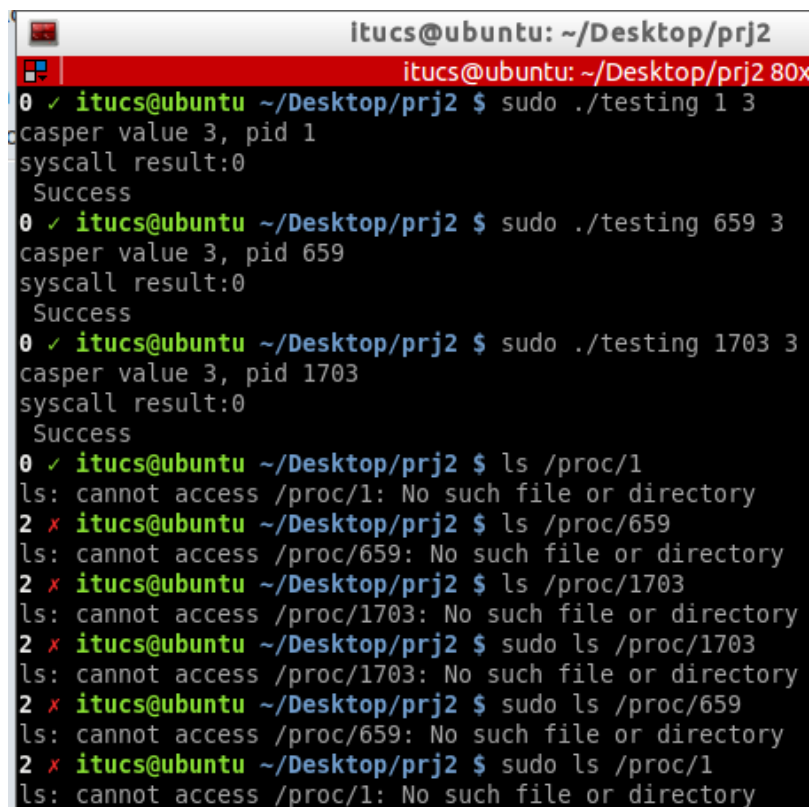
Two processes with different groups were selected to test **casper=2** condition. 659 pid does not have any common group with the bash, but all groups of 1703 pid are same with bash:



```
itucs@ubuntu: ~/Desktop/prj2
itucs@ubuntu: ~/Desktop/prj2 80x20
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ sudo ./testing 659 2
casper value 2, pid 659
syscall result:0
Success
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ sudo ./testing 1703 2
casper value 2, pid 1703
syscall result:0
Success
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ ls /proc/659
ls: cannot access /proc/659: No such file or directory
2 ✗ itucs@ubuntu ~/Desktop/prj2 $ ls /proc/1703
attr          cpuset        latency        mountstats     projid
autogroup      cwd           limits         net            root
auxv           environ       loginuid       ns             sched
cgroup         exe          map_files     oom_adj        sched
clear_refs     fd           maps          oom_score      sessi
cmdline        fdinfo       mem           oom_score_adj  setgr
comm           gid_map      mountinfo     pagemap        smaps
coredump_filter io           mounts        personality     stack
```

*Screenshot 11: Testing for casper=2, with processes 659 and 1703*

Testing for **casper=3**, for different processes with different user and group ids. All these processes are invisible, even with super user privileges.



```
itucs@ubuntu: ~/Desktop/prj2
itucs@ubuntu: ~/Desktop/prj2 80x20
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ sudo ./testing 1 3
casper value 3, pid 1
syscall result:0
Success
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ sudo ./testing 659 3
casper value 3, pid 659
syscall result:0
Success
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ sudo ./testing 1703 3
casper value 3, pid 1703
syscall result:0
Success
0 ✓ itucs@ubuntu ~/Desktop/prj2 $ ls /proc/1
ls: cannot access /proc/1: No such file or directory
2 ✗ itucs@ubuntu ~/Desktop/prj2 $ ls /proc/659
ls: cannot access /proc/659: No such file or directory
2 ✗ itucs@ubuntu ~/Desktop/prj2 $ ls /proc/1703
ls: cannot access /proc/1703: No such file or directory
2 ✗ itucs@ubuntu ~/Desktop/prj2 $ sudo ls /proc/1703
ls: cannot access /proc/1703: No such file or directory
2 ✗ itucs@ubuntu ~/Desktop/prj2 $ sudo ls /proc/659
ls: cannot access /proc/659: No such file or directory
2 ✗ itucs@ubuntu ~/Desktop/prj2 $ sudo ls /proc/1
ls: cannot access /proc/1: No such file or directory
```

*Screenshot 12: Testing for casper=3*