

Message Passing Interface

MPI Basics



Tom Murphy

Director of Contra Costa College
High Performance Computing Center





Preliminaries

answering: “What is a cluster”

- To set-up a cluster we must:
 - Configure the individual computers
 - Establish some form of communication between machines
 - Run the program(s) that exploit the above
- MPI is all about exploitation



MPI Basics
BWUPEP2011, UIUC, May 29 - June 10 2011





So what does MPI do?

“What does the coder do?”

- Simply stated:
 - MPI allows moving data between processes
 - Data that is needed
 - for a computation
 - or
 - from a computation
- Now just wait a second!
 - Shouldn't that be processors!





Simple or Complex?

for now it's simple

- MPI has 100+ very complex library calls
 - 52 Point-to-Point Communication
 - 16 Collective Communication
 - 30 Groups, Contexts, and Communicators
 - 16 Process Topologies
 - 13 Environmental Inquiry
 - 1 Profiling
- MPI only needs 6 very simple complex library calls






Six Basic MPI commands


Via three key fingers

- How do I start and stop




- MPI_Init
- MPI_Finalize

- Know thy self (and others)



- MPI_Comm_rank
- MPI_Comm_size

- Middle Finger - The Message Passing



- MPI_Send
- MPI_Recv





Essential MPI Organization

that sometimes get in the way

- Data Representation is Standardized
 - MPI data types
- Harnessing Processes for a Task
 - MPI Communicators
- Specifying a kind of message
 - MPI Tags
- How many: Processes and Processors
 - -np
 - -machinefile





Data Representation

Exact -> Integer Types

- Signed
 - MPI_CHAR
 - MPI_SHORT
 - MPI_INT
 - MPI_LONG
- Unsigned
 - MPI_UNSIGNED_CHAR
 - MPI_UNSIGNED_SHORT
 - MPI_UNSIGNED
 - MPI_UNSIGNED_LONG





Data Representation

Approximate -> Floating Point

- MPI_FLOAT
- MPI_DOUBLE
- MPI_LONG_DOUBLE



MPI Basics
BWUPEP2011, UIUC, May 29 - June 10 2011





Data Representation

Special Cases

- MPI_BYTE
 - Device independent
 - Exactly 8 bits
- MPI_PACKED
 - Allows non-contiguous data
 - MPI_PACK
 - MPI_UNPACK





Under the hood of the Six

How do I start and stop

- MPI_Init (int *argc, char ***argv)
 - We gotta change (int argc, char **argv)
- since
- MPI uses it to pass data to all machines
- MPI_Finalize ()





Under the hood of the Six

Know thyself (and others)

- MPI_Comm_rank
(MPI_Comm comm, int *rank)
- MPI_Comm_size
(MPI_Comm comm, int *size)





Under the hood of the Six

The actual message passing

- MPI_Send(
void* buf, int count, MPI_Datatype datatype,
int dest, int tag, MPI_Comm comm)
- MPI_Recv(
void* buf, int count, MPI_Datatype datatype,
int source, int tag, MPI_Comm comm,
MPI_Status *status)





MPI Hello World

Lets explore some code

- Fire up a qsub interactive shell on AC
 - `ssh <account>@ac.ncsa.uiuc.edu`
 - `cp ~tra5/mpihello.c .`
 - `qsub -I`
 - `mpdstartup`
 - `mpicc -o mpihello mpihello.c`
 - `mpirun -np 4 ./mpihello`





MPI Hello World

Lets explore some code

- Fire up a qsub interactive shell on sooner
 - ssh **yourusername**@sooner.oscer.ou.edu
 - cp -r ~hneeman/BWUPEP2011_exercises/MPIHello/ ~/BWUPEP2011_exercises/
 - cd ~/BWUPEP2011_exercises/MPIHello/C/
 - **Edit mpihello.bsub : replace e-mail address and username**
 - make_cmd
 - bsub < mpihello.bsub





MPI Hello World

A fugue in six parts

1. Including the Right Stuff
2. General Declarations
3. MPI Setup
4. Client-side Code
5. Server-side Code
6. The Grand Finale



MPI Basics
BWUPEP2011, UIUC, May 29 - June 10 2011





MPI Hello World

Part 1: Including the right stuff

```
#include <mpi.h>
#include <stdio.h>
#include <string.h>
```

```
#define SERVER 0
#define TAG 2
```





MPI Hello World

Part 2: General Declarations

```
int main(int argc, char **argv) {  
  
    int my_rank, world_size;  
    int destination=SERVER, source;  
    int tag=TAG, length;  
    char message[256], name[80];  
    MPI_Status status;
```





MPI Hello World

Part 3: MPI Setup

```
MPI_Init(&argc, &argv);  
  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &world_size);  
MPI_Get_processor_name(name, &length);  
  
sprintf(message, "Process %d (%d of %d) on %s!",  
        my_rank, my_rank+1, world_size, name);
```





MPI Hello World

Part 4: Client-side Code

```
if (my_rank != SERVER) {           // Client section
    fprintf(stderr, "Client: process %d\n", my_rank);
    MPI_Send(message, strlen(message)+1, MPI_CHAR,
              destination, tag, MPI_COMM_WORLD);
}
```





MPI Hello World

Part 5: Server-side Code

```
} else { // Server Section
    fprintf(stderr, "Server: process %d\n", my_rank);

    fprintf(stderr, "%s\n", message);
    for (source = 1; source < world_size ; source++) {
        MPI_Recv(message, 256, MPI_CHAR,
                 source, tag, MPI_COMM_WORLD, &status);
        fprintf(stderr, "%s\n", message);
    }
}
```





MPI Hello World

Part 6: The Grand Finale

```
fprintf(stderr, "Calling Finalize %d\n", my_rank);  
MPI_Finalize();  
}
```

