

Parallel Implementations of Gaussian Elimination

Vasilije Perović

Western Michigan University

vasilije.perovic@wmich.edu

January 27, 2012

Linear systems of equations

General form of a linear system of equations is given by

$$\begin{array}{ccccccc} a_{11}x_1 & + \cdots + & a_{1n}x_n & = & b_1 \\ a_{21}x_1 & + \cdots + & a_{2n}x_n & = & b_2 \\ \vdots & & \vdots & & \\ a_{m1}x_1 & + \cdots + & a_{mn}x_n & = & b_m \end{array}$$

where a_{ij} 's and b_i 's are known and we are solving for x_i 's.

$$A\mathbf{x} = \mathbf{b}$$

More compactly, we can rewrite system of linear equations in the form

$$A\mathbf{x} = \mathbf{b}$$

where

$$A = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_m \end{bmatrix}$$

Why study solutions of $A\mathbf{x} = \mathbf{b}$

- 1 One of the most fundamental problems in the field of scientific computing
- 2 Arises in many applications:

Why study solutions of $Ax = b$

- 1 One of the most fundamental problems in the field of scientific computing
- 2 Arises in many applications:
 - Chemical engineering
 - Interpolation
 - Structural analysis
 - Regression Analysis
 - Numerical ODEs and PDEs

Methods for solving $A\mathbf{x} = \mathbf{b}$

- 1 *Direct methods*
- 2 *Iterative methods*

Methods for solving $Ax = b$

- 1 *Direct methods* – obtain the exact solution (in real arithmetic) in finitely many operations
- 2 *Iterative methods* – generate sequence of approximations that converge in the limit to the solution

Methods for solving $Ax = b$

- ① *Direct methods* – obtain the exact solution (in real arithmetic) in finitely many operations
 - Gaussian elimination (LU factorization)
 - QR factorization
 - WZ factorization
- ② *Iterative methods* – generate sequence of approximations that converge in the limit to the solution
 - Jacobi iteration
 - Gauss-Seidal iteration
 - SOR method (successive over-relaxation)

Methods for solving $Ax = b$

- ① *Direct methods* – obtain the exact solution (in real arithmetic) in finitely many operations
 - Gaussian elimination (LU factorization)
 - QR factorization
 - WZ factorization
- ② *Iterative methods* – generate sequence of approximations that converge in the limit to the solution
 - Jacobi iteration
 - Gauss-Seidal iteration
 - SOR method (successive over-relaxation)

Gaussian Elimination

When solving $A\mathbf{x} = \mathbf{b}$ we will assume throughout this presentation that A is non-singular and A and \mathbf{b} are known

$$\begin{array}{rclcl} a_{11}x_1 & + \cdots + & a_{1n}x_n & = & b_1 \\ a_{21}x_1 & + \cdots + & a_{2n}x_n & = & b_2 \\ \vdots & & \vdots & & \\ a_{n1}x_1 & + \cdots + & a_{nn}x_n & = & b_n . \end{array}$$

Gaussian Elimination

Assuming that $a_{11} \neq 0$ we first subtract a_{21}/a_{11} times the first equation from the second equation to eliminate the coefficient x_1 in the second equation, and so on until the coefficients of x_1 in the last $n - 1$ rows have all been eliminated. This gives the modified system of equations

Gaussian Elimination

Assuming that $a_{11} \neq 0$ we first subtract a_{21}/a_{11} times the first equation from the second equation to eliminate the coefficient x_1 in the second equation, and so on until the coefficients of x_1 in the last $n - 1$ rows have all been eliminated. This gives the modified system of equations

$$\begin{array}{rclcl} a_{11}x_1 + & a_{12}x_2 & + \cdots + & a_{1n}x_n & = b_1 \\ & a_{22}^{(1)}x_2 & + \cdots + & a_{2n}^{(1)}x_n & = b_2^{(1)} \\ & \vdots & & \vdots & \vdots \\ & a_{n1}^{(1)}x_2 & + \cdots + & a_{nn}^{(1)}x_n & = b_n^{(1)} \end{array},$$

where

$$a_{ij}^{(1)} = a_{ij} - a_{1j} \frac{a_{i1}}{a_{11}} \quad b_i^{(1)} = b_i - b_1 \frac{a_{i1}}{a_{11}} \quad i, j = 1, 2, \dots, n.$$

Gaussian Elimination (*Forward Reduction*)

Applying the same process the last $n - 1$ equations of the modified system to eliminate coefficients of x_2 in the last $n - 2$ equations, and so on, until the entire system has been reduced to the (*upper*) *triangular form*

$$\begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ & a_{22}^{(1)} & \cdots & a_{2n}^{(1)} \\ & & \ddots & \vdots \\ & & & a_{nn}^{(n-1)} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \cdots \\ x_n \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2^{(1)} \\ \vdots \\ b_n^{(n-1)} \end{bmatrix}.$$

Gaussian Elimination (*Forward Reduction*)

Applying the same process the last $n - 1$ equations of the modified system to eliminate coefficients of x_2 in the last $n - 2$ equations, and so on, until the entire system has been reduced to the (*upper*) *triangular form*

$$\begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ & a_{22}^{(1)} & \cdots & a_{2n}^{(1)} \\ & & \ddots & \vdots \\ & & & a_{nn}^{(n-1)} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \cdots \\ x_n \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2^{(1)} \\ \vdots \\ b_n^{(n-1)} \end{bmatrix}.$$

- The superscripts indicate the number of times the elements had to be changed.

Gaussian Elimination – Example

- Perform the *forward reduction* the the system given below

$$\begin{bmatrix} 4 & -9 & 2 \\ 2 & -4 & 4 \\ -1 & 2 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 2 \\ 3 \\ 1 \end{bmatrix}$$

Gaussian Elimination – Example

- Perform the *forward reduction* the the system given below

$$\begin{bmatrix} 4 & -9 & 2 \\ 2 & -4 & 4 \\ -1 & 2 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 2 \\ 3 \\ 1 \end{bmatrix}$$

We start by writing down the augmented matrix for the given system:

$$\left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 2 & -4 & 4 & 3 \\ -1 & 2 & 2 & 1 \end{array} \right]$$

Gaussian Elimination – Example

We perform appropriate row operations to obtain:

$$\left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 2 & -4 & 4 & 3 \\ -1 & 2 & 2 & 1 \end{array} \right] \longrightarrow$$

Gaussian Elimination – Example

We perform appropriate row operations to obtain:

$$\left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 2 & -4 & 4 & 3 \\ -1 & 2 & 2 & 1 \end{array} \right] \xrightarrow[\begin{array}{l} R_2 - (\frac{2}{4})R_1 \rightarrow R_2 \\ R_3 - (\frac{-1}{4})R_1 \rightarrow R_3 \end{array}]{\hspace{1cm}} \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 0 & 0.5 & 3 & 2 \\ 0 & -0.25 & 2.5 & 1.5 \end{array} \right]$$

Gaussian Elimination – Example

We perform appropriate row operations to obtain:

$$\begin{aligned}
 \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 2 & -4 & 4 & 3 \\ -1 & 2 & 2 & 1 \end{array} \right] & \xrightarrow{\begin{array}{l} R_2 - \left(\frac{2}{4}\right)R_1 \rightarrow R_2 \\ R_3 - \left(\frac{-1}{4}\right)R_1 \rightarrow R_3 \end{array}} \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 0 & 0.5 & 3 & 2 \\ 0 & -0.25 & 2.5 & 1.5 \end{array} \right] \\
 & \xrightarrow{R_3 - \left(\frac{-0.25}{0.5}\right)R_2 \rightarrow R_3} \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 0 & 0.5 & 3 & 2 \\ 0 & 0 & 4 & 2.5 \end{array} \right]
 \end{aligned}$$

Gaussian Elimination – Example

Note that the row operations used to eliminate x_1 from the second and the third equations are equivalent to multiplying *on the left* the augmented matrix:

$$\begin{bmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0.25 & 0 & 1 \end{bmatrix} \cdot \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 2 & -4 & 4 & 3 \\ -1 & 2 & 2 & 1 \end{array} \right] = \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 0 & 0.5 & 3 & 2 \\ 0 & -0.25 & 2.5 & 1.5 \end{array} \right]$$

Gaussian Elimination – Example

Note that the row operations used to eliminate x_1 from the second and the third equations are equivalent to multiplying *on the left* the augmented matrix:

$$\underbrace{\begin{bmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0.25 & 0 & 1 \end{bmatrix}}_{L_1} \cdot \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 2 & -4 & 4 & 3 \\ -1 & 2 & 2 & 1 \end{array} \right] = \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 0 & 0.5 & 3 & 2 \\ 0 & -0.25 & 2.5 & 1.5 \end{array} \right]$$

Gaussian Elimination – Example

Note that the row operations used to eliminate x_1 from the second and the third equations are equivalent to multiplying *on the left* the augmented matrix:

$$\underbrace{\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0.5 & 1 \end{bmatrix}}_{L_2} \cdot \underbrace{\begin{bmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0.25 & 0 & 1 \end{bmatrix}}_{L_1} \cdot [A|\mathbf{b}] = \left[\begin{array}{ccc|c} 4 & -9 & 2 & 2 \\ 0 & 0.5 & 3 & 2 \\ 0 & 0 & 4 & 2.5 \end{array} \right]$$

Gaussian Elimination – Example

$$\underbrace{\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0.5 & 1 \end{bmatrix}}_{L_2} \cdot \underbrace{\begin{bmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0.25 & 0 & 1 \end{bmatrix}}_{L_1} \cdot [A|\mathbf{b}] = \begin{bmatrix} 4 & -9 & 2 & | & 2 \\ 0 & 0.5 & 3 & | & 2 \\ 0 & 0 & 4 & | & 2.5 \end{bmatrix}$$

$\hat{L} = L_2 \cdot L_1$

Thus, we can write

$$(L_2 \cdot L_1) \cdot A = \hat{L} \cdot A = U$$

LU factorization

In general, we can think of row operations as multiplying with matrices on the left, thus the i^{th} elimination step is equivalent as multiplication on the left by

$$L_i = \begin{bmatrix} 1 & & & & \\ & \ddots & & & \\ & & 1 & & \\ & & -l_{i+1,i} & & \\ & & \vdots & \ddots & \\ & & -l_{n,i} & & 1 \end{bmatrix} \cdot$$

LU factorization

Continuing in this fashion we obtain

$$\hat{L}A\mathbf{x} = \hat{L}\mathbf{b}, \quad \hat{L} = L_{n-1} \cdots L_2 L_1,$$

LU factorization

Continuing in this fashion we obtain

$$\hat{L}A\mathbf{x} = \hat{L}\mathbf{b}, \quad \hat{L} = L_{n-1} \cdots L_2 L_1,$$

$$A = \hat{L}^{-1}U = LU.$$

LU factorization

Continuing in this fashion we obtain

$$A = \hat{L}^{-1}U = LU.$$

Thus, the Gaussian elimination algorithm for solving $A\mathbf{x} = \mathbf{b}$ is mathematically equivalent to the three-step process:

- | | |
|---------------------------------|-----------------------------|
| 1. Factor | $A = LU$ |
| 2. Solve (forward substitution) | $L\mathbf{y} = \mathbf{b}$ |
| 3. Solve (back substitution) | $U\mathbf{x} = \mathbf{y}.$ |

LU factorization

Continuing in this fashion we obtain

$$A = \hat{L}^{-1}U = LU.$$

Thus, the Gaussian elimination algorithm for solving $A\mathbf{x} = \mathbf{b}$ is mathematically equivalent to the three-step process:

- | | |
|---------------------------------|-----------------------------|
| 1. Factor | $A = LU$ |
| 2. Solve (forward substitution) | $L\mathbf{y} = \mathbf{b}$ |
| 3. Solve (back substitution) | $U\mathbf{x} = \mathbf{y}.$ |

- Order of operations:

- $\frac{n^3}{3} + n^2 - \frac{n}{3}$ multiplications/divisions
- $\frac{n^3}{3} + \frac{n^2}{2} - \frac{5n}{6}.$

Need for partial pivoting

Apply LU factorization *without pivoting* to

$$A = \begin{bmatrix} 0.0001 & 1 \\ 1 & 1 \end{bmatrix}$$

in *three-decimal-digit* floating point arithmetic.

Need for partial pivoting

$$A = \begin{bmatrix} 0.0001 & 1 \\ 1 & 1 \end{bmatrix}$$

Solution: L and U are easily obtainable and are given by:

$$L = \begin{bmatrix} 1 & 0 \\ fl(1/10^{-4}) & 1 \end{bmatrix}, \quad fl(1/10^{-4}) \text{ rounds to } 10^4,$$

$$U = \begin{bmatrix} 10^{-4} & 1 \\ 0 & fl(1 - 10^4 \cdot 1) \end{bmatrix}, \quad fl(1 - 10^4 \cdot 1) \text{ rounds to } -10^4$$

$$\text{so } LU = \begin{bmatrix} 1 & 0 \\ 10^4 & 1 \end{bmatrix} \begin{bmatrix} 10^{-4} & 1 \\ 0 & -10^4 \end{bmatrix} = \begin{bmatrix} 10^{-4} & 1 \\ 1 & 0 \end{bmatrix}$$

$$\text{but } A = \begin{bmatrix} 10^{-4} & 1 \\ 1 & 1 \end{bmatrix}.$$

Need for partial pivoting

$$A = \begin{bmatrix} 10^{-4} & 1 \\ 1 & 1 \end{bmatrix} \quad \text{but} \quad LU = \begin{bmatrix} 10^{-4} & 1 \\ 1 & 0 \end{bmatrix}$$

Remark 1: Note that original a_{22} has been entirely “lost” from the computation by subtracting 10^4 from it. Thus if we were to use this LU factorization in order to solve a system there would be no way to guarantee an accurate answer. This is called *numerical instability*.

Need for partial pivoting

$$A = \begin{bmatrix} 10^{-4} & 1 \\ 1 & 1 \end{bmatrix} \quad \text{but} \quad LU = \begin{bmatrix} 10^{-4} & 1 \\ 1 & 0 \end{bmatrix}$$

Remark 1: Note that original a_{22} has been entirely “lost” from the computation by subtracting 10^4 from it. Thus if we were to use this LU factorization in order to solve a system there would be no way to guarantee an accurate answer. This is called *numerical instability*.


Remark 2: Suppose we attempted to solve system $A\mathbf{x} = [1, 2]^T$ for \mathbf{x} using this LU decomposition. The correct answer is $\mathbf{x} \approx [1, 1]^T$. Instead, if we were to stick with the three-digit-floating point arithmetic we would get an answer $\hat{\mathbf{x}} = [0, 1]^T$ which is completely erroneous.

Sequential Algorithm

```

for (k = 0; k < n-1; k++) { /* Forward elimination */
    r = max_col(a,k);
    if (k != r) exchange_row(a,b,r,k);
    for (i=k+1; i < n; i++) {
        l[i] = a[i][k]/a[k][k];
        for (j=k+1; j < n; j++)
            a[i][j] = a[i][j] - l[i] * a[k][j];
        b[i] = b[i] - l[i] * b[k];
    }
}
for (k = n-1; k >= 0; k--) { /* Backward substitution */
    sum = 0.0;
    for (j=k+1; j < n; j++)
        sum = sum + a[k][j] * x[j];
    x[k] = 1/a[k][k] * (b[k] - sum);
}
return x;
    
```

Row-Oriented Algorithm



- 1 Determination of the local pivot element,
- 2 Determination of the global pivot element,
- 3 Exchange of the pivot row,
- 4 Distribution of the pivot row,
- 5 Computation of the elimination factors,
- 6 Computation of the matrix elements.

Row-Oriented Algorithm

Remark 1: The computation of the solution vector \mathbf{x} in the backward substitution is inherently serial, since the values of x_k depend on each other and are computed one after another. Thus, in step k the processor P_q owning row k computes the value of x_k and sends the value to all other processors by a *single* broadcast operation.


Row-Oriented Algorithm

Remark 1: The computation of the solution vector \mathbf{x} in the backward substitution is inherently serial, since the values of x_k depend on each other and are computed one after another. Thus, in step k the processor P_q owning row k computes the value of x_k and sends the value to all other processors by a *single* broadcast operation.

Remark 2: Note that the data distribution is not quite adequate. For example, suppose we are at the i^{th} step and that $i > m \cdot n/p$ where m is a natural number. In that case, processors p_0, p_1, \dots, p_{m-1} are idle since all their work is done until the back substitution step at the very end. Hence there is an issue with load balancing.

Row-Oriented Algorithm

Remark 2: Note that the data distribution is not quite adequate. For example, suppose we are at the i^{th} step and that $i > m \cdot n/p$ where m is a natural number. In that case, processors p_0, p_1, \dots, p_{m-1} are idle since all their work is done until the back substitution step at the very end. Hence there is an issue with load balancing.



block-cyclic row distribution

Row-Oriented Algorithm

Remark 3: We could also consider *column-oriented* data distribution. In that case, at the k^{th} step the processor that owns k^{th} column would have all needed data to compute the new pivot. On the one hand, this would reduce communication between processors that we had when considering row-oriented data distribution. On the other hand, with column orientation pivot determination is all done serially. Thus, in case when $n \gg p$ (which is often case) row oriented data distribution might be more advantageous since the pivot determination is done in parallel

Pipelined Implementation

- 1 Challenges and possible approaches
 - *Send ahead* strategy

Pipelined Implementation

- 1 Challenges and possible approaches
 - *Send ahead* strategy
 - Challenges with pivoting

Pipelined Implementation

- 1 Challenges and possible approaches
 - *Send ahead* strategy
 - Challenges with pivoting
 - *Column pivoting*

Pipelined Implementation

- ① Challenges and possible approaches
 - *Send ahead* strategy
 - Challenges with pivoting
 - *Column pivoting*
- ② Advantages
 - It allows *asynchronous execution*: “processes can reduce their portions of the augmented matrix as soon as the pivot rows are available”

Pipelined Implementation

① Challenges and possible approaches

- *Send ahead* strategy
- Challenges with pivoting
- *Column pivoting*

② Advantages

- It allows *asynchronous execution*: “processes can reduce their portions of the augmented matrix as soon as the pivot rows are available”
- It allows processes to overlap communication and computation times.

Final Remarks

- Compute *communication* and *computation* times.

Final Remarks

- Compute *communication* and *computation* times.
 - Largely depends on what kind of distributed system used (usual configurations that have been studying are 1D and 2D meshes as well as hypercube).

Final Remarks

- Compute *communication* and *computation* times.
 - Largely depends on what kind of distributed system used (usual configurations that have been studying are 1D and 2D meshes as well as hypercube).
- Compute *isoefficiency functions* and determine scalability.

Final Remarks

- Compute *communication* and *computation* times.
 - Largely depends on what kind of distributed system used (usual configurations that have been studying are 1D and 2D meshes as well as hypercube).
- Compute *isoefficiency functions* and determine scalability.
 - Row (column) oriented algorithms, stable but not scalable.
 - Pipelined implementation is perfectly scalable.
- Structured systems
 - Symmetric/Hermitian.
 - Banded and upper (lower) triangular.
 - Positive definite (CG method).

Final Remarks

- Compute *communication* and *computation* times.
 - Largely depends on what kind of distributed system used (usual configurations that have been studying are 1D and 2D meshes as well as hypercube).
- Compute *isoefficiency functions* and determine scalability.
 - Row (column) oriented algorithms, stable but not scalable.
 - Pipelined implementation is perfectly scalable.
- Structured systems
 - Symmetric/Hermitian.
 - Banded and upper (lower) triangular.
 - Positive definite (CG method).
- *Sparse* systems and *iterative methods* (numerical PDEs)