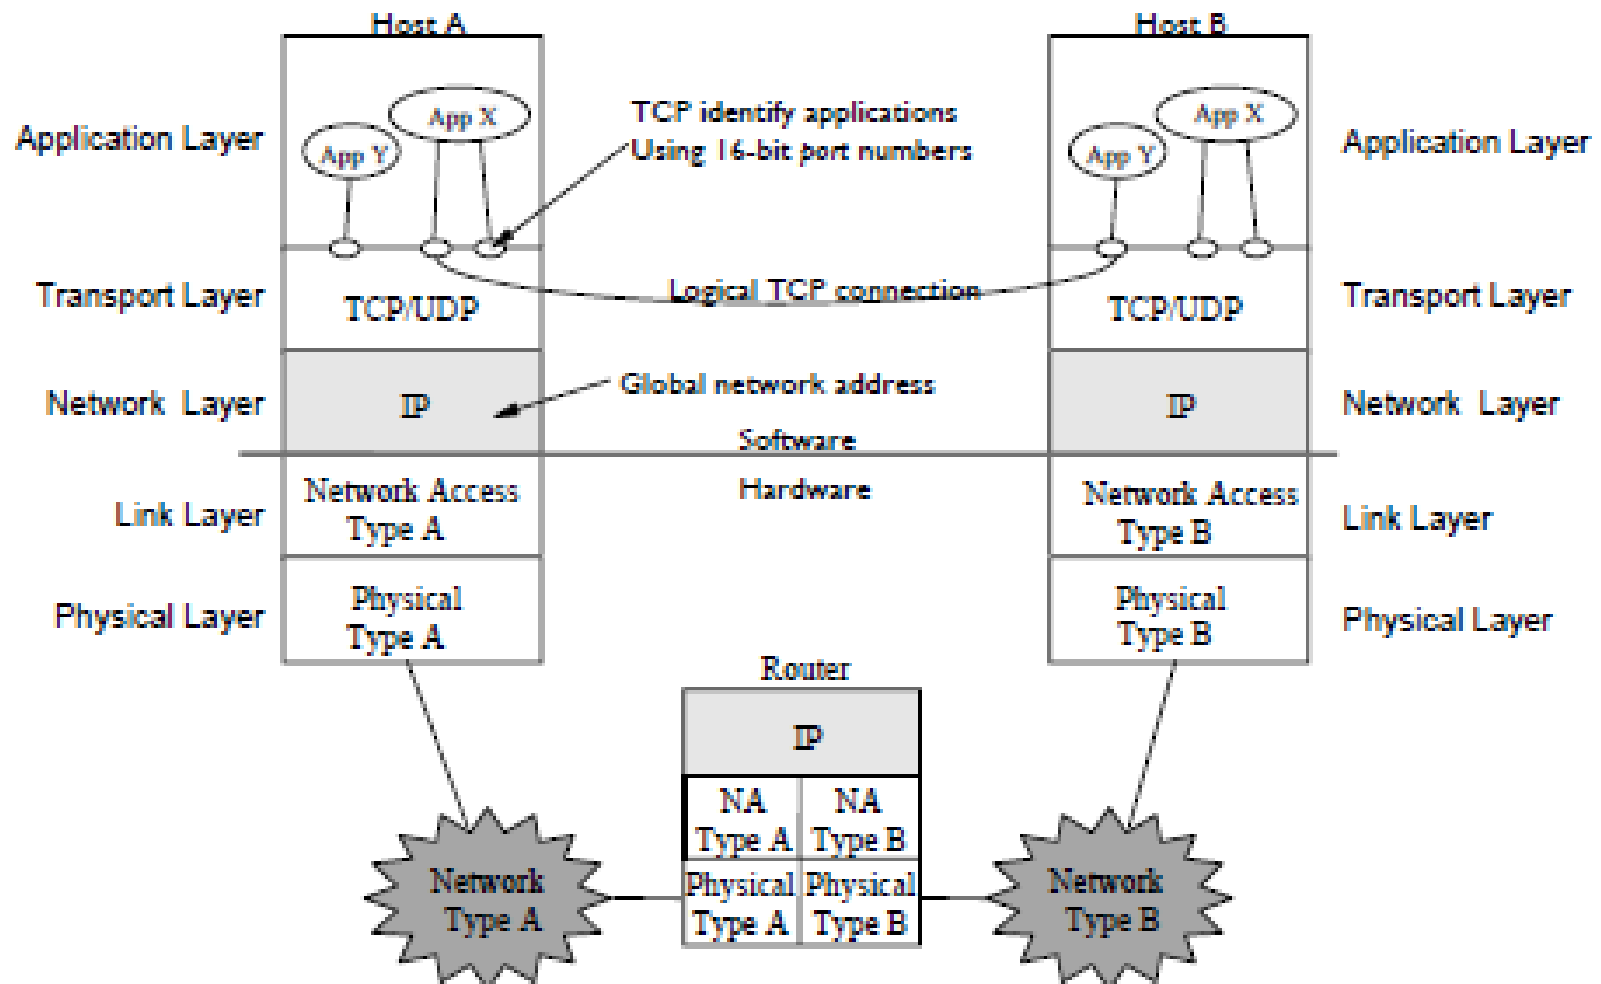
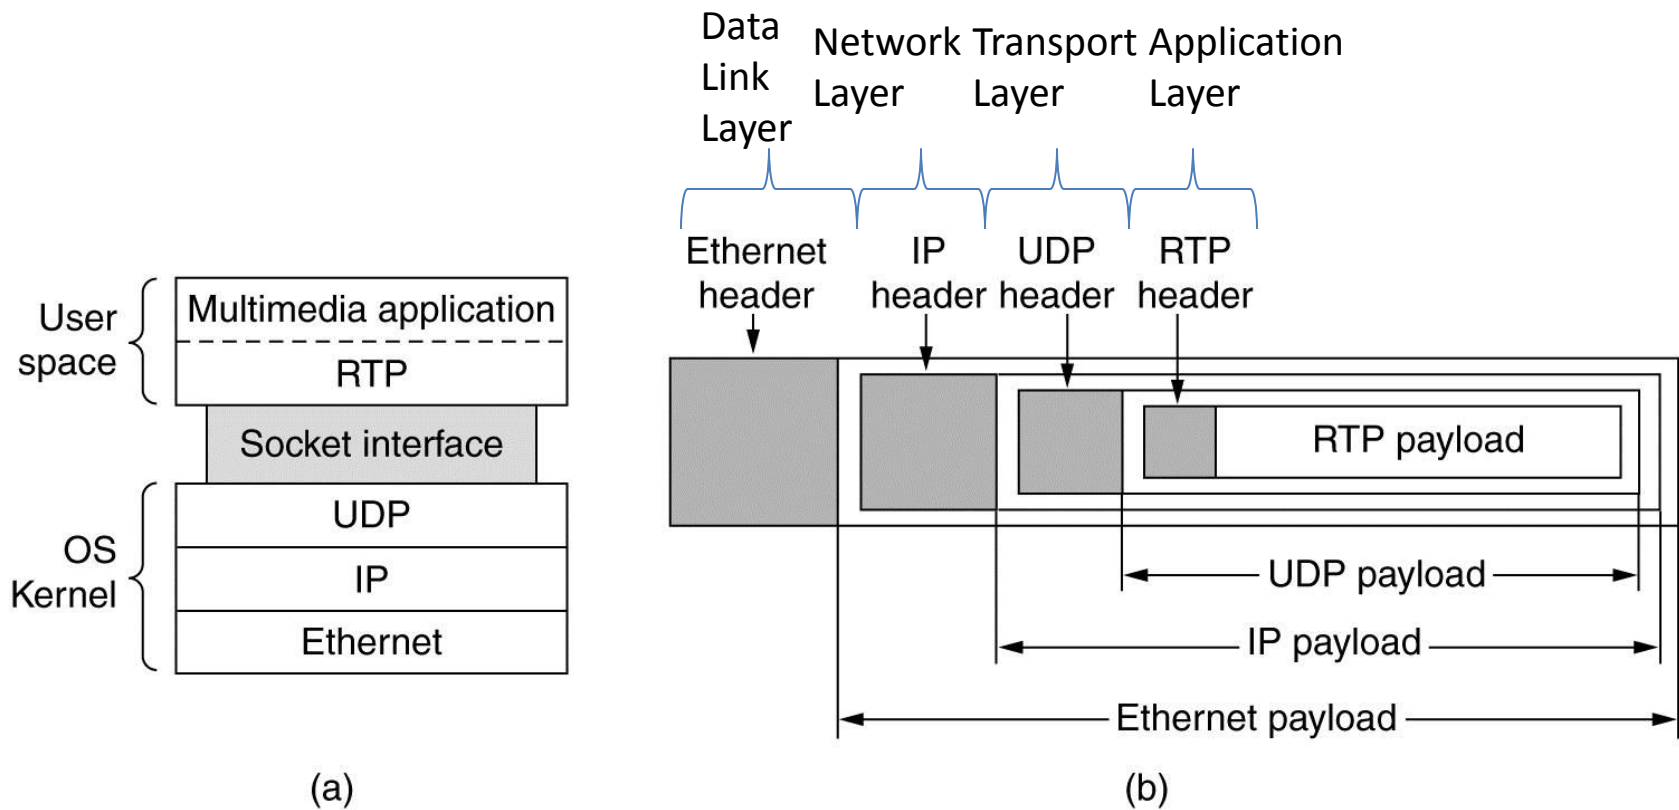


Multimedia Networking

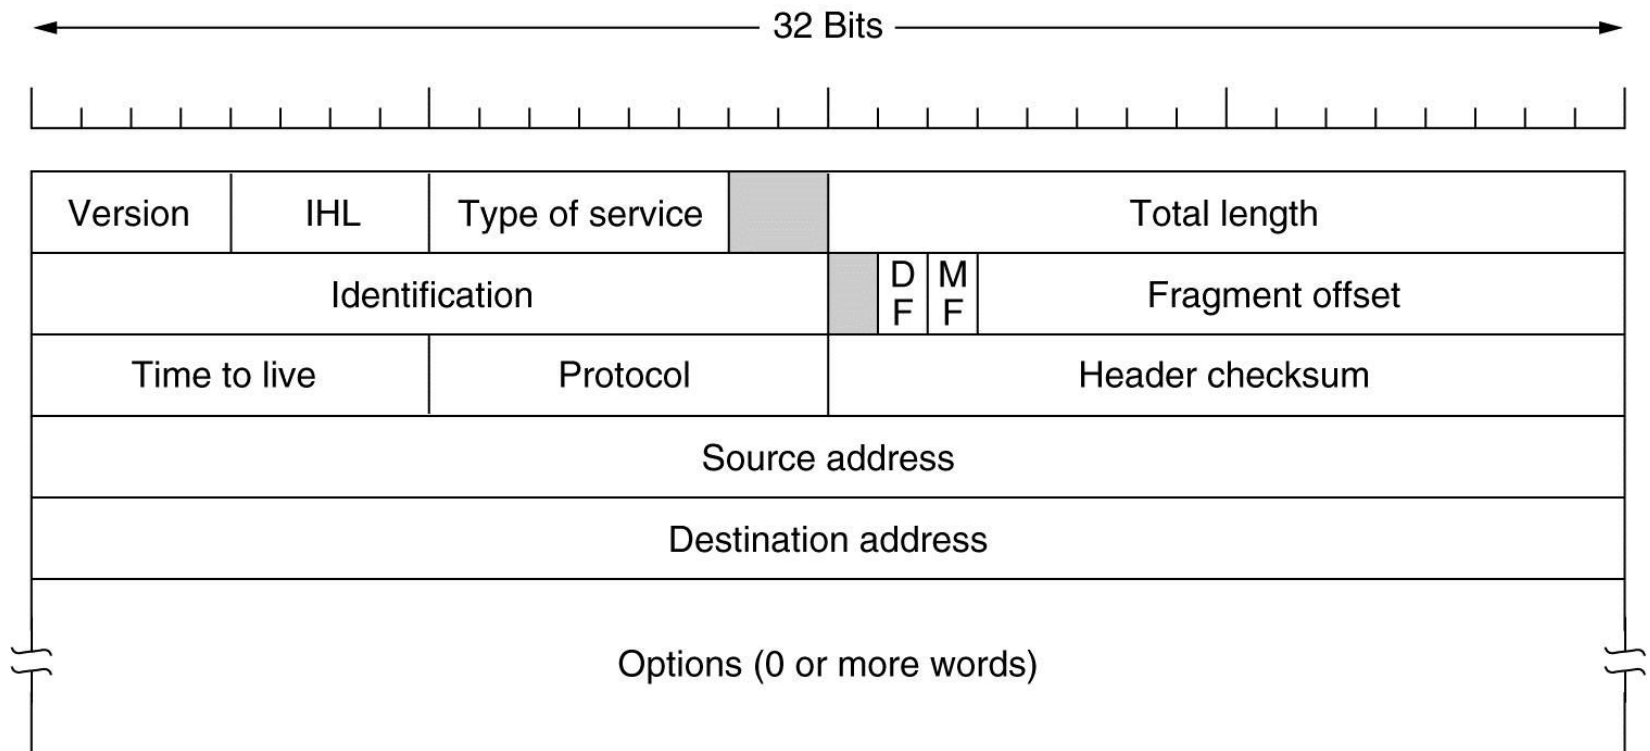
Internet Protocol Architecture



The embedding of protocol data units for multimedia transmission




The IP Protocol



UDP


- UDP is connectionless (no flow control, no error control, no retransmission req.
 - Just IP packet with source and destination ports specified
 - Demultiplexing multiple processes using ports
 - Client-server request reply.

The UDP header.



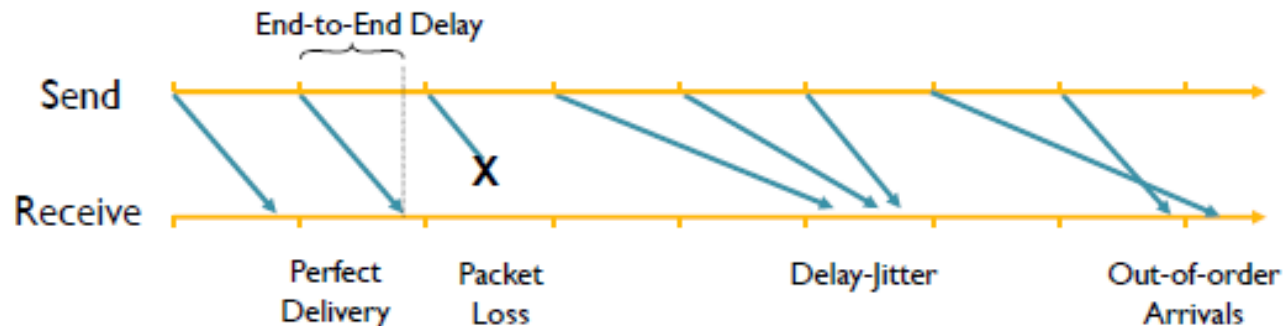
Remote Procedure Call

- Remote procedure call is just like a functional call, but execution takes place on a remote machine.
- Client is talking to server procedure through client stub which is local.
- Pointer passing is difficult. No support for global variables.



Network issues important for multimedia transmission

- Real-time interactive multimedia applications, e.g., Internet phone and real-time video conferencing are very sensitive to
 - Delay
 - Loss
 - Delay-Jitter
 - Out-of-Order Arrivals



Problems with internet

- Latency: Transmission processing delay + Queuing delays (routers) + Propagation delays + End system processing delay.
 - Voice comm.:
 - Roundtrip delay >50msec results in echo
 - One-way delay >250msec results in talker overlap
- Delay Jitters:
 - Variance of frame/packet delays, measure of smoothness of playback
 - Due to random queuing delays in the routers.
 - Might reduce jitter by buffering at the destination, but this increases average delay
- Packet Loss: Packet drop due to overflowed queues in the routers.
 - Also due to TTL expiring (packet not reaching destination for long)
- Out-of-Order Arrivals: No TCP
 - Packets taking different paths.
 - Random queue delays in the routers on different paths.
- Sync skew
 - Measure of multimedia data synchronization, i.e. audio and video
 - Lip synchronization is typically limited to 80msec (voice before video is less)

Current solutions


- Unlike TCP, UDP does not worry about congestion
 - No acknowledgements of correctly received frames=>don't have to wait for acks
 - No transmitter window for frames that increases slowly in size with correct transmissions
 - In TCP slow start
 - Send 1 byte frame receive +ve ack
 - Send 2 byte frame receive +ve ack
 - Send 4 byte frame receive +ve ack, etc.

Current solutions

- Remedies for jitter
 - Content may be buffered at client (increases delay)
 - Data may be prefetched or playback stopped (Interruptions make users unhappy!)
- Remedies for synchronization
 - Timestamp packets at the sender (so receiver knows when the packets should be played back) for each medium
- Remedies for changes in currently available bandwidth
 - Compression ratio (bit rate) dynamically adapted


Internet Phone Example

- Application needs seq. no. and timestamp
 - RTP could be used



Reducing jitter


- $p = \text{packaging period of the media at the sender}$



- Playing the samples as soon as they arrive
 - ensures minimal end-to-end latency
 - has large gap variation


Reducing jitter

- Enqueuing the first sample for some duration better ensures continuous playout.
 - But at the cost of higher playout latency throughout the conference.



Fixed playout delay


- When should it playout?
 $p - r$ or $p' - r$?
- Playout delay should not exceed some threshold, e.g.,
 - 250 ms for audio
 - 10-15 frames/s for video (66.6 ms - 100 ms)
- Tradeoff between *playout delay* vs. *packet loss*.
- For interactive services, such as internet phone, long delays can be bothersome if not intolerable



Adaptive playout delay

- Typically, end-to-end delay distribution of packets within a talk spurt
 - is not known.
 - can change over time.
- Estimate the *network delay* and *its variance*, and adjust the playout delay accordingly.
- Determine play out delay on a per-talkspurt basis.

Adaptive playout delay



- s_i = timestamp of the i th packet generated by the sender
- r_i = time packet i received by the receiver
- p_i = time packet i is played at the receiver
- d_i = average network delay
- v_i = average deviation of network delay

Adaptive playout delay

- End-to-end delay for the i th packet will be $r_i - s_i$. The average network delay of the i th packet, d_i , is estimated by a running average as

$$d_i = (1 - u)d_{i-1} + u(r_i - s_i)$$

where u is fixed constant (e.g., $u = 0.01$)

- The average deviation of delay is

$$v_i = (1 - u)v_{i-1} + u|r_i - s_i - d_i|$$

- Playout time allows for exceeding the deviation by Kv_i

$$p_i = s_i + d_i + Kv_i$$

where K is called a congestion estimator

Adaptive playout delay

- When should playout delay be changed?
 - For voice, wait for silent periods in between talkspurts
- No packet loss case:
 - receiver can determine that a packet is the first packet of the talk spurt.
 - e.g., $s_i - s_{i-1} > 20\text{ ms}$, then receiver knows ith packet start a new talkspurt
- Packet loss case:
 - then easily, within a talkspurt, two successive packets received may have timestamp that differ by more than 20 ms.
 - Use sequence numbers.

Continuous audio transmission

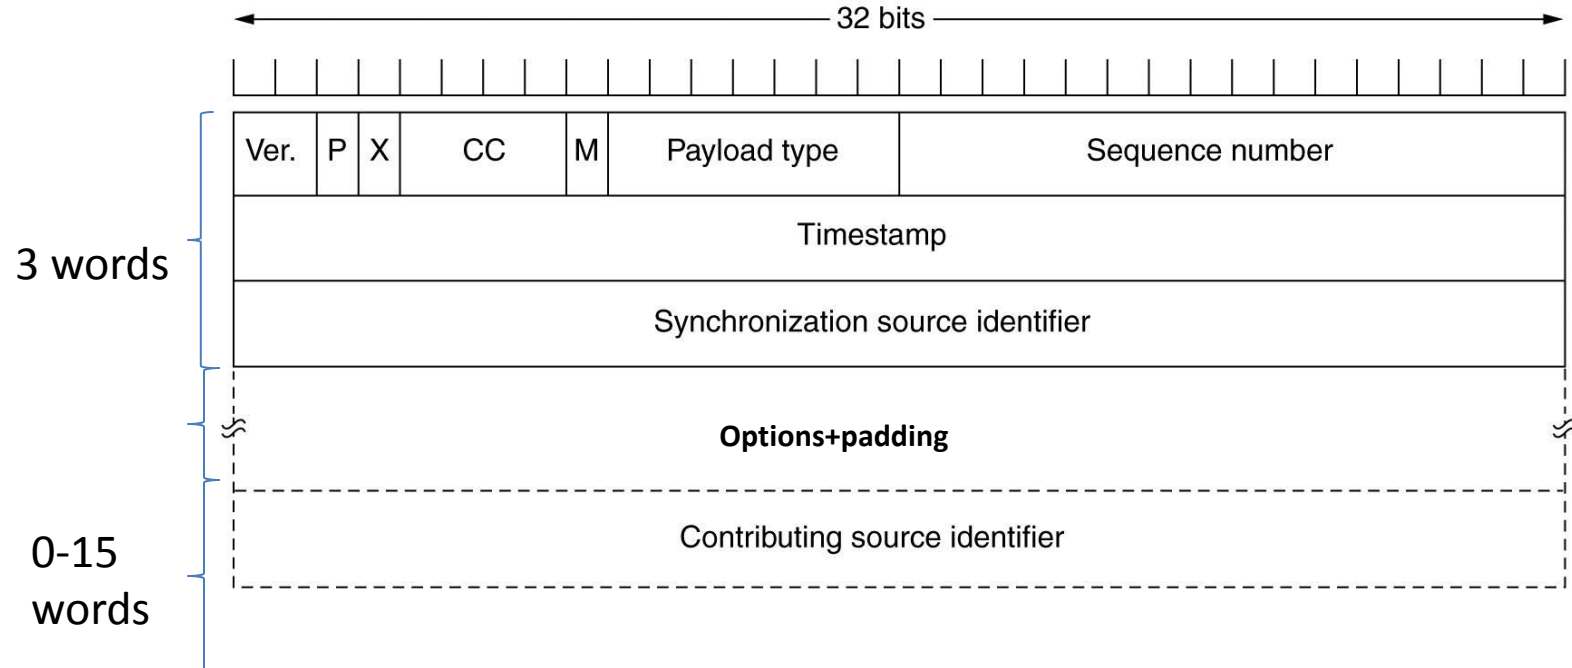
- Scheduling playout based on network delay estimates can give poor results
 - Need *E-to-E measurement* of delay-jitter .
 - Need synchronized clocks at receiver and transmitter
- Continuous multimedia: No silent periods (talkspurts) for many forms of multimedia
 - Music
 - Mixed audio streams
 - Video

Real Time Transport Protocol

- RTP is a transport protocol implemented in the application layer.
 - Runs on top of UDP (in the payload of UDP segments)
- RTP multiplexes several real time multimedia data streams onto a single stream of UDP packets.
- No flow control, error control, retransmissions, acks.
- However, timestamps, sequence numbers and info on how sources are mixed, etc.

RTP Header

- Time stamps: Enable synchronization between multiple streams and reduce jitter.
- Payload type: Encoding algorithm
- P: Padding applied to multiple of 4 bytes. X: Extension header present. CC: # contributing sources. M: application specific bit.
- Synchronization source id: which stream the packet belongs to.
- Contributing source id: the streams being mixed.



Payload Type

- Indicates the type of multimedia data carried by the packet.
 - Audio: PCM μ -law, GSM, MPEG Audio, etc.
 - Video: Motion JPEG, H.261, MPEG-1, MPEG-2, MPEG-4 H.264 AVC etc.
- Defines sampling rate of timestamp.
 - Audio
 - 8 kHz sampling rate (every 125 μ s)
 - 20 ms minimum inter-frame time.
 - Marker (M) bit set on the first packet after a silence period in which no packets are sent.

Sequence no.

- Allows the receiver to detect missing and out of-ordered packets.
- 16 bits long.
- Sender increments the sequence number by one for each transmitted packet.
- RTP does not do anything when a packet is detected to be lost
 - no negative acks, timeouts, retransmissions as in a sliding window protocol
 - instead, the application decides what should be done.

Timestamp

- Enables the receiver to playback samples at appropriate intervals and to enable different media streams to be synchronized.
- 32 bits long.
- A counter of “ticks”, where time between ticks depends on media format.
- Example:
 - Audio application that samples data every $125\ \mu\text{s}$ use this as its clock resolution.
 - Multiple audio samples can be transmitted together.
 - If packet is generated every 10 ms (each containing $10\ \text{ms}/125\ \mu\text{s} = 80$ samples), timestamp between successive packets would have time stamp differ by 80.

Timestamp vs. Sequence No

- Timestamp relates packet to real-time.
 - Timestamp values sampled from a media specific clock.
- Sequence number relates packet to other packets.
- Many packets can have the same timestamp but each should have different sequence number.

MPEG Example

- Out-of-order transmission.
 - Sequence numbers increase monotonically.
 - Timestamps reflect reference relationships.
- Large frames.
 - One video frame likely to be split into parts and packed into multiple RTP packets.
 - Timestamps associate packets together as part of the same frame, while seq. no distinguish packets from each other

Audio Silence example


- Consider audio data type.
- How can the receiver distinguish between packet loss and silence using the timestamp/seq. no mechanism ?
- A big jump in timestamp occurs after receiving no packets for a while.
 - Lost packet detected: next packet received will not have the awaited next seq. no.
 - Silent period detected: next packet received have the awaited next seq. no.

CSRC

- SSRC value for a contributor.
- Used by a mixer to identify the contributing sources of media.
- Mixers
 - Helps in resynchronizing the incoming audio packets and forward it to the low speed link.
 - Also combines several flows in a single new one.
 - Appears as a new source.

Mixer Example

Video and Audio Conference



SSRC

- 32-bit number that uniquely identifies the source of the RTP stream.
 - Not an IP or TCP address.
 - All packets with the same SSRC go together.
- Each sender picks a random SSRC.
- Purpose:
 - Independence from lower-layer protocols.
 - A single node with multiple input sources (e.g., several cameras) can distinguish those sources

Real Time Transport Control Protocol

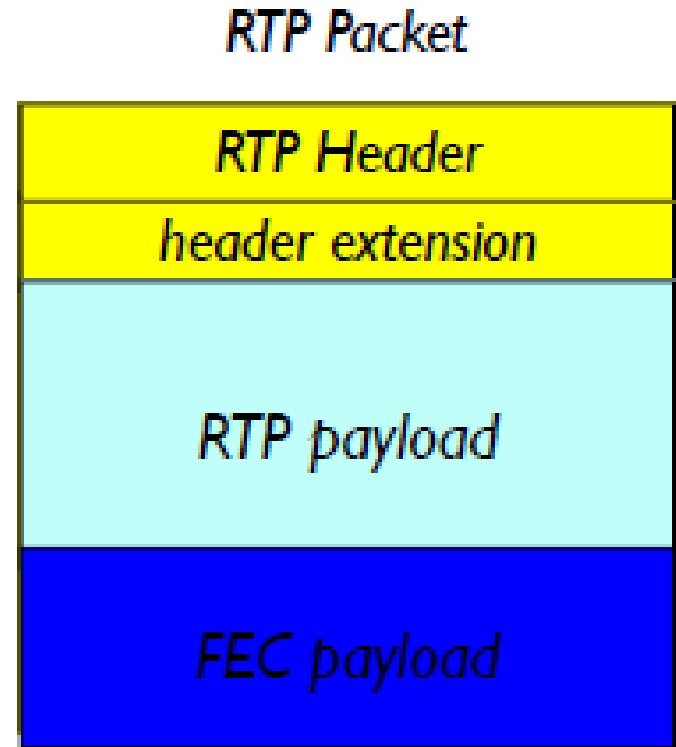
- Hand in hand with RTP.
- Provide feedback on network conditions such as delay, jitter, bandwidth, congestion, packet loss rates.
 - Adaptive encoding.
- Interstream synchronization
 - Different clocks with different granularities, drift rates.

Dealing with Packet Loss

- Packet is considered lost if it never arrives or arrives after its scheduled playout time.
- Sending packets over TCP not viable
 - Retransmission => unacceptable end-to-end delay.
 - Congestion control => slow start: transmission rate is reduced
- Packet loss rate of 1% - 20% is manageable (depending how voice is encoded and transmitted).
 - Forward Error Correction and Interleaving can help conceal loss packets.


Forward Error Correction

- Redundancy is introduced into the stream to enable the receiver to recover from errors due to loss.
- Forms of redundancy
 - Simple replication and retransmission of original data.
 - *k-way XOR*
 - Replication, recoding, and re-transmission of original data.



FEC Basic Concept

- Problem: If a sample is lost, how do we ensure that the redundant information necessary for the repair arrives?
 - In time.
 - Bandwidth of FEC?
 - Location of FEC in the stream?



FEC Method 1: n-way XOR

- Send a redundant encoded chunk after every n chunks.
 - Transmit the word-by-word XOR of groups of n chunks.
- A lost packet can be fully recovered at the receiver

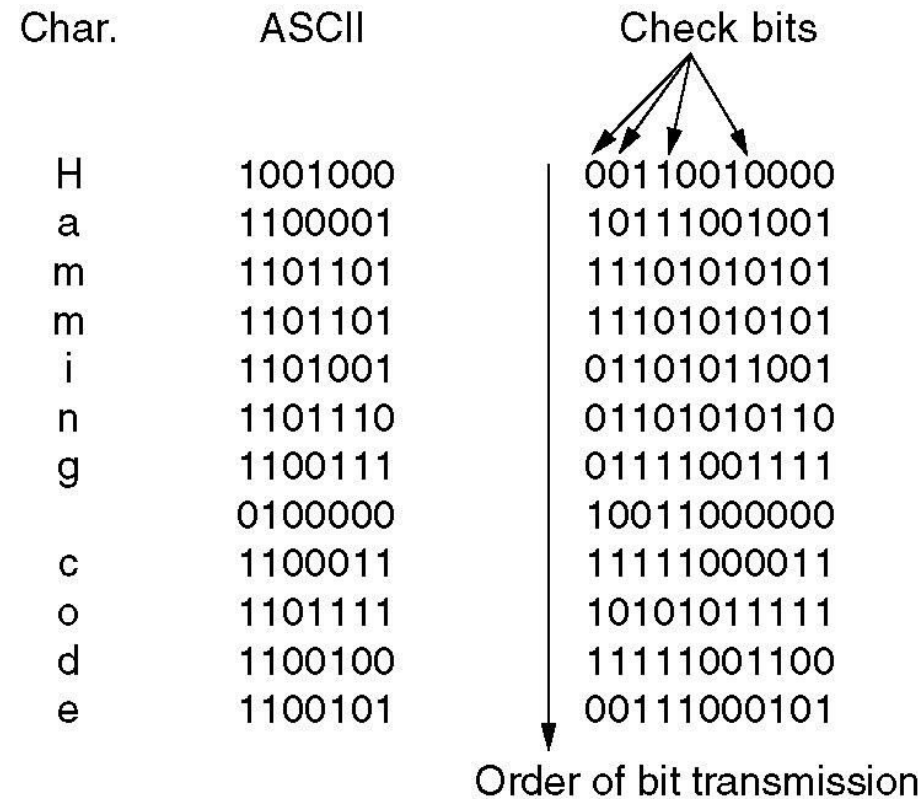
$$\begin{array}{ccccccc} 10100011 & \oplus & 10100011 & \oplus & 00000011 & = & 00000011 \\ \text{Sample 1} & & \text{Sample 2} & & \text{Sample 3} & & \text{Redundant Code} \\ & & \text{Lost!} & & & & \end{array}$$

$$\begin{array}{ccccccc} 10100011 & \oplus & 00000011 & \oplus & 00000011 & = & 10100011 \\ \text{Sample 1} & & \text{Redundant Code} & & \text{Sample 3} & & \text{Reconstructed} \\ & & & & & & \text{Sample 2} \end{array}$$

- Example: 3-way XOR
- n too large => higher prob. of packet loss.
- n too small => higher transmission rate and delay.

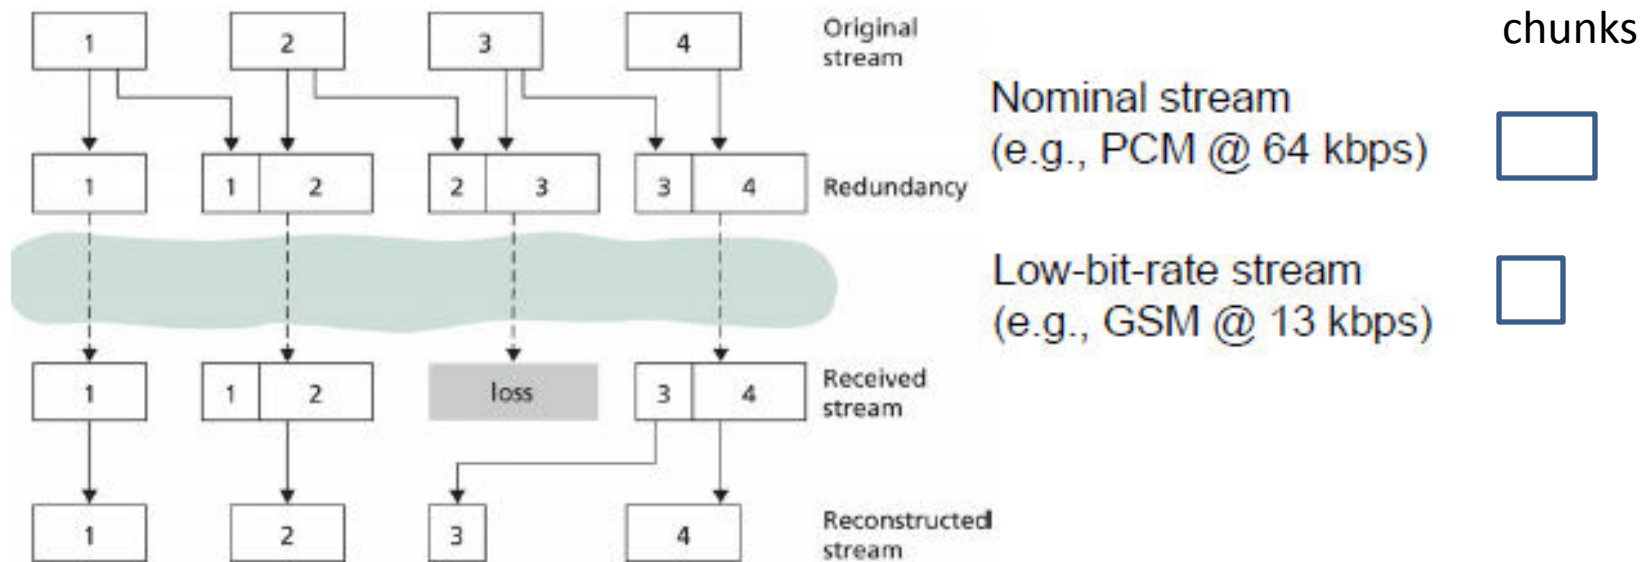
Hamming code to correct burst errors

- $\lceil \log_2 m \rceil$ check bits ($m=11$)
- Each check bit is added to yield zero XOR (even parity)
- A bit is checked by those check bits occurring in its expansion
 - E.g.
 - Data bit in position $7=1+2+4$ is checked by check bits in positions 1,2,4
 - Check bits 1,2,4 not satisfying even parity \Rightarrow data bit in position 7 is in error



FEC Method 2: Sending Lowbit-Rate Encoded Chunks

- Occasional low-bit-rate chunks in between high-bit-rate chunks yields higher overall audio quality in the face of packet losses.




Other FEC issues

- Obvious problem with FEC?
 - Could make congestion (and thus loss) worse.
- If bandwidth is to be held constant, then original media rate is reduced resulting in terms of lower original media quality.

Interleaving

- Interleave parts of consecutive chunks to get new units
- If a unit is lost, chunks are at least partially received.



- Multiple small gaps as opposed to a single large gap!
 - Can interpolate to close small gaps
 - Low overhead, but increase in latency.
 - Not for Internet phone, but for streaming stored video.